



Red Hat OpenShift Data Foundation 4.21

Configuring OpenShift Data Foundation Disaster Recovery for OpenShift Workloads

The OpenShift Data Foundation Disaster Recovery capabilities for Metropolitan and Regional regions is now General Available which also includes Disaster Recovery with stretch cluster.

Red Hat OpenShift Data Foundation 4.21 Configuring OpenShift Data Foundation Disaster Recovery for OpenShift Workloads

The OpenShift Data Foundation Disaster Recovery capabilities for Metropolitan and Regional regions is now General Available which also includes Disaster Recovery with stretch cluster.

Legal Notice

Copyright © Red Hat.

Except as otherwise noted below, the text of and illustrations in this documentation are licensed by Red Hat under the Creative Commons Attribution–Share Alike 3.0 Unported license . If you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, the Red Hat logo, JBoss, Hibernate, and RHCE are trademarks or registered trademarks of Red Hat, LLC. or its subsidiaries in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

XFS is a trademark or registered trademark of Hewlett Packard Enterprise Development LP or its subsidiaries in the United States and other countries.

The OpenStack[®] Word Mark and OpenStack logo are trademarks or registered trademarks of the Linux Foundation, used under license.

All other trademarks are the property of their respective owners.

Abstract

The intent of this solution guide is to detail the steps necessary to deploy OpenShift Data Foundation for disaster recovery with Advanced Cluster Management and stretch cluster to achieve a highly available storage infrastructure.

Table of Contents

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION	5
CHAPTER 1. INTRODUCTION TO OPENSIFT DATA FOUNDATION DISASTER RECOVERY	6
CHAPTER 2. DISASTER RECOVERY SUBSCRIPTION REQUIREMENT	7
CHAPTER 3. METRO-DR SOLUTION FOR OPENSIFT DATA FOUNDATION	8
3.1. COMPONENTS OF METRO-DR SOLUTION	8
3.2. UNSUPPORTED FEATURES FOR METRO-DR	9
3.3. METRO-DR DEPLOYMENT WORKFLOW	9
3.4. REQUIREMENTS FOR ENABLING METRO-DR	10
3.5. REQUIREMENTS FOR DEPLOYING RED HAT CEPH STORAGE STRETCH CLUSTER WITH ARBITER	12
3.5.1. Hardware requirements	12
3.5.2. Software requirements	13
3.5.3. Network configuration requirements	13
3.6. DEPLOYING RED HAT CEPH STORAGE	13
3.6.1. Node pre-deployment steps	13
3.6.2. Cluster bootstrapping and service deployment with cephadm utility	17
3.6.3. Configuring Red Hat Ceph Storage stretch mode	23
3.7. INSTALLING OPENSIFT DATA FOUNDATION ON MANAGED CLUSTERS	27
3.8. INSTALLING OPENSIFT DATA FOUNDATION MULTICLUSTER ORCHESTRATOR OPERATOR	30
3.9. CONFIGURING SSL ACCESS ACROSS CLUSTERS	31
3.10. CREATING DISASTER RECOVERY POLICY ON HUB CLUSTER	33
3.11. CONFIGURE DRCLUSTERS FOR FENCING AUTOMATION	35
3.11.1. Add node IP addresses to DRClusters	35
3.11.2. Add fencing annotations to DRClusters	36
3.12. CREATE SAMPLE APPLICATION FOR TESTING DISASTER RECOVERY SOLUTION	37
3.12.1. ApplicationSet-based applications	37
3.12.1.1. Creating ApplicationSet-based applications	37
3.12.1.2. Apply Data policy to sample ApplicationSet-based application	38
3.12.2. Deleting sample application	39
3.13. SUBSCRIPTION-BASED APPLICATION FAILOVER BETWEEN MANAGED CLUSTERS	40
3.14. APPLICATIONSET-BASED APPLICATION FAILOVER BETWEEN MANAGED CLUSTERS	42
3.15. RELOCATING SUBSCRIPTION-BASED APPLICATION BETWEEN MANAGED CLUSTERS	44
3.16. RELOCATING AN APPLICATIONSET-BASED APPLICATION BETWEEN MANAGED CLUSTERS	46
3.17. HUB RECOVERY USING RED HAT ADVANCED CLUSTER MANAGEMENT [TECHNOLOGY PREVIEW]	49
3.17.1. Configuring passive hub cluster	49
3.17.2. Switching to passive hub cluster	50
CHAPTER 4. REGIONAL-DR SOLUTION FOR OPENSIFT DATA FOUNDATION	53
4.1. COMPONENTS OF REGIONAL-DR SOLUTION	53
4.2. HYBRID CLOUD REGIONAL-DR SOLUTION	54
4.3. REGIONAL-DR DEPLOYMENT WORKFLOW	54
4.4. REQUIREMENTS FOR ENABLING REGIONAL-DR	55
4.4.1. Prerequisite for MTV-Migrated Virtual Machines	58
4.5. CREATING AN OPENSIFT DATA FOUNDATION CLUSTER ON MANAGED CLUSTERS	59
4.6. INSTALLING OPENSIFT DATA FOUNDATION MULTICLUSTER ORCHESTRATOR OPERATOR	61
4.6.1. [Optional] Connecting storage clusters with the ocs-provider-server ServiceExport	62
4.7. CONFIGURING SSL ACCESS ACROSS CLUSTERS	62
4.8. CREATING DISASTER RECOVERY POLICY ON HUB CLUSTER	64
4.9. DISASTER RECOVERY PROTECTION FOR CLONED AND RESTORED RBD VOLUMES	67
4.10. CREATE SAMPLE APPLICATION FOR TESTING DISASTER RECOVERY SOLUTION	67

4.10.1. ApplicationSet-based applications	68
4.10.1.1. Creating ApplicationSet-based applications	68
4.10.1.2. Apply Data policy to sample ApplicationSet-based application	69
4.10.2. Deleting sample application	72
4.11. VIRTUAL MACHINE GARBAGE COLLECTION AUTOMATION	73
4.12. SUBSCRIPTION-BASED APPLICATION FAILOVER BETWEEN MANAGED CLUSTERS	74
4.13. APPLICATIONSET-BASED APPLICATION FAILOVER BETWEEN MANAGED CLUSTERS	76
4.14. RELOCATING SUBSCRIPTION-BASED APPLICATION BETWEEN MANAGED CLUSTERS	78
4.15. RELOCATING AN APPLICATIONSET-BASED APPLICATION BETWEEN MANAGED CLUSTERS	79
4.16. DISASTER RECOVERY PROTECTION FOR DISCOVERED APPLICATIONS	81
4.16.1. Prerequisites for disaster recovery protection of discovered applications	81
4.16.2. Creating a sample discovered application	84
4.16.3. Enrolling a sample discovered application for disaster recovery protection	85
4.16.4. Discovered application failover and relocate	89
4.16.4.1. Failover disaster recovery protected discovered application	89
4.16.4.2. Relocate disaster recovery protected discovered application	90
4.16.5. Disable disaster recovery for protected applications	92
4.17. RECOVERING TO A REPLACEMENT CLUSTER WITH REGIONAL-DR	92
4.18. VIEWING RECOVERY POINT OBJECTIVE VALUES FOR DISASTER RECOVERY ENABLED APPLICATIONS	99
4.19. HUB RECOVERY USING RED HAT ADVANCED CLUSTER MANAGEMENT	100
4.19.1. Configuring passive hub cluster	100
4.19.2. Switching to passive hub cluster	101
4.20. ENABLING GRANULAR DISASTER RECOVERY FOR INDIVIDUAL OR GROUPS OF VIRTUAL MACHINES IN A NAMESPACE	103
4.21. ENABLING GRANULAR DISASTER RECOVERY FOR VIRTUAL MACHINES USING THE CLI	104
4.21.1. Key points about granular virtual machine protection	108
CHAPTER 5. DISASTER RECOVERY WITH STRETCH CLUSTER FOR OPENSIFT DATA FOUNDATION	109
5.1. REQUIREMENTS FOR ENABLING STRETCH CLUSTER	110
5.2. APPLYING TOPOLOGY ZONE LABELS TO OPENSIFT CONTAINER PLATFORM NODES	110
5.3. INSTALLING LOCAL STORAGE OPERATOR	111
5.4. INSTALLING RED HAT OPENSIFT DATA FOUNDATION OPERATOR	111
5.5. CREATING OPENSIFT DATA FOUNDATION CLUSTER	113
5.5.1. Creating NetworkFenceClass	117
5.6. VERIFYING OPENSIFT DATA FOUNDATION DEPLOYMENT	118
5.6.1. Verifying the state of the pods	118
5.6.2. Verifying the OpenShift Data Foundation cluster is healthy	121
5.6.3. Verifying the Multicloud Object Gateway is healthy	121
5.6.4. Verifying that the specific storage classes exist	122
5.7. INSTALL ZONE AWARE SAMPLE APPLICATION	122
5.7.1. Scaling the application after installation	124
5.7.2. Modify Deployment to be Zone Aware	125
5.8. RECOVERING OPENSIFT DATA FOUNDATION STRETCH CLUSTER	128
5.8.1. Understanding zone failure	129
5.8.2. Recovering zone-aware HA applications with RWX storage	129
5.8.3. Recovering HA applications with RWX storage	130
5.8.4. Recovering applications with RWO storage	130
5.8.5. Recovering StatefulSet pods	130
5.8.6. Fencing and tainting nodes to prevent volume corruption during non-graceful node shutdown	130
CHAPTER 6. DISABLING DISASTER RECOVERY FOR A DISASTER RECOVERY ENABLED APPLICATION	133
6.1. DISABLING DR MANAGED APPLICATIONS	133

6.2. DISABLING DR DISCOVERED APPLICATIONS	133
CHAPTER 7. UNINSTALLING DISASTER RECOVERY FROM HUB AND MANAGED CLUSTERS	135
7.1. UNINSTALLING DISASTER RECOVERY FROM HUB AND MANAGED CLUSTERS	135
7.1.1. Before uninstalling disaster recovery	135
7.1.2. Uninstalling disaster recovery for peer managed clusters	136
7.1.3. Uninstalling disaster recovery operators	137
CHAPTER 8. MONITORING DISASTER RECOVERY HEALTH	139
8.1. ENABLE MONITORING FOR DISASTER RECOVERY	139
8.2. ENABLING DISASTER RECOVERY DASHBOARD ON HUB CLUSTER	139
8.3. VIEWING HEALTH STATUS OF DISASTER RECOVERY REPLICATION RELATIONSHIPS	140
8.4. VIEWING HEALTH STATUS OF APPLICATIONSET-BASED APPLICATIONS	140
8.5. DISASTER RECOVERY METRICS	141
8.5.1. Last synchronization timestamp in seconds	141
8.5.2. Policy schedule interval in seconds	142
8.5.3. Last synchronization duration in seconds	142
8.5.4. Total bytes transferred from most recent synchronization	142
8.5.5. Workload protection status	143
8.6. DISASTER RECOVERY ALERTS	143
CHAPTER 9. TROUBLESHOOTING DISASTER RECOVERY	146
9.1. TROUBLESHOOTING METRO-DR	146
9.1.1. A statefulset application stuck after failover	146
9.1.2. DR policies protect all applications in the same namespace	146
9.1.3. Relocate or fallback might be stuck in Initiating state	147
9.2. TROUBLESHOOTING REGIONAL-DR	147
9.2.1. rbd-mirror daemon health is in warning state	147
9.2.2. volsync-rsync-src pod is in error state as it is unable to resolve the destination hostname	148
9.2.3. Cleanup and data sync for ApplicationSet workloads remain stuck after older primary managed cluster is recovered post failover	148
9.2.4. Secondary transition failed as PVC is potentially in use by a pod	150
9.2.5. vm-recipe related known issues	150
9.2.5.1. ramendr.openshift.io/k8s-resource-selector label is not automatically applied when a PVC is added to an already protected workload NS	150
9.2.5.2. DRPolicy without flattening applied to a cloned volume leads to errors	150
9.3. TROUBLESHOOTING 2-SITE STRETCH CLUSTER WITH ARBITER	150
9.3.1. Recovering workload pods stuck in ContainerCreating state post zone recovery	151

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION

We appreciate your input on our documentation. Do let us know how we can make it better.

To give feedback, create a Jira ticket:

1. Log in to the [Jira](#).
2. Click **Create** in the top navigation bar
3. Enter a descriptive title in the Summary field.
4. Enter your suggestion for improvement in the Description field. Include links to the relevant parts of the documentation.
5. Select **Documentation** in the Components field.
6. Click Create at the bottom of the dialogue.

CHAPTER 1. INTRODUCTION TO OPENSIFT DATA FOUNDATION DISASTER RECOVERY

Disaster recovery (DR) is the ability to recover and continue business critical applications from natural or human created disasters. It is a component of the overall business continuance strategy of any major organization as designed to preserve the continuity of business operations during major adverse events.

The OpenShift Data Foundation DR capability enables DR across multiple Red Hat OpenShift Container Platform clusters, and is categorized as follows:

- **Metro-DR**
Metro-DR ensures business continuity during the unavailability of a data center with no data loss. In the public cloud these would be similar to protecting from an Availability Zone failure.
- **Regional-DR**
Regional-DR ensures business continuity during the unavailability of a geographical region, accepting some loss of data in a predictable amount. In the public cloud this would be similar to protecting from a region failure.
- **Disaster Recovery with stretch cluster**
Stretch cluster solution ensures business continuity with no-data loss disaster recovery protection with OpenShift Data Foundation based synchronous replication in a single OpenShift cluster, stretched across two data centers with low latency and one arbiter node.

Zone failure in Metro-DR and region failure in Regional-DR is usually expressed using the terms, **Recovery Point Objective (RPO)** and **Recovery Time Objective (RTO)**

- **RPO** is a measure of how frequently you take backups or snapshots of persistent data. In practice, the RPO indicates the amount of data that will be lost or need to be reentered after an outage.
- **RTO** is the amount of downtime a business can tolerate. The RTO answers the question, "How long can it take for our system to recover after we are notified of a business disruption?"

The intent of this guide is to detail the Disaster Recovery steps and commands necessary to be able to failover an application from one OpenShift Container Platform cluster to another and then relocate the same application to the original primary cluster.

CHAPTER 2. DISASTER RECOVERY SUBSCRIPTION REQUIREMENT

Disaster Recovery features supported by Red Hat OpenShift Data Foundation require all of the following prerequisites to successfully implement a disaster recovery solution:

- A valid Red Hat OpenShift Data Foundation Advanced entitlement
- A valid Red Hat Advanced Cluster Management for Kubernetes subscription

Any Red Hat OpenShift Data Foundation Cluster containing PVs participating in active replication either as a source or destination requires OpenShift Data Foundation Advanced entitlement. This subscription should be active on both source and destination clusters.

To know how subscriptions for OpenShift Data Foundation work, see [knowledgebase article on OpenShift Data Foundation subscriptions](#).



IMPORTANT

OpenShift Data Foundation deployed with Multus networking is not supported for Regional Disaster Recovery (Regional-DR) setups.

CHAPTER 3. METRO-DR SOLUTION FOR OPENSIFT DATA FOUNDATION

This section of the guide provides details of the Metro Disaster Recovery (Metro-DR) steps and commands necessary to be able to failover an application from one OpenShift Container Platform cluster to another and then failback the same application to the original primary cluster. In this case the OpenShift Container Platform clusters will be created or imported using Red Hat Advanced Cluster Management (RHACM) and have distance limitations between the OpenShift Container Platform clusters of less than 10ms RTT latency.

The persistent storage for applications is provided by an external Red Hat Ceph Storage (RHCS) cluster stretched between the two locations with the OpenShift Container Platform instances connected to this storage cluster. An arbiter node with a storage monitor service is required at a third location (different location than where OpenShift Container Platform instances are deployed) to establish quorum for the RHCS cluster in the case of a site outage. This third location can be in the range of ~100ms RTT from the storage cluster connected to the OpenShift Container Platform instances.

This is a general overview of the Metro DR steps required to configure and execute OpenShift Disaster Recovery (ODR) capabilities using OpenShift Data Foundation and RHACM across two distinct OpenShift Container Platform clusters separated by distance. In addition to these two clusters called managed clusters, a third OpenShift Container Platform cluster is required that will be the Red Hat Advanced Cluster Management (RHACM) hub cluster.



IMPORTANT

You can now easily set up Metropolitan disaster recovery solutions for workloads based on OpenShift virtualization technology using OpenShift Data Foundation. For more information, see the [knowledgebase article](#).

3.1. COMPONENTS OF METRO-DR SOLUTION

Metro-DR is composed of Red Hat Advanced Cluster Management for Kubernetes, Red Hat Ceph Storage and OpenShift Data Foundation components to provide application and data mobility across OpenShift Container Platform clusters.

Red Hat Advanced Cluster Management for Kubernetes

Red Hat Advanced Cluster Management (RHACM) provides the ability to manage multiple clusters and application lifecycles. Hence, it serves as a control plane in a multi-cluster environment.

RHACM is split into two parts:

- RHACM Hub: components that run on the multi-cluster control plane.
- Managed clusters: components that run on the clusters that are managed.

For more information about this product, see [RHACM documentation](#) and the [RHACM “Manage Applications” documentation](#).

Red Hat Ceph Storage

Red Hat Ceph Storage is a massively scalable, open, software-defined storage platform that combines the most stable version of the Ceph storage system with a Ceph management platform, deployment utilities, and support services. It significantly lowers the cost of storing enterprise data and helps organizations manage exponential data growth. The software is a robust and modern petabyte-scale storage platform for public or private cloud deployments.

For more product information, see [Red Hat Ceph Storage](#) .

OpenShift Data Foundation

OpenShift Data Foundation provides the ability to provision and manage storage for stateful applications in an OpenShift Container Platform cluster. It is backed by Ceph as the storage provider, whose lifecycle is managed by Rook in the OpenShift Data Foundation component stack and Ceph-CSI provides the provisioning and management of Persistent Volumes for stateful applications.

OpenShift DR

OpenShift DR is a disaster recovery orchestrator for stateful applications across a set of peer OpenShift clusters which are deployed and managed using RHACM and provides cloud-native interfaces to orchestrate the life-cycle of an application's state on Persistent Volumes. These include:

- Protecting an application and its state relationship across OpenShift clusters
- Failing over an application and its state to a peer cluster
- Relocate an application and its state to the previously deployed cluster

OpenShift DR is split into three components:

- **ODF Multicluster Orchestrator:** Installed on the multi-cluster control plane (RHACM Hub), it orchestrates configuration and peering of OpenShift Data Foundation clusters for Metro and Regional DR relationships.
- **OpenShift DR Hub Operator:** Automatically installed as part of ODF Multicluster Orchestrator installation on the hub cluster to orchestrate failover or relocation of DR enabled applications.
- **OpenShift DR Cluster Operator:** Automatically installed on each managed cluster that is part of a Metro and Regional DR relationship to manage the lifecycle of all PVCs of an application.

3.2. UNSUPPORTED FEATURES FOR METRO-DR

The following features related to external mode are not supported when using Metro-DR:

- StorageClasses using non-default RADOS namespace
- User created StorageClasses, even when using default RADOS namespace
- Multiple StorageClasses

3.3. METRO-DR DEPLOYMENT WORKFLOW

This section provides an overview of the steps required to configure and deploy Metro-DR capabilities using the latest versions of Red Hat OpenShift Data Foundation, Red Hat Ceph Storage (RHCS) and Red Hat Advanced Cluster Management for Kubernetes (RHACM) version 2.10 or later, across two distinct OpenShift Container Platform clusters. In addition to two managed clusters, a third OpenShift Container Platform cluster will be required to deploy the Advanced Cluster Management.

To configure your infrastructure, perform the below steps in the order given:

1. Ensure requirements across the Hub, Primary and Secondary Openshift Container Platform clusters that are part of the DR solution are met. See [Requirements for enabling Metro-DR](#).

2. Ensure you meet the requirements for deploying Red Hat Ceph Storage stretch cluster with arbiter. See [Requirements for deploying Red Hat Ceph Storage](#) .
3. Deploy and configure Red Hat Ceph Storage stretch mode. For instructions on enabling Ceph cluster on two different data centers using stretched mode functionality, see [Deploying Red Hat Ceph Storage](#).
4. Install OpenShift Data Foundation operator and create a storage system on Primary and Secondary managed clusters. See [Installing OpenShift Data Foundation on managed clusters](#) .
5. Install the ODF Multiclustler Orchestrator on the Hub cluster. See [Installing ODF Multiclustler Orchestrator on Hub cluster](#).
6. Configure SSL access between the Hub, Primary and Secondary clusters. See [Configuring SSL access across clusters](#).
7. Create a DRPolicy resource for use with applications requiring DR protection across the Primary and Secondary clusters. See [Creating Disaster Recovery Policy on Hub cluster](#) .

**NOTE**

The Metro-DR solution can only have one DRpolicy.

8. Testing your disaster recovery solution with:
 - a. **Subscription-based** application:
 - Create sample applications. See [Creating sample application](#) .
 - Test failover and relocate operations using the sample application between managed clusters. See [Subscription-based application failover](#) and [relocating subscription-based application](#).
 - b. **ApplicationSet-based** application:
 - Create sample applications. See [Creating ApplicationSet-based applications](#) .
 - Test failover and relocate operations using the sample application between managed clusters. See [ApplicationSet-based application failover](#) and [relocating ApplicationSet-based application](#).
 - c. **Discovered applications**
 - Ensure all requirements mentioned in Prerequisites is addressed. See [Prerequisites for disaster recovery protection of discovered applications](#)
 - Create a sample discovered application. See [Creating a sample discovered application](#)
 - Enroll the discovered application. See [Enrolling a sample discovered application for disaster recovery protection](#)
 - Test failover and relocate. See [Discovered application failover and relocate](#)

3.4. REQUIREMENTS FOR ENABLING METRO-DR

The prerequisites to installing a disaster recovery solution supported by Red Hat OpenShift Data Foundation are as follows:

- You must have the following OpenShift clusters that have network reachability between them:
 - **Hub cluster** where Red Hat Advanced Cluster Management (RHACM) for Kubernetes operator are installed.
 - **Primary managed cluster** where OpenShift Data Foundation is running.
 - **Secondary managed cluster** where OpenShift Data Foundation is running.



NOTE

For configuring hub recovery setup, you need a 4th cluster which acts as the passive hub. The primary managed cluster (Site-1) can be co-situated with the active RHACM hub cluster while the passive hub cluster is situated along with the secondary managed cluster (Site-2). Alternatively, the active RHACM hub cluster can be placed in a neutral site (Site-3) that is not impacted by the failures of either of the primary managed cluster at Site-1 or the secondary cluster at Site-2. In this situation, if a passive hub cluster is used it can be placed with the secondary cluster at Site-2. For more information, see [Configuring passive hub cluster for hub recovery](#).

Hub recovery is a Technology Preview feature and is subject to Technology Preview support limitations.

- Ensure that RHACM operator and MultiClusterHub is installed on the Hub cluster. See [RHACM installation guide](#) for instructions.
After the operator is successfully installed, the web console automatically reloads to apply the changes. During this process, a temporary error message might appear on the page and this is expected and disappears after the refresh completes.



IMPORTANT

Ensure that application traffic routing and redirection are configured appropriately.

- On the Hub cluster
 - Navigate to **All Clusters → Infrastructure → Clusters**.
 - Import or create the **Primary managed cluster** and the **Secondary managed cluster** using the RHACM console.
 - Choose the appropriate options for your environment.

After the managed clusters are successfully created or imported, you can see the list of clusters that were imported or created on the console. For instructions, see [Creating a cluster](#) and [Importing a target managed cluster to the hub cluster](#).

**WARNING**

The OpenShift Container Platform managed clusters and the Red Hat Ceph Storage (RHCS) nodes have distance limitations. The network latency between the sites must be below 10 milliseconds round-trip time (RTT).


3.5. REQUIREMENTS FOR DEPLOYING RED HAT CEPH STORAGE STRETCH CLUSTER WITH ARBITER

Red Hat Ceph Storage is an open-source enterprise platform that provides unified software-defined storage on standard, economical servers and disks. With block, object, and file storage combined into one platform, Red Hat Ceph Storage efficiently and automatically manages all your data, so you can focus on the applications and workloads that use it.

This section provides a basic overview of the Red Hat Ceph Storage deployment. For more complex deployment, refer to the [official documentation guide for Red Hat Ceph Storage](#).

**NOTE**

Only Flash media is supported since it runs with **min_size=1** when degraded. Use stretch mode only with all-flash OSDs. Using all-flash OSDs minimizes the time needed to recover once connectivity is restored, thus minimizing the potential for data loss.

**IMPORTANT**

Erasured coded pools cannot be used with stretch mode.

3.5.1. Hardware requirements

For information on minimum hardware requirements for deploying Red Hat Ceph Storage, see [Minimum hardware recommendations for containerized Ceph](#).

Table 3.1. Physical server locations and Ceph component layout for Red Hat Ceph Storage cluster deployment:

Node name	Datacenter	Ceph components
ceph1	DC1	OSD+MON+MGR
ceph2	DC1	OSD+MON
ceph3	DC1	OSD+MDS+RGW
ceph4	DC2	OSD+MON+MGR
ceph5	DC2	OSD+MON

Node name	Datacenter	Ceph components
ceph6	DC2	OSD+MDS+RGW
ceph7	DC3	MON

3.5.2. Software requirements

Use the latest software version of **Red Hat Ceph Storage 8**

For more information on the supported Operating System versions for Red Hat Ceph Storage, see knowledgebase article on [Red Hat Ceph Storage: Supported configurations](#) .

3.5.3. Network configuration requirements

The recommended Red Hat Ceph Storage configuration is as follows:

- You must have two separate networks, one public network and one private network.
- You must have three different datacenters that support VLANS and subnets for Ceph's private and public network for all datacenters.



NOTE

You can use different subnets for each of the datacenters.

- The latencies between the two datacenters running the Red Hat Ceph Storage Object Storage Devices (OSDs) cannot exceed 10 ms RTT. For the **arbiter** datacenter, this was tested with values as high up to 100 ms RTT to the other two OSD datacenters.

Here is an example of a basic network configuration that we have used in this guide:

- **DC1: Ceph public/private network:**10.0.40.0/24
- **DC2: Ceph public/private network:**10.0.40.0/24
- **DC3: Ceph public/private network:**10.0.40.0/24

For more information on the required network environment, see [Ceph network configuration](#) .

3.6. DEPLOYING RED HAT CEPH STORAGE

3.6.1. Node pre-deployment steps

Before installing the Red Hat Ceph Storage Ceph cluster, perform the following steps to fulfill all the requirements needed.

1. Register all the nodes to the Red Hat Network or Red Hat Satellite and subscribe to a valid pool:

```
subscription-manager register
subscription-manager subscribe --pool=8a8XXXXXX9e0
```

2. Enable access for all the nodes in the Ceph cluster for the following repositories:

- **rhel9-for-x86_64-baseos-rpms**
- **rhel9-for-x86_64-appstream-rpms**

```
subscription-manager repos --disable="*" --enable="rhel9-for-x86_64-baseos-rpms" --
enable="rhel9-for-x86_64-appstream-rpms"
```

3. Update the operating system RPMs to the latest version and reboot if needed:

```
dnf update -y
reboot
```

4. Select a node from the cluster to be your bootstrap node. **ceph1** is our bootstrap node in this example going forward.

Only on the bootstrap node **ceph1**, enable the **ansible-2.9-for-rhel-9-x86_64-rpms** and **rhceph-6-tools-for-rhel-9-x86_64-rpms** repositories:

```
subscription-manager repos --enable="ansible-2.9-for-rhel-9-x86_64-rpms" --
enable="rhceph-6-tools-for-rhel-9-x86_64-rpms"
```

5. Configure the **hostname** using the bare/short hostname in all the hosts.

```
hostnamectl set-hostname <short_name>
```

6. Verify the hostname configuration for deploying Red Hat Ceph Storage with cephadm.

```
$ hostname
```

Example output:

```
ceph1
```

7. Modify `/etc/hosts` file and add the fqdn entry to the 127.0.0.1 IP by setting the DOMAIN variable with our DNS domain name.

```
DOMAIN="example.domain.com"

cat <<EOF >/etc/hosts
127.0.0.1 $(hostname).${DOMAIN} $(hostname) localhost localhost.localdomain localhost4
localhost4.localdomain4
::1 $(hostname).${DOMAIN} $(hostname) localhost6 localhost6.localdomain6
EOF
```

8. Check the long hostname with the **fqdn** using the **hostname -f** option.

```
$ hostname -f
```

Example output:

```
ceph1.example.domain.com
```

**NOTE**

To know more about why these changes are required, see [Fully Qualified Domain Names vs Bare Host Names](#).

9. Run the following steps on the bootstrap node. In our example, the bootstrap node is **ceph1**.

- a. Install the **cephadm-ansible** RPM package:

```
$ sudo dnf install -y cephadm-ansible
```

**IMPORTANT**

To run the ansible playbooks, you must have **ssh** passwordless access to all the nodes that are configured to the Red Hat Ceph Storage cluster. Ensure that the configured user (for example, **deployment-user**) has root privileges to invoke the **sudo** command without needing a password.

- b. To use a custom key, configure the selected user (for example, **deployment-user**) ssh config file to specify the id/key that will be used for connecting to the nodes via ssh:

```
cat <<EOF > ~/.ssh/config
Host ceph*
  User deployment-user
  IdentityFile ~/.ssh/ceph.pem
EOF
```

- c. Build the ansible inventory

```
cat <<EOF > /usr/share/cephadm-ansible/inventory
ceph1
ceph2
ceph3
ceph4
ceph5
ceph6
ceph7
[admin]
ceph1
ceph4
EOF
```

**NOTE**

Here, the Hosts (**Ceph1** and **Ceph4**) belonging to two different data centers are configured as part of the [admin] group on the inventory file and are tagged as **_admin** by **cephadm**. Each of these admin nodes receive the admin ceph keyring during the bootstrap process so that when one data center is down, we can check using the other available admin node.

- d. Verify that **ansible** can access all nodes using the ping module before running the pre-flight playbook.

—

```
$ ansible -i /usr/share/cephadm-ansible/inventory -m ping all -b
```

Example output:

```
ceph6 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/libexec/platform-python"
  },
  "changed": false,
  "ping": "pong"
}
ceph4 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/libexec/platform-python"
  },
  "changed": false,
  "ping": "pong"
}
ceph3 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/libexec/platform-python"
  },
  "changed": false,
  "ping": "pong"
}
ceph2 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/libexec/platform-python"
  },
  "changed": false,
  "ping": "pong"
}
ceph5 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/libexec/platform-python"
  },
  "changed": false,
  "ping": "pong"
}
ceph1 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/libexec/platform-python"
  },
  "changed": false,
  "ping": "pong"
}
ceph7 | SUCCESS => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/libexec/platform-python"
  },
  "changed": false,
  "ping": "pong"
}
```

- e. Navigate to the **/usr/share/cephadm-ansible** directory.

- f. Run ansible-playbook with relative file paths.

```
$ ansible-playbook -i /usr/share/cephadm-ansible/inventory /usr/share/cephadm-ansible/cephadm-preflight.yml --extra-vars "ceph_origin=rhcs"
```

The preflight playbook Ansible playbook configures the RHCS **dnf** repository and prepares the storage cluster for bootstrapping. It also installs podman, lvm2, chronyd, and cephadm. The default location for **cephadm-ansible** and **cephadm-preflight.yml** is **/usr/share/cephadm-ansible**. For additional information, see [Running the preflight playbook](#)

3.6.2. Cluster bootstrapping and service deployment with cephadm utility

The cephadm utility installs and starts a single Ceph Monitor daemon and a Ceph Manager daemon for a new Red Hat Ceph Storage cluster on the local node where the cephadm bootstrap command is run.

In this guide we are going to bootstrap the cluster and deploy all the needed Red Hat Ceph Storage services in one step using a cluster specification yaml file.

If you find issues during the deployment, it may be easier to troubleshoot the errors by dividing the deployment into two steps:

1. Bootstrap
2. Service deployment



NOTE

For additional information on the bootstrapping process, see [Bootstrapping a new storage cluster](#).

Procedure

1. Create json file to authenticate against the container registry using a json file as follows:

```
$ cat <<EOF > /root/registry.json
{
  "url":"registry.redhat.io",
  "username":"User",
  "password":"Pass"
}
EOF
```

2. Create a **cluster-spec.yaml** that adds the nodes to the Red Hat Ceph Storage cluster and also sets specific labels for where the services should run following table 3.1.

```
cat <<EOF > /root/cluster-spec.yaml
service_type: host
addr: 10.0.40.78 ## <XXX.XXX.XXX.XXX>
hostname: ceph1 ## <ceph-hostname-1>
location:
  root: default
  datacenter: DC1
labels:
```

```
- osd
- mon
- mgr
---
service_type: host
addr: 10.0.40.35
hostname: ceph2
location:
  datacenter: DC1
labels:
  - osd
  - mon
---
service_type: host
addr: 10.0.40.24
hostname: ceph3
location:
  datacenter: DC1
labels:
  - osd
  - mds
  - rgw
---
service_type: host
addr: 10.0.40.185
hostname: ceph4
location:
  root: default
  datacenter: DC2
labels:
  - osd
  - mon
  - mgr
---
service_type: host
addr: 10.0.40.88
hostname: ceph5
location:
  datacenter: DC2
labels:
  - osd
  - mon
---
service_type: host
addr: 10.0.40.66
hostname: ceph6
location:
  datacenter: DC2
labels:
  - osd
  - mds
  - rgw
---
service_type: host
addr: 10.0.40.221
hostname: ceph7
```

```

labels:
  - mon
---
service_type: mon
placement:
  label: "mon"
---
service_type: mds
service_id: cephfs
placement:
  label: "mds"
---
service_type: mgr
service_name: mgr
placement:
  label: "mgr"
---
service_type: osd
service_id: all-available-devices
service_name: osd.all-available-devices
placement:
  label: "osd"
spec:
  data_devices:
    all: true
---
service_type: rgw
service_id: objectgw
service_name: rgw.objectgw
placement:
  count: 2
  label: "rgw"
spec:
  rgw_frontend_port: 8080
EOF

```

3. Retrieve the IP for the NIC with the Red Hat Ceph Storage public network configured from the bootstrap node. After substituting **10.0.40.0** with the subnet that you have defined in your ceph public network, execute the following command.

```
$ ip a | grep 10.0.40
```

Example output:

```
10.0.40.78
```


4. Run the **cephadm** bootstrap command as the **root** user on the node that will be the initial Monitor node in the cluster. The **IP_ADDRESS** option is the node's IP address that you are using to run the **cephadm bootstrap** command.

**NOTE**

If you have configured a different user instead of **root** for passwordless SSH access, then use the **--ssh-user=** flag with the **cephadm bootstrap** command.

If you are using non default/id_rsa ssh key names, then use **--ssh-private-key** and **--ssh-public-key** options with **cephadm** command.

```
$ cephadm bootstrap --ssh-user=deployment-user --mon-ip 10.0.40.78 --apply-spec /root/cluster-spec.yaml --registry-json /root/registry.json
```

**IMPORTANT**

If the local node uses fully-qualified domain names (FQDN), then add the **--allow-fqdn-hostname** option to **cephadm bootstrap** on the command line.

Once the bootstrap finishes, you will see the following output from the previous cephadm bootstrap command:

You can access the Ceph CLI with:

```
sudo /usr/sbin/cephadm shell --fsid dd77f050-9afe-11ec-a56c-029f8148ea14 -c /etc/ceph/ceph.conf -k /etc/ceph/ceph.client.admin.keyring
```

Consider enabling telemetry to help improve Ceph:

```
ceph telemetry on
```

For more information see:

<https://docs.ceph.com/docs/pacific/mgr/telemetry/>

5. Verify the status of Red Hat Ceph Storage cluster deployment using the Ceph CLI client from ceph1:

```
$ ceph -s
```

Example output:

```
cluster:
  id: 3a801754-e01f-11ec-b7ab-005056838602
  health: HEALTH_OK

services:
  mon: 5 daemons, quorum ceph1,ceph2,ceph4,ceph5,ceph7 (age 4m)
  mgr: ceph1.khuuot(active, since 5m), standbys: ceph4.zotfsp
  osd: 12 osds: 12 up (since 3m), 12 in (since 4m)
  rgw: 2 daemons active (2 hosts, 1 zones)

data:
  pools: 5 pools, 107 pgs
```



```
objects: 191 objects, 5.3 KiB
usage: 105 MiB used, 600 GiB / 600 GiB avail
      105 active+clean
```



NOTE

It may take several minutes for all the services to start.

It is normal to get a global recovery event while you do not have any OSDs configured.

You can use **ceph orch ps** and **ceph orch ls** to further check the status of the services.

- Verify if all the nodes are part of the **cephadm** cluster.

```
$ ceph orch host ls
```

Example output:

```
HOST  ADDR      LABELS STATUS
ceph1 10.0.40.78 _admin osd mon mgr
ceph2 10.0.40.35 osd mon
ceph3 10.0.40.24 osd mds rgw
ceph4 10.0.40.185 osd mon mgr
ceph5 10.0.40.88 osd mon
ceph6 10.0.40.66 osd mds rgw
ceph7 10.0.40.221 mon
```



NOTE

You can run Ceph commands directly from the host because **ceph1** was configured in the **cephadm-ansible** inventory as part of the [admin] group. The Ceph admin keys were copied to the host during the **cephadm bootstrap** process.

- Check the current placement of the Ceph monitor services on the datacenters.

```
$ ceph orch ps | grep mon | awk '{print $1 " " $2}'
```

Example output:

```
mon.ceph1 ceph1
mon.ceph2 ceph2
mon.ceph4 ceph4
mon.ceph5 ceph5
mon.ceph7 ceph7
```

- Check the current placement of the Ceph manager services on the datacenters.

```
$ ceph orch ps | grep mgr | awk '{print $1 " " $2}'
```

Example output:

```
mgr.ceph2.ycgwyz ceph2
mgr.ceph5.kremtt ceph5
```

- Check the ceph osd crush map layout to ensure that each host has one OSD configured and its status is **UP**. Also, double-check that each node is under the right datacenter bucket as specified in table 3.1

```
$ ceph osd tree
```

Example output:

ID	CLASS	WEIGHT	TYPE	NAME	STATUS	REWEIGHT	PRI-AFF
-1		0.87900	root	default			
-16		0.43950	datacenter	DC1			
-11		0.14650	host	ceph1			
2	ssd	0.14650	osd.2	up	1.00000	1.00000	
-3		0.14650	host	ceph2			
3	ssd	0.14650	osd.3	up	1.00000	1.00000	
-13		0.14650	host	ceph3			
4	ssd	0.14650	osd.4	up	1.00000	1.00000	
-17		0.43950	datacenter	DC2			
-5		0.14650	host	ceph4			
0	ssd	0.14650	osd.0	up	1.00000	1.00000	
-9		0.14650	host	ceph5			
1	ssd	0.14650	osd.1	up	1.00000	1.00000	
-7		0.14650	host	ceph6			
5	ssd	0.14650	osd.5	up	1.00000	1.00000	

- Create and enable a new RDB block pool.

```
$ ceph osd pool create 32 32
$ ceph osd pool application enable rbdpool rbd
```



NOTE

The number 32 at the end of the command is the number of PGs assigned to this pool. The number of PGs can vary depending on several factors like the number of OSDs in the cluster, expected % used of the pool, etc. You can use the following calculator to determine the number of PGs needed: [Ceph Placement Groups \(PGs\) per Pool Calculator](#).

- Verify that the RBD pool has been created.

```
$ ceph osd lspools | grep rbdpool
```

Example output:

```
3 rbdpool
```

- Verify that MDS services are active and have located one service on each datacenter.

■

```
$ ceph orch ps | grep mds
```

Example output:

```
mds.cephfs.ceph3.cjpbqo  ceph3      running (17m)  117s ago  17m   16.1M   -
16.2.9
mds.cephfs.ceph6.lqmgqt  ceph6      running (17m)  117s ago  17m   16.1M   -
16.2.9
```

13. Create the CephFS volume.

```
$ ceph fs volume create cephfs
```



NOTE

The **ceph fs volume create** command also creates the needed data and meta CephFS pools. For more information, see [Configuring and Mounting Ceph File Systems](#).

14. Check the **Ceph** status to verify how the MDS daemons have been deployed. Ensure that the state is active where **ceph6** is the primary MDS for this filesystem and **ceph3** is the secondary MDS.

```
$ ceph fs status
```

Example output:

```
cephfs - 0 clients
=====
RANK STATE      MDS          ACTIVITY  DNS  INOS  DIRS  CAPS
0  active cephfs.ceph6.ggjywj Reqs: 0/s  10  13   12   0
    POOL      TYPE  USED AVAIL
cephfs.cephfs.meta metadata 96.0k 284G
cephfs.cephfs.data data    0 284G
    STANDBY MDS
cephfs.ceph3.ogcqkl
```

15. Verify that RGW services are active.

```
$ ceph orch ps | grep rgw
```

Example output:

```
rgw.objectgw.ceph3.kkxgb ceph3 *:8080 running (7m) 3m ago 7m 52.7M -
16.2.9
rgw.objectgw.ceph6.xmnpah ceph6 *:8080 running (7m) 3m ago 7m 53.3M -
16.2.9
```

3.6.3. Configuring Red Hat Ceph Storage stretch mode

Once the Red Hat Ceph Storage cluster is fully deployed using **cephadm**, use the following procedure to configure the stretch cluster mode. The new stretch mode is designed to handle the 2-site case.

Procedure

1. Check the current election strategy being used by the monitors with the `ceph mon dump` command. By default in a ceph cluster, the connectivity is set to classic.

```
ceph mon dump | grep election_strategy
```

Example output:

```
dumped monmap epoch 9
election_strategy: 1
```

2. Change the monitor election to connectivity.

```
ceph mon set election_strategy connectivity
```

3. Run the previous `ceph mon dump` command again to verify the `election_strategy` value.

```
$ ceph mon dump | grep election_strategy
```

Example output:

```
dumped monmap epoch 10
election_strategy: 3
```

To know more about the different election strategies, see [Configuring monitor election strategy](#).

4. Set the location for all our Ceph monitors:

```
ceph mon set_location ceph1 datacenter=DC1
ceph mon set_location ceph2 datacenter=DC1
ceph mon set_location ceph4 datacenter=DC2
ceph mon set_location ceph5 datacenter=DC2
ceph mon set_location ceph7 datacenter=DC3
```

5. Verify that each monitor has its appropriate location.

```
$ ceph mon dump
```

Example output:

```
epoch 17
fsid dd77f050-9afe-11ec-a56c-029f8148ea14
last_changed 2022-03-04T07:17:26.913330+0000
created 2022-03-03T14:33:22.957190+0000
min_mon_release 16 (pacific)
election_strategy: 3
0: [v2:10.0.143.78:3300/0,v1:10.0.143.78:6789/0] mon.ceph1; crush_location
{datacenter=DC1}
1: [v2:10.0.155.185:3300/0,v1:10.0.155.185:6789/0] mon.ceph4; crush_location
{datacenter=DC2}
2: [v2:10.0.139.88:3300/0,v1:10.0.139.88:6789/0] mon.ceph5; crush_location
{datacenter=DC2}
```

```
3: [v2:10.0.150.221:3300/0,v1:10.0.150.221:6789/0] mon.ceph7; crush_location
{datacenter=DC3}
4: [v2:10.0.155.35:3300/0,v1:10.0.155.35:6789/0] mon.ceph2; crush_location
{datacenter=DC1}
```

6. Create a CRUSH rule that makes use of this OSD crush topology by installing the **ceph-base** RPM package in order to use the **crushtool** command:

```
$ dnf -y install ceph-base
```

To know more about CRUSH ruleset, see [Ceph CRUSH ruleset](#).

7. Get the compiled CRUSH map from the cluster:

```
$ ceph osd getcrushmap > /etc/ceph/crushmap.bin
```

8. Decompile the CRUSH map and convert it to a text file in order to be able to edit it:

```
$ crushtool -d /etc/ceph/crushmap.bin -o /etc/ceph/crushmap.txt
```

9. Add the following rule to the CRUSH map by editing the text file **/etc/ceph/crushmap.txt** at the end of the file.

```
$ vim /etc/ceph/crushmap.txt
```

```
rule stretch_rule {
    id 1
    type replicated
    min_size 1
    max_size 10
    step take default
    step choose firstn 0 type datacenter
    step chooseleaf firstn 2 type host
    step emit
}
# end crush map
```

This example is applicable for active applications in both OpenShift Container Platform clusters.



NOTE

The rule **id** has to be unique. In the example, we only have one more crush rule with id 0 hence we are using id 1. If your deployment has more rules created, then use the next free id.

The CRUSH rule declared contains the following information:

- **Rule name**
 - Description: A unique whole name for identifying the rule.
 - Value: **stretch_rule**

- **id**
 - Description: A unique whole number for identifying the rule.
 - Value: **1**
- **type**
 - Description: Describes a rule for either a storage drive replicated or erasure-coded.
 - Value: **replicated**
- **min_size**
 - Description: If a pool makes fewer replicas than this number, CRUSH will not select this rule.
 - Value: 1
- **max_size**
 - Description: If a pool makes more replicas than this number, CRUSH will not select this rule.
 - Value: 10
- **step take default**
 - Description: Takes the root bucket called **default**, and begins iterating down the tree.
- **step choose firstn 0 type datacenter**
 - Description: Selects the datacenter bucket, and goes into its subtrees.
- **step chooseleaf firstn 2 type host**
 - Description: Selects the number of buckets of the given type. In this case, it is two different hosts located in the datacenter it entered at the previous level.
- **step emit**
 - Description: Outputs the current value and empties the stack. Typically used at the end of a rule, but may also be used to pick from different trees in the same rule.

10. Compile the new CRUSH map from the file **/etc/ceph/crushmap.txt** and convert it to a binary file called **/etc/ceph/crushmap2.bin**:

```
$ crushtool -c /etc/ceph/crushmap.txt -o /etc/ceph/crushmap2.bin
```

11. Inject the new crushmap we created back into the cluster:

```
$ ceph osd setcrushmap -i /etc/ceph/crushmap2.bin
```

Example output:

```
17
```

**NOTE**

The number 17 is a counter and it will increase (18,19, and so on) depending on the changes you make to the crush map.

12. Verify that the stretched rule created is now available for use.

```
ceph osd crush rule ls
```

Example output:

```
replicated_rule
stretch_rule
```

13. Enable the stretch cluster mode.

```
$ ceph mon enable_stretch_mode ceph7 stretch_rule datacenter
```

In this example, **ceph7** is the arbiter node, **stretch_rule** is the crush rule we created in the previous step and **datacenter** is the dividing bucket.

14. Verify all our pools are using the **stretch_rule** CRUSH rule we have created in our Ceph cluster:

```
$ for pool in $(rados lspools);do echo -n "Pool: ${pool}; ";ceph osd pool get ${pool}
crush_rule;done
```

Example output:

```
Pool: device_health_metrics; crush_rule: stretch_rule
Pool: cephfs.cephfs.meta; crush_rule: stretch_rule
Pool: cephfs.cephfs.data; crush_rule: stretch_rule
Pool: .rgw.root; crush_rule: stretch_rule
Pool: default.rgw.log; crush_rule: stretch_rule
Pool: default.rgw.control; crush_rule: stretch_rule
Pool: default.rgw.meta; crush_rule: stretch_rule
Pool: rbdpool; crush_rule: stretch_rule
```

This indicates that a working Red Hat Ceph Storage stretched cluster with arbiter mode is now available.

3.7. INSTALLING OPENSIFT DATA FOUNDATION ON MANAGED CLUSTERS

To configure storage replication between the two OpenShift Container Platform clusters, OpenShift Data Foundation operator must be installed first on each managed cluster.

Prerequisites

- Ensure that you have met the hardware requirements for OpenShift Data Foundation external deployments. For a detailed description of the hardware requirements, see [External mode requirements](#).

Procedure

1. Install and configure the latest **OpenShift Data Foundation** cluster on each of the managed clusters.
2. After installing the operator, create a StorageSystem using the option **Full deployment** type and **Connect with external storage platform** where your **Backing storage type** is **Red Hat Ceph Storage**.

For detailed instructions, refer to [Deploying OpenShift Data Foundation in external mode](#).

Use the following flags with the **ceph-external-cluster-details-exporter.py** script.

- a. At a minimum, you must use the following three flags with the **ceph-external-cluster-details-exporter.py** script:

--rbd-data-pool-name

With the name of the RBD pool that was created during RHCS deployment for OpenShift Container Platform. For example, the pool can be called **rbdpool**.

--rgw-endpoint

Provide the endpoint in the format **<ip_address>:<port>**. It is the RGW IP of the RGW daemon running on the same site as the OpenShift Container Platform cluster that you are configuring.

--run-as-user

With a different client name for each site.

- b. The following flags are **optional** if default values were used during the RHCS deployment:

--cephfs-filesystem-name

With the name of the CephFS filesystem we created during RHCS deployment for OpenShift Container Platform, the default filesystem name is **cephfs**.

--cephfs-data-pool-name

With the name of the CephFS data pool we created during RHCS deployment for OpenShift Container Platform, the default pool is called **cephfs.data**.

--cephfs-metadata-pool-name

With the name of the CephFS metadata pool we created during RHCS deployment for OpenShift Container Platform, the default pool is called **cephfs.meta**.

- c. Run the following command on the bootstrap node **ceph1**, to get the IP for the RGW endpoints in datacenter1 and datacenter2:

```
ceph orch ps | grep rgw.objectgw
```

Example output:

```
rgw.objectgw.ceph3.mecpzm ceph3 *:8080    running (5d)  31s ago  7w   204M
- 16.2.7-112.el8cp
rgw.objectgw.ceph6.mecpzm ceph6 *:8080    running (5d)  31s ago  7w   204M
- 16.2.7-112.el8cp
```

```
host ceph3.example.com
host ceph6.example.com
```


Example output:

```
ceph3.example.com has address 10.0.40.24
ceph6.example.com has address 10.0.40.66
```

- d. Run the **ceph-external-cluster-details-exporter.py** with the parameters that are configured for the first OpenShift Container Platform managed cluster **cluster1** on bootstrapped node **ceph1**.

```
python3 ceph-external-cluster-details-exporter.py --rbd-data-pool-name rbdpool --cephfs-filesystem-name cephfs --cephfs-data-pool-name cephfs.cephfs.data --cephfs-metadata-pool-name cephfs.cephfs.meta --<rgw-endpoint> XXX.XXX.XXX.XXX:8080 --run-as-user client.odf.cluster1 > ocp-cluster1.json
```



NOTE

Modify the <rgw-endpoint> XXX.XXX.XXX.XXX according to your environment.

- e. Run the **ceph-external-cluster-details-exporter.py** with the parameters that are configured for the first OpenShift Container Platform managed cluster **cluster2** on bootstrapped node **ceph1**.

```
python3 ceph-external-cluster-details-exporter.py --rbd-data-pool-name rbdpool --cephfs-filesystem-name cephfs --cephfs-data-pool-name cephfs.cephfs.data --cephfs-metadata-pool-name cephfs.cephfs.meta --rgw-endpoint XXX.XXX.XXX.XXX:8080 --run-as-user client.odf.cluster2 > ocp-cluster2.json
```



NOTE

Modify the <rgw-endpoint> XXX.XXX.XXX.XXX according to your environment.

- Save the two files generated in the bootstrap cluster (ceph1) **ocp-cluster1.json** and **ocp-cluster2.json** to your local machine.
 - Use the contents of file **ocp-cluster1.json** on the OpenShift Container Platform console on **cluster1** where external OpenShift Data Foundation is being deployed.
 - Use the contents of file **ocp-cluster2.json** on the OpenShift Container Platform console on **cluster2** where external OpenShift Data Foundation is being deployed.
3. Review the settings and then select **Create StorageSystem**.
 4. Validate the successful deployment of OpenShift Data Foundation on each managed cluster with the following command:

```
$ oc get storagecluster -n openshift-storage ocs-external-storagecluster -o jsonpath='{.status.phase}'
```

For the Multicloud Gateway (MCG):

```
$ oc get noobaa -n openshift-storage noobaa -o jsonpath='{.status.phase}'{"\n"}
```

Wait for the status result to be Ready for both queries on the **Primary managed cluster** and the **Secondary managed cluster**.

5. Verify the storage cluster is healthy.
 - a. On the OpenShift Web Console, navigate to **Storage → Data Foundation → Storage System**.
 - b. In the **Status** card of the **Overview** tab, click **Storage System** and then click the storage system link from the pop up that appears.
 - c. In the **Status** card of the **Block and File** tab, verify that the *Storage Cluster* has a green tick.
6. Enable read affinity for RBD and CephFS volumes to be served from the nearest datacenter.
 - a. On the Primary managed cluster, label all the nodes.

```
$ oc label nodes --all metro-dr.openshift-storage.topology.io/datacenter=DC1
```

Execute the following commands to enable read affinity:

```
$ oc patch storageclusters.ocs.openshift.io -n openshift-storage ocs-external-storagecluster -p '{"spec":{"csi":{"readAffinity":{"enabled":true,"crushLocationLabels":["metro-dr.openshift-storage.topology.io/datacenter"]}}}}' --type=merge
```

```
$ oc delete po -n openshift-storage -l 'app in (openshift-storage.cephfs.csi.ceph.com-ctrlplugin,openshift-storage.rbd.csi.ceph.com-ctrlplugin)'
```

- b. On the Secondary managed cluster, label all the nodes:

```
$ oc label nodes --all metro-dr.openshift-storage.topology.io/datacenter=DC2
```

Execute the following commands to enable read affinity:

```
$ oc patch storageclusters.ocs.openshift.io -n openshift-storage ocs-external-storagecluster -p '{"spec":{"csi":{"readAffinity":{"enabled":true,"crushLocationLabels":["metro-dr.openshift-storage.topology.io/datacenter"]}}}}' --type=merge
```

```
$ oc delete po -n openshift-storage -l 'app in (openshift-storage.cephfs.csi.ceph.com-ctrlplugin,openshift-storage.rbd.csi.ceph.com-ctrlplugin)'
```

3.8. INSTALLING OPENSIFT DATA FOUNDATION MULTICLUSTER ORCHESTRATOR OPERATOR

OpenShift Data Foundation Multicluster Orchestrator is a controller that is installed from OpenShift Container Platform's OperatorHub on the Hub cluster.

Prerequisites

- The **OpenShift GitOps** operator must be installed. For instructions, see [Installing Red Hat OpenShift GitOps Operator in web console](#).

Procedure

1. On the **Hub cluster**, navigate to **Ecosystem → Software Catalog** and use the keyword filter to search for **ODF Multicluster Orchestrator**.
2. Click **ODF Multicluster Orchestrator** tile.
3. Keep all default settings and click **Install**.
Ensure that the operator resources are installed in **openshift-operators** project and available to all namespaces.



NOTE

The **ODF Multicluster Orchestrator** also installs the **OpenShift DR Hub Operator** on the RHACM hub cluster as a dependency.

4. Verify that the operator **Pods** are in a **Running** state. The **OpenShift DR Hub operator** is also installed at the same time in **openshift-operators** namespace.

```
$ oc get pods -n openshift-operators
```

Example output:

NAME	READY	STATUS	RESTARTS	AGE
odf-multicluster-console-6845b795b9-blxrn	1/1	Running	0	4d20h
odfmo-controller-manager-f9d9dfb59-jbrsd	1/1	Running	0	4d20h
ramen-hub-operator-6fb887f885-fss4w	2/2	Running	0	4d20h

3.9. CONFIGURING SSL ACCESS ACROSS CLUSTERS

Configure network (SSL) access between the **primary** and **secondary clusters** so that metadata can be stored on the alternate cluster in a Multicloud Gateway (MCG) **object bucket** using a secure transport protocol and in the **Hub cluster** for verifying access to the object buckets.



NOTE

If all of your OpenShift clusters are deployed using a signed and valid set of certificates for your environment then this section can be skipped.

Procedure

1. Extract the ingress certificate for the Primary managed cluster and save the output to **primary.crt**.

```
$ oc get cm default-ingress-cert -n openshift-config-managed -o jsonpath="{['data']['ca-bundle.crt']}" > primary.crt
```

2. Extract the ingress certificate for the Secondary managed cluster and save the output to **secondary.crt**.

—

```
$ oc get cm default-ingress-cert -n openshift-config-managed -o jsonpath="{['data']['ca-bundle.crt']}" > secondary.crt
```

3. Create a new **ConfigMap** file to hold the remote cluster's certificate bundle with filename **cm-clusters-crt.yaml**.



NOTE

There could be more or less than three certificates for each cluster as shown in this example file. Also, ensure that the certificate contents are correctly indented after you copy and paste from the **primary.crt** and **secondary.crt** files that were created before.

```
apiVersion: v1
data:
  ca-bundle.crt: |
    -----BEGIN CERTIFICATE-----
    <copy contents of cert1 from primary.crt here>
    -----END CERTIFICATE-----

    -----BEGIN CERTIFICATE-----
    <copy contents of cert2 from primary.crt here>
    -----END CERTIFICATE-----

    -----BEGIN CERTIFICATE-----
    <copy contents of cert3 primary.crt here>
    -----END CERTIFICATE-----

    -----BEGIN CERTIFICATE-----
    <copy contents of cert1 from secondary.crt here>
    -----END CERTIFICATE-----

    -----BEGIN CERTIFICATE-----
    <copy contents of cert2 from secondary.crt here>
    -----END CERTIFICATE-----

    -----BEGIN CERTIFICATE-----
    <copy contents of cert3 from secondary.crt here>
    -----END CERTIFICATE-----
kind: ConfigMap
metadata:
  name: user-ca-bundle
  namespace: openshift-config
```

4. Create the **ConfigMap** on the **Primary managed cluster**, **Secondary managed cluster**, and the **Hub cluster**.

```
$ oc create -f cm-clusters-crt.yaml
```

Example output:

```
configmap/user-ca-bundle created
```

5. Patch default proxy resource on the **Primary managed cluster**, **Secondary managed cluster**, and the **Hub cluster**.

```
$ oc patch proxy cluster --type=merge --patch='{ "spec": { "trustedCA": { "name": "user-ca-bundle" } } }
```

Example output:

```
proxy.config.openshift.io/cluster patched
```

3.10. CREATING DISASTER RECOVERY POLICY ON HUB CLUSTER

Openshift Disaster Recovery Policy (DRPolicy) resource specifies OpenShift Container Platform clusters participating in the disaster recovery solution and the desired replication interval. DRPolicy is a cluster scoped resource that users can apply to applications that require Disaster Recovery solution.

The **ODF MultiCluster Orchestrator** Operator facilitates the creation of each **DRPolicy** and the corresponding **DRClusters** through the **Multicluster Web console**.

Prerequisites

- Ensure that there is a minimum set of two managed clusters.

Procedure


1. On the **OpenShift console**, navigate to **All Clusters** → **Data Services** → **Disaster recovery**.
2. On the Overview tab, click **Create a disaster recovery policy** or you can navigate to Policies tab and click **Create DRPolicy**.
3. Enter **Policy name**. Ensure that each DRPolicy has a unique name (for example: **ocp4perf1-ocp4perf2**).
4. Select two clusters from the list of managed clusters to which this new policy will be associated with.
5. **Replication policy** is automatically set to **sync** based on the OpenShift clusters selected.
6. Click **Create**.
7. Verify that the **DRPolicy** is created successfully. Run this command on the **Hub cluster** for each of the DRPolicy resources created, where `<drpolicy_name>` is replaced with your unique name.

```
$ oc get drpolicy <drpolicy_name> -o jsonpath='{.status.conditions[].reason}'
```

Example output:

```
Succeeded
```

When a DRPolicy is created, along with it, two DRCluster resources are also created. It could take up to 10 minutes for all three resources to be validated and for the status to show as **Succeeded**.

**NOTE**

Editing of **SchedulingInterval**, **ReplicationClassSelector**, **VolumeSnapshotClassSelector** and **DRClusters** field values are not supported in the DRPolicy.

8. Verify the object bucket access from the **Hub cluster** to both the **Primary managed cluster** and the **Secondary managed cluster**.

- a. Get the names of the **DRClusters** on the Hub cluster.

```
$ oc get drclusters
```

Example output:

```
NAME      AGE
ocp4perf1 4m42s
ocp4perf2 4m42s
```

- b. Check S3 access to each bucket created on each managed cluster. Use the DRCluster validation command, where `<drcluster_name>` is replaced with your unique name.

**NOTE**

Editing of **Region** and **S3ProfileName** field values are non supported in DRClusters.

```
$ oc get drcluster <drcluster_name> -o jsonpath='{.status.conditions[2].reason}'
```

Example output:

```
Succeeded
```

**NOTE**

Make sure to run commands for both **DRClusters** on the **Hub cluster**.

9. Verify that the **OpenShift DR Cluster operator** installation was successful on the **Primary managed cluster** and the **Secondary managed cluster**.

```
$ oc get csv,pod -n openshift-dr-system | egrep 'odr|ramen'
```

Example output:

```
clusterserviceversion.operators.coreos.com/odr-cluster-operator.v4.21.0-rhodef Openshift
DR Cluster Operator 4.21.0-rhodef odr-cluster-operator.v4.21.0-rhodef Succeeded
pod/ramen-dr-cluster-operator-77bf74849c-rfn85 2/2 Running 0 7h20m
```

You can also verify that **OpenShift DR Cluster Operator** is installed successfully in the **Ecosystem → Software Catalog** of each managed cluster.

10. Verify that the secret is propagated correctly on the Primary managed cluster and the Secondary managed cluster.

```
oc get secrets -n openshift-dr-system | grep Opaque
```

Match the output with the s3SecretRef from the Hub cluster:

```
oc get cm -n openshift-operators ramen-hub-operator-config -oyaml
```

3.11. CONFIGURE DRCLUSTERS FOR FENCING AUTOMATION

This configuration is required for enabling fencing prior to application failover. In order to prevent writes to the persistent volume from the cluster which is hit by a disaster, OpenShift DR instructs Red Hat Ceph Storage (RHCS) to fence the nodes of the cluster from the RHCS external storage. This section guides you on how to add the IPs or the IP Ranges for the nodes of the DRCluster.

3.11.1. Add node IP addresses to DRClusters

1. Find the **IP addresses** for all of the OpenShift nodes in the managed clusters by running this command in the **Primary managed cluster** and the **Secondary managed cluster**.

```
$ oc get nodes -o jsonpath='{range .items[*]}{.status.addresses[?(@.type=="ExternalIP")].address}{"\n"}{end}'
```

Example output:

```
10.70.56.118
10.70.56.193
10.70.56.154
10.70.56.242
10.70.56.136
10.70.56.99
```

Once you have the **IP addresses** then the **DRCluster** resources can be modified for each managed cluster.

2. Find the **DRCluster** names on the Hub Cluster.

```
$ oc get drcluster
```

Example output:

```
NAME      AGE
ocp4perf1 5m35s
ocp4perf2 5m35s
```

3. Edit each **DRCluster** to add your unique IP addresses after replacing **<drcluster_name>** with your unique name.

```
$ oc edit drcluster <drcluster_name>
```

```
apiVersion: ramendr.openshift.io/v1alpha1
```

```

kind: DRCluster
metadata:
[...]
spec:
  s3ProfileName: s3profile-<drcluster_name>-ocs-external-storagecluster
  ## Add this section
  cidrs:
    - <IP_Address1>/32
    - <IP_Address2>/32
    - <IP_Address3>/32
    - <IP_Address4>/32
    - <IP_Address5>/32
    - <IP_Address6>/32
  [...]

```

Example output:

```
drcluster.ramendr.openshift.io/ocp4perf1 edited
```



NOTE

There could be more than six IP addresses.

Modify this **DRCluster** configuration also for **IP addresses** on the **Secondary managed clusters** in the peer DRCluster resource (e.g., ocp4perf2).

3.11.2. Add fencing annotations to DRClusters

Add the following annotations to all the DRCluster resources. These annotations include details needed for the **NetworkFence** resource created later in these instructions (prior to testing application failover).



NOTE

Replace <drcluster_name> with your unique name.

```
$ oc edit drcluster <drcluster_name>
```

```

apiVersion: ramendr.openshift.io/v1alpha1
kind: DRCluster
metadata:
  ## Add this section
  annotations:
    drcluster.ramendr.openshift.io/storage-clusterid: openshift-storage
    drcluster.ramendr.openshift.io/storage-driver: openshift-storage.rbd.csi.ceph.com
    drcluster.ramendr.openshift.io/storage-secret-name: rook-csi-rbd-provisioner
    drcluster.ramendr.openshift.io/storage-secret-namespace: openshift-storage
  [...]

```

Example output:

```
drcluster.ramendr.openshift.io/ocp4perf1 edited
```


Make sure to add these annotations for both **DRCluster** resources (for example: **ocp4perf1** and **ocp4perf2**).

3.12. CREATE SAMPLE APPLICATION FOR TESTING DISASTER RECOVERY SOLUTION

OpenShift Data Foundation disaster recovery (DR) solution supports disaster recovery for ApplicationSet-based applications that are managed by RHACM. For more details, see [ApplicationSet documentation](#).

The following section details how to create an application and apply a DRPolicy to an application.

- [ApplicationSet-based applications](#)

OpenShift users that do not have cluster-admin permissions cannot create ApplicationSet-based applications.



IMPORTANT

Discovered applications using RBD block volumes and Ceph FileSystem are not supported when using Metro-DR.

3.12.1. ApplicationSet-based applications

3.12.1.1. Creating ApplicationSet-based applications

Prerequisite

- Ensure that the Red Hat OpenShift GitOps operator is installed on all three clusters: **Hub cluster**, **Primary managed cluster** and **Secondary managed cluster**. For instructions, see [Installing Red Hat OpenShift GitOps Operator in web console](#).
- On the Hub cluster, ensure that both Primary and Secondary managed clusters are registered to GitOps. For registration instructions, see [Registering managed clusters to GitOps](#). Then check if the Placement used by **GitOpsCluster** resource to register both managed clusters, has the tolerations to deal with cluster unavailability. You can verify if the following tolerations are added to the Placement using the command **oc get placement <placement-name> -n openshift-gitops -o yaml**.

```
tolerations:
- key: cluster.open-cluster-management.io/unreachable
  operator: Exists
- key: cluster.open-cluster-management.io/unavailable
  operator: Exists
```

In case the tolerations are not added, see [Configuring application placement tolerations for Red Hat Advanced Cluster Management and OpenShift GitOps](#).

- Ensure that you have created the **ClusterRoleBinding** yaml on both the **Primary** and **Secondary managed** clusters. For instruction, see the [Prerequisites chapter in RHACM documentation](#).

Procedure

1. On the Hub cluster, navigate to **All Clusters → Applications** and click **Create application**.
2. Choose the application type as **Argo CD ApplicationSet - Pull model**
3. In the **General** step, enter your **Application set name**.
4. Select **Argo server openshift-gitops** and **Requeue time** as **180** seconds.
5. Click **Next**.
6. In the Repository location for resources section, select **Repository type Git**.
7. Enter the Git repository URL for the sample application, the github Branch and Path where the resources busybox Pod and PVC will be created.
 - a. Use the sample application repository as <https://github.com/red-hat-storage/ocm-ramen-samples>
 - b. Select **Revision** as **main**
 - c. Choose Path as **workloads/deployment/odr-metro-rbd**.
8. Enter **Remote namespace** value. (example, busybox-sample) and click **Next**.
9. Choose the **Sync policy** settings as per your requirement or go with the default selections, and then click **Next**.
You can choose one or more options.
10. In Label expressions, add a label `<name>` with its value set to the **managed cluster** name.
11. Click **Next**.
12. Review the setting details and click **Submit**.

3.12.1.2. Apply Data policy to sample ApplicationSet-based application

Prerequisites

- Ensure that both managed clusters referenced in the Data policy are reachable. If not, the application will not be protected for disaster recovery until both clusters are online.

Procedure

1. On the Hub cluster, navigate to **All Clusters → Applications**.
2. Click the Actions menu at the end of application to view the list of available actions.
3. Click **Manage disaster recovery**.
4. Click **Enroll application**.
5. Select **Policy name** and click **Next**.
6. Select an **Application resource** and then use **PVC label selector** to select **PVC label** for the selected application resource.

**NOTE**

You can select more than one PVC label for the selected application resources.

7. Click **Next**.
8. In the **Enroll managed application** modal, review the policy configuration details and click **Assign**. The newly assigned Data policy details are displayed on the **Manage disaster recovery** modal.
9. Verify that you can view the assigned policy details on the Applications page.
 - a. On the Applications page, navigate to the **DR Status** column and view the status. The status will either be **healthy** or **critical**. You can click the status to view the last synced time for application volumes and the DR policy assigned.


**NOTE**

It may take a few minutes after disaster recovery is assigned for the **DR Status** to move from **critical** to **healthy**.

- b. Failover and relocate status can be viewed in the **DR Status** column when initiated. These statuses can be clicked to view more details about the action.
DR Status shows either **healthy**, **warning**, or **critical** when failover or relocate is not taking place.
10. After you apply DRPolicy to the applications, confirm whether the **ClusterDataProtected** is set to **True** in the drpc yaml output.

3.12.2. Deleting sample application

This section provides instructions for deleting the sample application **busybox** using the RHACM console.

**IMPORTANT**

When deleting a DR protected application, access to both clusters that belong to the DRPolicy is required. This is to ensure that all protected API resources and resources in the respective S3 stores are cleaned up as part of removing the DR protection. If access to one of the clusters is not healthy, deleting the **DRPlacementControl** resource for the application, on the hub, would remain in the Deleting state.


Prerequisites

- These instructions to delete the sample application should not be executed until the failover and relocate testing is completed and the application is ready to be removed from RHACM and the managed clusters.

Procedure

1. On the RHACM console, navigate to **Applications**.
2. Search for the sample application to be deleted (for example, **busybox**).

3. Click the Action Menu (⋮) next to the application you want to delete.
4. Click **Delete application**.
When the **Delete application** is selected a new screen will appear asking if the application related resources should also be deleted.
5. Select **Remove application related resources** checkbox to delete the Subscription and PlacementRule.
6. Click **Delete**. This will delete the busybox application on the Primary managed cluster (or whatever cluster the application was running on).
7. In addition to the resources deleted using the RHACM console, delete the **DRPlacementControl** if it is not auto-deleted after deleting the **busybox** application.
 - a. Log in to the OpenShift Web console for the Hub cluster and navigate to Installed Operators for the project **busybox-sample**.
For ApplicationSet applications, select the project as **openshift-gitops**.
 - b. Click **OpenShift DR Hub Operator** and then click the **DRPlacementControl** tab.
 - c. Click the Action Menu (⋮) next to the **busybox** application DRPlacementControl that you want to delete.
 - d. Click **Delete DRPlacementControl**.
 - e. Click **Delete**.



NOTE

This process can be used to delete any application with a **DRPlacementControl** resource.

3.13. SUBSCRIPTION-BASED APPLICATION FAILOVER BETWEEN MANAGED CLUSTERS

Perform a failover when a managed cluster becomes unavailable, due to any reason. This failover method is application-based.

Prerequisites

- If your setup has active and passive RHACM hub clusters, see [Hub recovery using Red Hat Advanced Cluster Management](#).
- When the primary cluster is in a state other than **Ready**, check the actual status of the cluster as it might take some time to update.
 1. Navigate to the **RHACM console → Infrastructure → Clusters → Cluster list** tab.
 2. Check the status of both the managed clusters individually before performing failover operation.
However, failover operation can still be performed when the cluster you are failing over to is in a **Ready** state.

Procedure

1. Enable fencing on the **Hub cluster**.

- a. Open CLI terminal and edit the **DRCluster resource**, where `<drcluster_name>` is your unique name.

CAUTION

Once the managed cluster is fenced, **all** communication from applications to the OpenShift Data Foundation external storage cluster will fail and some **Pods** will be in an unhealthy state (for example: **CreateContainerError**, **CrashLoopBackOff**) on the cluster that is now fenced.

```
$ oc edit drcluster <drcluster_name>
```

```
apiVersion: ramendr.openshift.io/v1alpha1
kind: DRCluster
metadata:
[...]
spec:
  ## Add this line
  clusterFence: Fenced
  cidrs:
  [...]
  [...]
```

Example output:

```
drcluster.ramendr.openshift.io/ocp4perf1 edited
```

- b. Verify the fencing status on the **Hub cluster** for the **Primary managed cluster**, replacing `<drcluster_name>` is your unique identifier.

```
$ oc get drcluster.ramendr.openshift.io <drcluster_name> -o jsonpath='{.status.phase}'
{"\n"}
```

Example output:

```
Fenced
```

- c. Login to your Ceph cluster and verify that the IPs that belong to the OpenShift Container Platform cluster nodes are now in the blocklist.

```
$ ceph osd blocklist ls
```

Example output

```
cidr:10.1.161.1:0/32 2028-10-30T22:30:03.585634+0000
cidr:10.1.161.14:0/32 2028-10-30T22:30:02.483561+0000
cidr:10.1.161.51:0/32 2028-10-30T22:30:01.272267+0000
cidr:10.1.161.63:0/32 2028-10-30T22:30:05.099655+0000
cidr:10.1.161.129:0/32 2028-10-30T22:29:58.335390+0000
cidr:10.1.161.130:0/32 2028-10-30T22:29:59.861518+0000
```

2. On the Hub cluster, navigate to **Applications**.
3. Click the **Actions** menu at the end of application row to view the list of available actions.
4. Click **Failover application**.
5. After the **Failover application** modal is shown, select **policy** and **target cluster** to which the associated application will failover in case of a disaster.
6. Click the **Select subscription group** dropdown to verify the default selection or modify this setting.
By default, the subscription group that replicates for the application resources is selected.
7. Check the status of the **Failover readiness**.
 - If the status is **Ready** with a green tick, it indicates that the target cluster is ready for failover to start. Proceed to step 7.
 - If the status is **Unknown** or **Not ready**, then wait until the status changes to **Ready**.
8. Click **Initiate**. The busybox application is now failing over to the **Secondary-managed cluster**.
9. In the **DR status** column, you can view the progress of the failover. The modal shows the progress of the failover as **Preparing** → **Failover** → **Restoring** → **Clean up**.
10. Verify that the activity status shows as **FailedOver** for the application.
 - a. Navigate to the **Applications** → **Overview** tab.
 - b. In the **Data policy** column, click the **policy** link for the application you applied the policy to.
 - c. On the **Data policy** popover, click the **View more details** link.

3.14. APPLICATIONSET-BASED APPLICATION FAILOVER BETWEEN MANAGED CLUSTERS

Perform a failover when a managed cluster becomes unavailable, due to any reason. This failover method is application-based.

Prerequisites

- If your setup has active and passive RHACM hub clusters, see [Hub recovery using Red Hat Advanced Cluster Management](#).
- When the primary cluster is in a state other than **Ready**, check the actual status of the cluster as it might take some time to update.
 1. Navigate to the **RHACM console** → **Infrastructure** → **Clusters** → **Cluster list** tab.
 2. Check the status of both the managed clusters individually before performing failover operation.
However, failover operation can still be performed when the cluster you are failing over to is in a **Ready** state.

Procedure

1. Enable fencing on the **Hub cluster**.

- a. Open CLI terminal and edit the **DRCluster resource**, where `<drcluster_name>` is your unique name.

CAUTION

Once the managed cluster is fenced, **all** communication from applications to the OpenShift Data Foundation external storage cluster will fail and some **Pods** will be in an unhealthy state (for example: **CreateContainerError**, **CrashLoopBackOff**) on the cluster that is now fenced.

```
$ oc edit drcluster <drcluster_name>
```

```
apiVersion: ramendr.openshift.io/v1alpha1
kind: DRCluster
metadata:
[...]
spec:
  ## Add this line
  clusterFence: Fenced
  cidrs:
  [...]
  [...]
```

Example output:

```
drcluster.ramendr.openshift.io/ocp4perf1 edited
```

- b. Verify the fencing status on the **Hub cluster** for the **Primary managed cluster**, replacing `<drcluster_name>` is your unique identifier.

```
$ oc get drcluster.ramendr.openshift.io <drcluster_name> -o jsonpath='{.status.phase}'
{"\n"}
```

Example output:

```
Fenced
```

- c. Login to your Ceph cluster and verify that the IPs that belong to the OpenShift Container Platform cluster nodes are now in the blocklist.

```
$ ceph osd blocklist ls
```

Example output

```
cidr:10.1.161.1:0/32 2028-10-30T22:30:03.585634+0000
cidr:10.1.161.14:0/32 2028-10-30T22:30:02.483561+0000
cidr:10.1.161.51:0/32 2028-10-30T22:30:01.272267+0000
cidr:10.1.161.63:0/32 2028-10-30T22:30:05.099655+0000
cidr:10.1.161.129:0/32 2028-10-30T22:29:58.335390+0000
cidr:10.1.161.130:0/32 2028-10-30T22:29:59.861518+0000
```

2. On the Hub cluster, navigate to **Applications**.
3. Click the **Actions** menu at the end of application row to view the list of available actions.
4. Click **Failover application**.
5. When the **Failover application** modal is shown, verify the details presented are correct and check the status of the Failover **readiness**. If the status is **Ready** with a green tick, it indicates that the target cluster is ready for failover to start.
6. Click **Initiate**. The busybox resources are now created on the target cluster.
7. In the **DR status** column, you can view the progress of the failover. The modal shows the progress of the failover as **Preparing → Failover → Restoring → Clean up**.
8. Verify that the activity status shows as **FailedOver** for the application.
 - a. Navigate to the **Applications → Overview** tab.
 - b. In the **Data policy** column, click the **policy** link for the application you applied the policy to.
 - c. On the **Data policy** popover, verify that you can see one or more policy names and the ongoing activities associated with the policy in use with the application.

3.15. RELOCATING SUBSCRIPTION-BASED APPLICATION BETWEEN MANAGED CLUSTERS

Relocate an application to its preferred location when all managed clusters are available.

Prerequisite

- If your setup has active and passive RHACM hub clusters, see [Hub recovery using Red Hat Advanced Cluster Management](#).
- When the primary cluster is in a state other than **Ready**, check the actual status of the cluster as it might take some time to update. Relocate can only be performed when both primary and preferred clusters are up and running.
 1. Navigate to **RHACM console → Infrastructure → Clusters → Cluster list** tab.
 2. Check the status of both the managed clusters individually before performing relocate operation.
- Verify that applications were cleaned up from the cluster before unfencing it:
 - Ensure that either there are no pods, or that no pods are using PVCs.
 - PVCs remain in terminating state.

Procedure

1. Disable fencing on the Hub cluster.
 - a. Edit the **DRCluster resource** for this cluster, replacing <drcluster_name> with a unique name.


```
$ oc edit drcluster <drcluster_name>
```

```
apiVersion: ramendr.openshift.io/v1alpha1
kind: DRCluster
metadata:
[...]
spec:
  cidrs:
  [...]
  ## Modify this line
  clusterFence: Unfenced
  [...]
  [...]
```

Example output:

```
drcluster.ramendr.openshift.io/ocp4perf1 edited
```

- b. Gracefully reboot OpenShift Container Platform nodes that were **Fenced**. A reboot is required to resume the I/O operations after unfencing to avoid any further recovery orchestration failures. Reboot all nodes of the cluster by following the steps in the procedure, [Rebooting a node gracefully](#).



NOTE

Make sure that all the nodes are initially cordoned and drained before you reboot and perform uncordon operations on the nodes.

- c. After all OpenShift nodes are rebooted and are in a **Ready** status, verify that all Pods are in a healthy state by running this command on the Primary managed cluster (or whatever cluster has been Unfenced).

```
oc get pods -A | egrep -v 'Running|Completed'
```

Example output:

NAMESPACE	READY	STATUS	RESTARTS	NAME	AGE
-----------	-------	--------	----------	------	-----

The output for this query should be zero Pods before proceeding to the next step.



IMPORTANT

If there are Pods still in an unhealthy status because of severed storage communication, troubleshoot and resolve before continuing. Because the storage cluster is external to OpenShift, it also has to be properly recovered after a site outage for OpenShift applications to be healthy.

Alternatively, you can use the OpenShift Web Console dashboards and Overview tab to assess the health of applications and the external ODF storage cluster. The detailed OpenShift Data Foundation dashboard is found by navigating to **Storage → Data Foundation**.

- d. Verify that the **Unfenced** cluster is in a healthy state. Validate the fencing status in the Hub cluster for the Primary-managed cluster, replacing <drcluster_name> with a unique name.

```
$ oc get drcluster.ramendr.openshift.io <drcluster_name> -o jsonpath='{.status.phase}{"\n"}'
```

Example output:

```
Unfenced
```

- e. Login to your Ceph cluster and verify that the IPs that belong to the OpenShift Container Platform cluster nodes are NOT in the blocklist.

```
$ ceph osd blocklist ls
```

Ensure that you do not see the IPs added during fencing.

2. On the Hub cluster, navigate to **Applications**.
3. Click the **Actions** menu at the end of application row to view the list of available actions.
4. Click **Relocate application**.
5. When the **Relocate application** modal is shown, select **policy** and **target cluster** to which the associated application will relocate to in case of a disaster.
6. By default, the subscription group that will deploy the application resources is selected. Click the **Select subscription group** dropdown to verify the default selection or modify this setting.
7. Check the status of the **Relocation readiness**.
 - If the status is **Ready** with a green tick, it indicates that the target cluster is ready for relocation to start. Proceed to step 7.
 - If the status is **Unknown** or **Not ready**, then wait until the status changes to **Ready**.
8. Click **Initiate**. The busybox resources are now created on the target cluster.
9. In the **DR status** column, you can view the progress of the failover. The modal shows the progress of the failover as **Preparing** → **Syncing** → **Restoring** → **Clean up**.
10. Verify that the activity status shows as **Relocated** for the application.
 - a. Navigate to the **Applications** → **Overview** tab.
 - b. In the **Data policy** column, click the **policy** link for the application you applied the policy to.
 - c. On the **Data policy** popover, click the **View more details** link.

3.16. RELOCATING AN APPLICATIONSET-BASED APPLICATION BETWEEN MANAGED CLUSTERS

Relocate an application to its preferred location when all managed clusters are available.

Prerequisite

- If your setup has active and passive RHACM hub clusters, see [Hub recovery using Red Hat Advanced Cluster Management](#).
- When the primary cluster is in a state other than **Ready**, check the actual status of the cluster as it might take some time to update. Relocate can only be performed when both primary and preferred clusters are up and running.
 1. Navigate to **RHACM console → Infrastructure → Clusters → Cluster list** tab.
 2. Check the status of both the managed clusters individually before performing relocate operation.
- Verify that applications were cleaned up from the cluster before unfencing it:
 - Ensure that either there are no pods, or that no pods are using PVCs.
 - PVCs remain in terminating state.

Procedure

1. Disable fencing on the Hub cluster.
 - a. Edit the **DRCluster** resource for this cluster, replacing <drcluster_name> with a unique name.

```
$ oc edit drcluster <drcluster_name>

apiVersion: ramendr.openshift.io/v1alpha1
kind: DRCluster
metadata:
[...]
spec:
  cidrs:
  [...]
  ## Modify this line
  clusterFence: Unfenced
  [...]
  [...]
```

Example output:

```
drcluster.ramendr.openshift.io/ocp4perf1 edited
```

- b. Gracefully reboot OpenShift Container Platform nodes that were **Fenced**. A reboot is required to resume the I/O operations after unfencing to avoid any further recovery orchestration failures. Reboot all nodes of the cluster by following the steps in the procedure, [Rebooting a node gracefully](#).



NOTE

Make sure that all the nodes are initially cordoned and drained before you reboot and perform uncordon operations on the nodes.

- c. After all OpenShift nodes are rebooted and are in a **Ready** status, verify that all Pods are in a healthy state by running this command on the Primary managed cluster (or whatever cluster has been Unfenced).

```
oc get pods -A | egrep -v 'Running|Completed'
```

Example output:

```

NAMESPACE          NAME
READY STATUS    RESTARTS    AGE
```

The output for this query should be zero Pods before proceeding to the next step.



IMPORTANT

If there are Pods still in an unhealthy status because of severed storage communication, troubleshoot and resolve before continuing. Because the storage cluster is external to OpenShift, it also has to be properly recovered after a site outage for OpenShift applications to be healthy.

Alternatively, you can use the OpenShift Web Console dashboards and Overview tab to assess the health of applications and the external ODF storage cluster. The detailed OpenShift Data Foundation dashboard is found by navigating to **Storage → Data Foundation**.

- d. Verify that the **Unfenced** cluster is in a healthy state. Validate the fencing status in the Hub cluster for the Primary-managed cluster, replacing <drcluster_name> with a unique name.

```
$ oc get drcluster.ramendr.openshift.io <drcluster_name> -o jsonpath='{.status.phase}'
```

Example output:

```
Unfenced
```

- e. Login to your Ceph cluster and verify that the IPs that belong to the OpenShift Container Platform cluster nodes are NOT in the blocklist.

```
$ ceph osd blocklist ls
```

Ensure that you do not see the IPs added during fencing.

2. On the Hub cluster, navigate to **Applications**.
3. Click the **Actions** menu at the end of application row to view the list of available actions.
4. Click **Relocate application**.
5. When the **Relocate application** modal is shown, select **policy** and **target cluster** to which the associated application will relocate to in case of a disaster.
6. Click **Initiate**. The busybox resources are now created on the target cluster.

7. Close the modal window and track the status using the **DR Status** column on the Applications page. The status moves from **Preparing** → **Syncing** → **Restoring** → **Clean up**.
8. Verify that the activity status shows as **Relocated** for the application.
 - a. Navigate to the **Applications** → **Overview** tab.
 - b. In the **DR Status** column, click the **policy** link for the application you applied the policy to.
 - c. On the **Data policy** popover, verify that you can see one or more policy names and the relocation status associated with the policy in use with the application.



IMPORTANT

Disaster recovery protection for discovered applications is a Technology Preview feature and is subject to Technology Preview support limitations. Technology Preview features are not supported with Red Hat production service level agreements (SLAs) and might not be functionally complete. Red Hat does not recommend using them in production. These features provide early access to upcoming product features, enabling customers to test functionality and provide feedback during the development process.

For more information, see [Technology Preview Features Support Scope](#).

3.17. HUB RECOVERY USING RED HAT ADVANCED CLUSTER MANAGEMENT [TECHNOLOGY PREVIEW]

When your setup has active and passive Red Hat Advanced Cluster Management for Kubernetes (RHACM) hub clusters, and in case where the active hub is down, you can use the passive hub to failover or relocate the disaster recovery protected workloads.



IMPORTANT

Hub recovery for Metro-DR is a Technology Preview feature and is subject to Technology Preview support limitations. Technology Preview features are not supported with Red Hat production service level agreements (SLAs) and might not be functionally complete. Red Hat does not recommend using them in production. These features provide early access to upcoming product features, enabling customers to test functionality and provide feedback during the development process.

For more information, see [Technology Preview Features Support Scope](#).

3.17.1. Configuring passive hub cluster

To perform hub recovery in case the active hub is down or unreachable, follow the procedure in this section to configure the passive hub cluster and then failover or relocate the disaster recovery protected workloads.

Procedure

1. Ensure that RHACM operator and **MultiClusterHub** is installed on the passive hub cluster. See [RHACM installation guide](#) for instructions.
After the operator is successfully installed, the web console automatically reloads to apply the changes. During this process, a temporary error message might appear on the page and this is expected and disappears after the refresh completes.

2. Before hub recovery, configure backup and restore. See [Backup and restore](#) topics of *RHACM Business continuity* guide.
3. Install the multicluster orchestrator (MCO) operator along with Red Hat OpenShift GitOps operator on the passive RHACM hub prior to the restore. For instructions to restore your RHACM hub, see [Installing OpenShift Data Foundation Multicluster Orchestrator operator](#).
4. Ensure that **.spec.cleanupBeforeRestore** is set to **None** for the **Restore.cluster.open-cluster-management.io** resource. For details, see [Restoring passive resources while checking for backups](#) chapter of RHACM documentation.
5. If SSL access across clusters was configured manually during setup, then re-configure SSL access across clusters. For instructions, see [Configuring SSL access across clusters](#) chapter.
6. On the passive hub, add label to **openshift-operators** namespace to enable basic monitoring of **VolumeSynchronizationDelay** alert using this command. For alert details, see [Disaster recovery alerts](#) chapter.

```
$ oc label namespace openshift-operators openshift.io/cluster-monitoring='true'
```

3.17.2. Switching to passive hub cluster

Use this procedure when active hub is down or unreachable.


Procedure

1. During the restore procedure, to avoid eviction of resources when **ManifestWorks** are not regenerated correctly, you can enlarge the **AppliedManifestWork** eviction grace period. On the passive hub cluster, check for existing global **KlusterletConfig**.
 - If global **KlusterletConfig** exists then edit and set the value for **appliedManifestWorkEvictionGracePeriod** parameter to a larger value. For example, 24 hours or more.
 - If global **KlusterletConfig** does not exist, then create the **Klusterletconfig** using the following yaml:

```
apiVersion: config.open-cluster-management.io/v1alpha1
kind: KlusterletConfig
metadata:
  name: global
spec:
  appliedManifestWorkEvictionGracePeriod: "24h"
```

The configuration will be propagated to all the managed clusters automatically.

2. Restore the backups on the passive hub cluster. For information, see [Restoring a hub cluster from backup](#).



IMPORTANT

Recovering a failed hub to its passive instance will only restore applications and their DR protected state to its last scheduled backup. Any application that was DR protected after the last scheduled backup would need to be protected again on the new hub.

3. Verify that the restore is complete.

```
$ oc -n <restore-namespace> wait restore <restore-name> --
for=jsonpath='{.status.phase}'=Finished --timeout=120s
```

4. Verify that the Primary and Secondary managed clusters are successfully imported into the RHACM console and they are accessible. If any of the managed clusters are down or unreachable then they will not be successfully imported.
5. Wait until DRPolicy validation succeeds.
6. Verify that the **DRPolicy** is created successfully. Run this command on the **Hub cluster** for each of the DRPolicy resources created, where *<drpolicy_name>* is replaced with a unique name.

```
$ oc get drpolicy <drpolicy_name> -o jsonpath='{.status.conditions[].reason}'{"\n"}
```

Example output:

```
Succeeded
```

7. Refresh the RHACM console to make the DR monitoring dashboard tab accessible if it was enabled on the Active hub cluster.
8. Verify the DRPC output using the following command on the new hub cluster:

```
$ oc get drpc -A -o wide
```

If **PROGRESSION** shows a status of **PAUSED**, administrative intervention is required to unpause it. **PROGRESSION** enters **PAUSED** state under the following conditions:

- **Cluster Query Failure:** None of the clusters were successfully queried during the DRPC reconciliation. This situation can occur during hub recovery.
- **Action Mismatch:** The DRPC action differs from the queried VRG action.
- **Cluster Mismatch:** The DRPC action and the VRG action are the same, but the Primary VRG is found in a different cluster than the one expected by the DRPC.



IMPORTANT

If you cannot diagnose and resolve the cause of the pause, contact [Red Hat Customer Support](#).

If **PROGRESSION** is in either **Completed** or **Cleaning up**, it is safe to proceed.

9. Edit the global **KlusterletConfig** on the new hub and remove the parameter **appliedManifestWorkEvictionGracePeriod** and its value.
10. Depending on whether the active hub cluster, or both the active hub cluster along with the primary managed cluster had been down, follow the next steps based on your scenario:
 - a. If only the active hub cluster had been down, and if the managed clusters are still accessible, no further action is required.

- b. If the primary managed cluster had been down, along with the active hub cluster, you need to fail over the workloads from the primary managed cluster to the secondary managed cluster.

For failover instructions, based on your workload type, see [Subscription-based applications](#) or [ApplicationSet-based applications](#).

11. Verify that the failover is successful. If the Primary managed cluster is also down, then the PROGRESSION status for the workload would be in **Cleaning Up** phase until the down Primary managed cluster is back online and successfully imported into the RHACM console.

On the passive hub cluster, run the following command to check the PROGRESSION status.

```
$ oc get drpc -o wide -A
```


CHAPTER 4. REGIONAL-DR SOLUTION FOR OPENSIFT DATA FOUNDATION

This section of the guide provides details of the Regional Disaster Recovery (Regional-DR) steps and commands necessary to be able to failover an application from one OpenShift Container Platform cluster to another and then failback the same application to the original primary cluster. In this case the OpenShift Container Platform clusters will be created or imported using Red Hat Advanced Cluster Management (RHACM).

Though Regional-DR solution is built on Asynchronous data replication and hence could have a potential data loss but provides the protection against a broad set of failures.

This is a general overview of the Regional-DR steps required to configure and execute OpenShift Disaster Recovery (ODR) capabilities using OpenShift Data Foundation and RHACM across two distinct OpenShift Container Platform clusters separated by distance. In addition to these two clusters called managed clusters, a third OpenShift Container Platform cluster is required that will be the Red Hat Advanced Cluster Management (RHACM) hub cluster.



NOTE

Regional-DR is currently supported only on VMware, bare metal, and hybrid cloud (one cluster in the cloud) environments.



IMPORTANT

You can now easily set up Regional disaster recovery solutions for workloads based on OpenShift virtualization technology using OpenShift Data Foundation. For more information, see the [knowledgebase article](#).

4.1. COMPONENTS OF REGIONAL-DR SOLUTION

Regional-DR is composed of Red Hat Advanced Cluster Management for Kubernetes and OpenShift Data Foundation components to provide application and data mobility across Red Hat OpenShift Container Platform clusters.

Red Hat Advanced Cluster Management for Kubernetes

Red Hat Advanced Cluster Management (RHACM) provides the ability to manage multiple clusters and application lifecycles. Hence, it serves as a control plane in a multi-cluster environment.

RHACM is split into two parts:

- RHACM Hub: components that run on the multi-cluster control plane.
- Managed clusters: components that run on the clusters that are managed.

For more information about this product, see [RHACM documentation](#) and the [RHACM “Manage Applications” documentation](#).

OpenShift Data Foundation

OpenShift Data Foundation provides the ability to provision and manage storage for stateful applications in an OpenShift Container Platform cluster.

OpenShift Data Foundation is backed by Ceph as the storage provider, whose lifecycle is managed by Rook in the OpenShift Data Foundation component stack. Ceph-CSI provides the provisioning and management of Persistent Volumes for stateful applications.

OpenShift Data Foundation stack is now enhanced with the following abilities for disaster recovery:

- Enable RBD block pools for mirroring across OpenShift Data Foundation instances (clusters)
- Ability to mirror specific images within an RBD block pool
- Provides csi-addons to manage per Persistent Volume Claim (PVC) mirroring

OpenShift DR

OpenShift DR is a set of orchestrators to configure and manage stateful applications across a set of peer OpenShift clusters which are managed using RHACM and provides cloud-native interfaces to orchestrate the life-cycle of an application's state on Persistent Volumes. These include:

- Protecting an application and its state relationship across OpenShift clusters
- Failing over an application and its state to a peer cluster
- Relocate an application and its state to the previously deployed cluster

OpenShift DR is split into three components:

- **ODF Multicluster Orchestrator:** Installed on the multi-cluster control plane (RHACM Hub), it orchestrates configuration and peering of OpenShift Data Foundation clusters for Metro and Regional DR relationships
- **OpenShift DR Hub Operator:** Automatically installed as part of ODF Multicluster Orchestrator installation on the hub cluster to orchestrate failover or relocation of DR enabled applications.
- **OpenShift DR Cluster Operator:** Automatically installed on each managed cluster that is part of a Metro and Regional DR relationship to manage the lifecycle of all PVCs of an application.

4.2. HYBRID CLOUD REGIONAL-DR SOLUTION

OpenShift Data Foundation supports pairing an on-premise OpenShift Data Foundation deployment with a secondary cluster in the cloud, enabling hybrid disaster recovery and workload flexibility. This capability allows on-prem environments such as Bare Metal and vSphere to integrate with cloud platforms Azure and Azure Red Hat OpenShift (ARO), expanding OpenShift Data Foundation's support for hybrid deployment scenarios.

4.3. REGIONAL-DR DEPLOYMENT WORKFLOW

This section provides an overview of the steps required to configure and deploy Regional-DR capabilities using the latest version of Red Hat OpenShift Data Foundation across two distinct OpenShift Container Platform clusters. In addition to two managed clusters, a third OpenShift Container Platform cluster will be required to deploy the Red Hat Advanced Cluster Management (RHACM).

To configure your infrastructure, perform the below steps in the order given:

1. Ensure requirements across the three: Hub, Primary and Secondary OpenShift Container Platform clusters that are part of the DR solution are met. See [Requirements for enabling Regional-DR](#).

2. Install OpenShift Data Foundation operator and create a storage system on Primary and Secondary managed clusters. See [Creating OpenShift Data Foundation cluster on managed clusters](#).
3. Install the ODF Multiclustor Orchestrator on the Hub cluster. See [Installing ODF Multiclustor Orchestrator on Hub cluster](#).
4. Configure SSL access between the Hub, Primary and Secondary clusters. See [Configuring SSL access across clusters](#).
5. Create a DRPolicy resource for use with applications requiring DR protection across the Primary and Secondary clusters. See [Creating Disaster Recovery Policy on Hub cluster](#).


**NOTE**

There can be more than a single policy.

6. Testing your disaster recovery solution with:
 - a. **ApplicationSet-based** application:
 - Create sample applications. See [Creating ApplicationSet-based applications](#).
 - Test failover and relocate operations using the sample application between managed clusters. See [ApplicationSet-based application failover](#) and [relocating ApplicationSet-based application](#).
 - b. **Discovered applications**
 - Ensure all requirements mentioned in Prerequisites is addressed. See [Prerequisites for disaster recovery protection of discovered applications](#)
 - Create a sample discovered application. See [Creating a sample discovered application](#)
 - Enroll the discovered application. See [Enrolling a sample discovered application for disaster recovery protection](#)
 - Test failover and relocate. See [Discovered application failover and relocate](#)

4.4. REQUIREMENTS FOR ENABLING REGIONAL-DR

The prerequisites to installing a disaster recovery solution supported by Red Hat OpenShift Data Foundation are as follows:


**IMPORTANT**

If you plan to protect Virtual Machines (VMs) that were migrated using the Migration Toolkit for Virtualization (MTV), you must ensure that all residual migration resources are completely cleaned up. Failure to do so will cause disaster recovery workflows to stall or fail.

Before proceeding with DR configuration, review the mandatory cleanup prerequisites in [Prerequisite for MTV-Migrated Virtual Machines](#).

- You must have three OpenShift clusters that have network reachability between them:

- **Hub cluster** where Red Hat Advanced Cluster Management (RHACM) for Kubernetes operator is installed.
- **Primary managed cluster** where OpenShift Data Foundation is running.
- **Secondary managed cluster** where OpenShift Data Foundation is running.




NOTE

For configuring hub recovery setup, you need a 4th cluster which acts as the passive hub. The primary managed cluster (Site-1) can be co-situated with the active RHACM hub cluster while the passive hub cluster is situated along with the secondary managed cluster (Site-2). Alternatively, the active RHACM hub cluster can be placed in a neutral site (Site-3) that is not impacted by the failures of either of the primary managed cluster at Site-1 or the secondary cluster at Site-2. In this situation, if a passive hub cluster is used it can be placed with the secondary cluster at Site-2. For more information, see [Configuring passive hub cluster for hub recovery](#).

- Ensure that RHACM operator and MultiClusterHub is installed on the Hub cluster. See [RHACM installation guide](#) for instructions and [Submariner networking requirements table](#) for network prerequisites.

After the operator is successfully installed, the web console automatically reloads to apply the changes. During this process, a temporary error message might appear on the page and this is expected and disappears after the refresh completes.



IMPORTANT

Ensure that application traffic routing and redirection are configured appropriately.

- On the Hub cluster
 - Navigate to **All Clusters → Infrastructure → Clusters**.
 - Import or create the **Primary managed cluster** and the **Secondary managed cluster** using the RHACM console.
 - Choose the appropriate options for your environment.
For instructions, see [Creating clusters](#) and [Cluster import](#).
- Connect the private OpenShift cluster and service networks using the RHACM Submariner add-ons. Verify that the two clusters have non-overlapping service and cluster private networks. Otherwise, ensure that the Globalnet is enabled during the Submariner add-ons installation. Run the following command for each of the managed clusters to determine if Globalnet needs to be enabled. The following example is for non-overlapping cluster and service networks so Globalnet would not be enabled.

```
$ oc get networks.config.openshift.io cluster -o json | jq .spec
```

Example output for Primary cluster:

```
{
  "clusterNetwork": [
    {
```

```

        "cidr": "10.5.0.0/16",
        "hostPrefix": 23
    },
    ],
    "externalIP": {
        "policy": {}
    },
    "networkType": "OVNKubernetes",
    "serviceNetwork": [
        "10.15.0.0/16"
    ]
}

```

Example output for Secondary cluster:

```

{
  "clusterNetwork": [
    {
      "cidr": "10.6.0.0/16",
      "hostPrefix": 23
    }
  ],
  "externalIP": {
    "policy": {}
  },
  "networkType": "OVNKubernetes",
  "serviceNetwork": [
    "10.16.0.0/16"
  ]
}

```

Additionally, if Submariner and [OpenShift Data Foundation](#) are already installed on the managed clusters, use the OpenShift Data Foundation command line interface (CLI) tool to get additional information about the clusters. This information can determine the need for enabling Globalnet during submariner installation based on the clusters service and private networks.

Download the the OpenShift Data Foundation CLI tool from the [customer portal](#).

Run the following command on one of the two managed clusters, where **PeerManagedClusterName(ClusterID)** is the name of the peer OpenShift Data Foundation cluster:


```
$ odf get dr-prereq <PeerManagedClusterName(ClusterID)>
```

If submariner is not installed with Globalnet on clusters with overlapping services, the following output shows:

```

Info: Submariner is installed.
Info: Globalnet is required.
Info: Globalnet is not enabled.

```

**NOTE**

This requires Submariner to be uninstalled and then reinstalled with Globalnet enabled.

If submariner is installed with Globalnet on clusters with overlapping services, the following output shows:

```
Info: Submariner is installed.
Info: Globalnet is required.
Info: Globalnet is enabled.
```

If submariner is installed without Globalnet on clusters with non-overlapping services, the following output shows:

```
Info: Submariner is installed.
Info: Globalnet is not required.
Info: Globalnet is not enabled.
```

If submariner is installed with Globalnet on clusters with non-overlapping services, the following output shows:

```
Info: Submariner is installed.
Info: Globalnet is not required.
Info: Globalnet is enabled.
```

For more information, see [Submariner documentation](#).

4.4.1. Prerequisite for MTV-Migrated Virtual Machines

When Virtual Machines (VMs) are migrated using the Migration Toolkit for Virtualization (MTV), additional migration-related resources are created. If these resources are not properly cleaned up, they can block Disaster Recovery (DR) workflows.

During and after an MTV migration, the following pods may retain mounts on the VM's PersistentVolumeClaim (PVC):

- **virt-launcher pod:** This is expected and required for any running VM.
- **MTV migration / validation pods:** Created as part of the MTV migration workflow (for example, disk validation or import pods).

After migration completes, MTV migration or validation pods typically move to a **Completed** state. However, a **Completed** state does not guarantee that PVC mounts are released. If the associated MTV migration plan is not archived, these pods may continue to retain PVC mounts. This prevents PVC cleanup and causes DR workflows to fail or stall when failover or relocate operations are initiated.

Impact on DR / RDR workflows

If MTV migration artifacts are not fully cleaned up, you may observe the following issues:

- PVCs or **VolumeReplication** resources stuck in **Terminating**
- DR relocate or failover stuck in **WaitingForCleanup**

- DR resources not created on the target cluster
- Applications remaining in a **Relocating** or **FailingOver** state

Required Actions Before Enabling DR

Before enrolling an MTV-migrated VM for DR protection or initiating a relocate/failover operation, ensure that:

1. All MTV migration plans associated with the VM are archived. Deleting a migration plan without archiving it is not sufficient and can leave residual resources behind. Archiving the migration plan is the supported and required method to finalize the migration and release related resources.
2. No MTV migration or validation pods are holding PVC mounts.
3. VM PVCs are mounted only by the expected **virt-launcher** pod.

For detailed instructions on how to archive a migration plan, see [Migrating virtual machines by using the Red Hat OpenShift web console](#) in the *Installing and using the Migration Toolkit for Virtualization* guide.

IMPORTANT

When you are protecting virtual machines (VMs) with DR, you must ensure that the target VMs have successfully transitioned to the **Running** state before triggering any DR protection workflows, failover, or relocate operations.

Do not enroll or initiate DR operations for a VM that is in any of the following intermediate or pending phases:

- **Provisioning**
- **Scheduling**
- **CreatingDataVolume**
- **Starting**

During these initial stages, PersistentVolumeClaim (PVC) creation, storage attachment, and disk initialization are still actively underway. Triggering DR actions too early can result in the system unintentionally protecting temporary or intermediate volumes, which leads to replication inconsistencies, stuck states, or complete recovery failures on the target cluster.

4.5. CREATING AN OPENSIFT DATA FOUNDATION CLUSTER ON MANAGED CLUSTERS

In order to configure storage replication between the two OpenShift Container Platform clusters, create an OpenShift Data Foundation storage system after you install the OpenShift Data Foundation operator.

NOTE

Refer to OpenShift Data Foundation deployment guides and instructions that are specific to your infrastructure.

Procedure

1. Install and configure the latest **OpenShift Data Foundation** cluster on each of the managed clusters.
For information about the OpenShift Data Foundation deployment, refer to your [infrastructure specific deployment guides](#).

2. Validate the successful deployment of OpenShift Data Foundation on each managed cluster with the following command:

```
$ oc get storagecluster -n openshift-storage ocs-storagecluster -o jsonpath='{.status.phase}'
```

For the Multicloud Gateway (MCG):

```
$ oc get noobaa -n openshift-storage noobaa -o jsonpath='{.status.phase}'
```

If the status result is **Ready** for both queries on the **Primary managed cluster** and the **Secondary managed cluster**, then continue with the next step.

3. Verify the status of the storage cluster:
 - a. On the OpenShift Web Console, navigate to **Storage → Data Foundation → Storage System**.
 - b. In the **Status** card of the **Overview** tab, click **Storage System** and then click the storage system link from the pop up that appears.
 - c. In the **Status** card of the **Block and File** tab, verify that the *Storage Cluster* has a green tick.
4. [Optional] If Globalnet was enabled when Submariner was installed, then edit the **StorageCluster** after the OpenShift Data Foundation install finishes.
For Globalnet networks, manually edit the **StorageCluster** yaml to add the **clusterID** and set **enabled** to **true**. Replace *<clustername>* with your RHACM imported or newly created managed cluster name. Edit the **StorageCluster** on both the Primary managed cluster and the Secondary managed cluster.




WARNING

Do not make this change in the **StorageCluster** unless you enabled Globalnet when Submariner was installed.

```
$ oc edit storagecluster -o yaml -n openshift-storage
```

```
spec:
  network:
    multiClusterService:
      clusterID: <clustername>
      enabled: true
```


**IMPORTANT**

If **multiClusterService** is enabled, it can not be disabled or undone as it failover the MONs and restarts the OSDs with GlobalNet IP addresses which can not be changed once assigned.

5. After the above changes are made,
 - a. Wait for the OSD pods to restart and OSD services to be created.
 - b. Wait for all MONS to failover.
 - c. Ensure that the MONS and OSD services are exported.

```
$ oc get serviceexport -n openshift-storage
```

NAME	AGE
rook-ceph-mon-d	4d14h
rook-ceph-mon-e	4d14h
rook-ceph-mon-f	4d14h
rook-ceph-osd-0	4d14h
rook-ceph-osd-1	4d14h
rook-ceph-osd-2	4d14h

- d. Ensure that cluster is in a **Ready** state and cluster health has a green tick indicating **Health ok**. Verify using step 3.

4.6. INSTALLING OPENSIFT DATA FOUNDATION MULTICLUSTER ORCHESTRATOR OPERATOR

OpenShift Data Foundation Multicluster Orchestrator is a controller that is installed from OpenShift Container Platform's OperatorHub on the Hub cluster.

Prerequisites

- The **OpenShift GitOps** operator must be installed. For instructions, see [Installing Red Hat OpenShift GitOps Operator in web console](#).

Procedure

1. On the **Hub cluster**, navigate to **Ecosystem → Software Catalog** and use the keyword filter to search for **ODF Multicluster Orchestrator**.
2. Click **ODF Multicluster Orchestrator** tile.
3. Keep all default settings and click **Install**.
Ensure that the operator resources are installed in **openshift-operators** project and available to all namespaces.

**NOTE**

The **ODF Multicluster Orchestrator** also installs the **Openshift DR Hub Operator** on the RHACM hub cluster as a dependency.

- Verify that the operator **Pods** are in a **Running** state. The **OpenShift DR Hub operator** is also installed at the same time in **openshift-operators** namespace.

```
$ oc get pods -n openshift-operators
```

Example output:

NAME	READY	STATUS	RESTARTS	AGE
odf-multicluster-console-6845b795b9-blxrn	1/1	Running	0	4d20h
odfmo-controller-manager-f9d9dfb59-jbrsd	1/1	Running	0	4d20h
ramen-hub-operator-6fb887f885-fss4w	2/2	Running	0	4d20h

4.6.1. [Optional] Connecting storage clusters with the ocs-provider-server ServiceExport

[Optional] If Globalnet was enabled when Submariner was installed, use this section to connect the **StorageClusters** with the **ocs-provider-server** ServiceExport.

Procedure

- Create a **ServiceExport** called **ocs-provider-server**. Copy the following yaml to **ocs-provider-server.yaml** and create the resource on both the **Primary managed cluster** and the **Secondary managed cluster**.

```
apiVersion: multicluster.x-k8s.io/v1alpha1
kind: ServiceExport
metadata:
  name: ocs-provider-server
  namespace: openshift-storage
```

- Add an annotation to the **StorageCluster**. To add the annotation, run the following command on both the **Primary managed cluster** and the **Secondary managed cluster**.



NOTE

The **<managedcluster_name>** refers to the **ManagedCluster** name in RHACM for the **Primary managed cluster** or the **Secondary managed cluster** name.

```
$ oc annotate storagecluster ocs-storagecluster -n openshift-storage ocs.openshift.io/api-server-exported-address=<managedcluster_name>.ocs-provider-server.openshift-storage.svc.clusterset.local:50051
```

4.7. CONFIGURING SSL ACCESS ACROSS CLUSTERS

Configure network (SSL) access between the **primary and secondary clusters** so that metadata can be stored on the alternate cluster in a Multicloud Gateway (MCG) **object bucket** using a secure transport protocol and in the **Hub cluster** for verifying access to the object buckets.



NOTE

If all of your OpenShift clusters are deployed using a signed and valid set of certificates for your environment then this section can be skipped.

Procedure

1. Extract the ingress certificate for the Primary managed cluster and save the output to **primary.crt**.

```
$ oc get cm default-ingress-cert -n openshift-config-managed -o jsonpath="{['data']['ca-bundle\.crt']}" > primary.crt
```

2. Extract the ingress certificate for the Secondary managed cluster and save the output to **secondary.crt**.

```
$ oc get cm default-ingress-cert -n openshift-config-managed -o jsonpath="{['data']['ca-bundle\.crt']}" > secondary.crt
```

3. Create a new **ConfigMap** file to hold the remote cluster's certificate bundle with filename **cm-clusters.crt.yaml**.



NOTE

There could be more or less than three certificates for each cluster as shown in this example file. Also, ensure that the certificate contents are correctly indented after you copy and paste from the **primary.crt** and **secondary.crt** files that were created before.

```
apiVersion: v1
data:
  ca-bundle.crt: |
    -----BEGIN CERTIFICATE-----
    <copy contents of cert1 from primary.crt here>
    -----END CERTIFICATE-----

    -----BEGIN CERTIFICATE-----
    <copy contents of cert2 from primary.crt here>
    -----END CERTIFICATE-----

    -----BEGIN CERTIFICATE-----
    <copy contents of cert3 primary.crt here>
    -----END CERTIFICATE-----

    -----BEGIN CERTIFICATE-----
    <copy contents of cert1 from secondary.crt here>
    -----END CERTIFICATE-----

    -----BEGIN CERTIFICATE-----
    <copy contents of cert2 from secondary.crt here>
    -----END CERTIFICATE-----

    -----BEGIN CERTIFICATE-----
    <copy contents of cert3 from secondary.crt here>
```

```
-----END CERTIFICATE-----
kind: ConfigMap
metadata:
  name: user-ca-bundle
  namespace: openshift-config
```

4. Create the **ConfigMap** on the **Primary managed cluster**, **Secondary managed cluster**, and the **Hub cluster**.

```
$ oc create -f cm-clusters-crt.yaml
```

Example output:

```
configmap/user-ca-bundle created
```

5. Patch default proxy resource on the **Primary managed cluster**, **Secondary managed cluster**, and the **Hub cluster**.

```
$ oc patch proxy cluster --type=merge --patch='{"spec":{"trustedCA":{"name":"user-ca-bundle"}}}'
```

Example output:

```
proxy.config.openshift.io/cluster patched
```

4.8. CREATING DISASTER RECOVERY POLICY ON HUB CLUSTER

OpenShift Disaster Recovery Policy (DRPolicy) resource specifies OpenShift Container Platform clusters participating in the disaster recovery solution and the desired replication interval. DRPolicy is a cluster scoped resource that users can apply to applications that require Disaster Recovery solution.

The **ODF MultiCluster Orchestrator** Operator facilitates the creation of each **DRPolicy** and the corresponding **DRClusters** through the **Multicluster Web console**.

Prerequisites

- Ensure that there is a minimum set of two managed clusters.

Procedure

1. On the **OpenShift console**, navigate to **All Clusters → Data Services → Disaster recovery**.
2. On the Overview tab, click **Create a disaster recovery policy** or you can navigate to Policies tab and click **Create DRPolicy**.
3. Enter **Policy name**. Ensure that each DRPolicy has a unique name (for example: **ocp4bos1-ocp4bos2-5m**).
4. Select two clusters from the list of managed clusters to which this new policy will be associated with.
5. **Replication policy** is automatically set to **Asynchronous** based on the OpenShift clusters selected and a **Sync schedule** option will become available.

6. Set **Sync schedule**.**IMPORTANT**

For every desired replication interval a new **DRPolicy** must be created with a unique name (such as: **ocp4bos1-ocp4bos2-10m**). The same clusters can be selected but the **Sync schedule** can be configured with a different replication interval in minutes/hours/days. The minimum is one minute.

7. **Optional:** Expand **Advanced settings** and check the **Enable disaster recovery support for restored and cloned PersistentVolumeClaims (For Data Foundation only)** checkbox. See [Disaster recovery protection for cloned and restored RBD volumes](#) for more details.

**NOTE**

This option should only be used with discovered applications.

8. Click **Create**.


9. Verify that the **DRPolicy** is created successfully. Run this command on the **Hub cluster** for each of the DRPolicy resources created, where `<drpolicy_name>` is replaced with your unique name.

```
$ oc get drpolicy <drpolicy_name> -o jsonpath='{.status.conditions[].reason}'{"\n"}
```

Example output:

```
Succeeded
```

When a DRPolicy is created, along with it, two DRCluster resources are also created. It could take up to 10 minutes for all three resources to be validated and for the status to show as **Succeeded**.

**NOTE**

Editing of **SchedulingInterval**, **ReplicationClassSelector**, **VolumeSnapshotClassSelector** and **DRClusters** field values are not supported in the DRPolicy.

10. Verify the object bucket access from the **Hub cluster** to both the **Primary managed cluster** and the **Secondary managed cluster**.

- a. Get the names of the **DRClusters** on the Hub cluster.

```
$ oc get drclusters
```

Example output:

```
NAME      AGE
ocp4bos1  4m42s
ocp4bos2  4m42s
```

- b. Check S3 access to each bucket created on each managed cluster. Use the DRCluster validation command, where `<drcluster_name>` is replaced with your unique name.

**NOTE**

Editing of **Region** and **S3ProfileName** field values are non supported in DRClusters.

```
$ oc get drcluster <drcluster_name> -o jsonpath='{.status.conditions[2].reason}'
```

Example output:

```
Succeeded
```

**NOTE**

Make sure to run commands for both **DRClusters** on the **Hub cluster**.

11. Verify that the **OpenShift DR Cluster operator** installation was successful on the **Primary managed cluster** and the **Secondary managed cluster**.

```
$ oc get csv,pod -n openshift-dr-system | egrep 'odr|ramen'
```

Example output:

```
clusterserviceversion.operators.coreos.com/odr-cluster-operator.v4.21.0-rhodef Openshift
DR Cluster Operator 4.21.0-rhodef odr-cluster-operator.v4.21.0-rhodef Succeeded
pod/ramen-dr-cluster-operator-77bf74849c-rfn85 2/2 Running 0 7h20m
```

You can also verify that **OpenShift DR Cluster Operator** is installed successfully in the **Ecosystem → Software Catalog** of each managed cluster.

12. Verify that the **StorageClusterPeer** is in a **Peered** state on the **Primary managed cluster** and the **Secondary managed cluster**.

**NOTE**

`<managedcluster_name>` refers to the **ManagedCluster** name in RHACM. On the **Primary managed cluster** use the **Secondary managed cluster** name. On the **Secondary managed cluster** use the **Primary managed cluster** name.

```
$ oc get storageclusterpeer <managedcluster_name>-peer -n openshift-storage -oyaml | yq
'.status.state'
```

Example output:

```
Peered
```

**NOTE**

When the initial DRPolicy is created, the VolSync operator is installed automatically in the **volsync-system** project on each managed cluster. VolSync is used to set up volume replication between two clusters to protect CephFS-based PVCs. The replication feature is enabled by default.

13. Verify that the status of the OpenShift Data Foundation mirroring **daemon** health on the **Primary managed cluster** and the **Secondary managed cluster**.

```
$ oc get cephblockpoolradosnamespaces ocs-storagecluster-cephblockpool-builtin-implicit -n openshift-storage -o jsonpath='{.status.mirroringStatus.summary}'
```

Example output:

```
{"daemon_health":"OK","group_health":"OK","group_states":
{"health":"OK","image_health":"OK","image_states":{},"states":{}}
```

CAUTION

It could take up to 10 minutes for the **daemon_health** and **health** to go from **Warning** to **OK**. If the status does not become **OK** eventually, then use the RHACM console to verify that the Submariner connection between managed clusters is still in a healthy state. Do not proceed until all values are **OK**.

4.9. DISASTER RECOVERY PROTECTION FOR CLONED AND RESTORED RBD VOLUMES

Cloned and restored RBD volumes are copy-on-write clones of their corresponding data sources and remain linked to the parent volume to optimize space usage and Ceph performance. This also applies to cloned virtual machines that use cloned RBD-backed PersistentVolumeClaims (PVCs). To enable Regional Disaster Recovery, the RBD images backing such volumes must be flattened by separating them from the parent image.

Selecting the advanced option in Step 7 of [Creating Disaster Recovery Policy on Hub cluster](#) instructs CephCSI driver to perform necessary operations on the cloned or restored RBD volumes to prepare them for Disaster Recovery protection. The necessary operation has the following consequences:

- Space consumption may increase up to the total capacity of all cloned and restored RBD volumes being protected. It is recommended to make sure there is enough space available in the cluster to avoid the cluster getting full.
- Ceph cluster resource usage increases.
- The time taken for the operation to complete depends on the number of the RBD volumes, size of each volume, Ceph cluster resource availability, and client IO going on the RBD volume. Therefore, the amount of time required for completion of the operation cannot be determined.

4.10. CREATE SAMPLE APPLICATION FOR TESTING DISASTER RECOVERY SOLUTION

OpenShift Data Foundation disaster recovery (DR) solution supports disaster recovery for ApplicationSet-based applications that are managed by RHACM. For more details, see [ApplicationSet](#) documentation.

The following section details how to create an application and apply a DRPolicy to an application.

- [ApplicationSet-based applications](#)

OpenShift users that do not have cluster-admin permissions cannot create ApplicationSet-based applications.



NOTE

Regional-DR now supports leveraging RBD and CephFS volumes using non-default replica-2 storage classes that are managed by OpenShift Data Foundation. Support has also been extended to compression-enabled storage classes.

Multi-volume consistency for disaster recovery

By default, Red Hat OpenShift Data Foundation Disaster Recovery (DR) provides crash-consistent multi-volume consistency groups for Regional-DR to be used by applications that are deployed across multiple volumes.



IMPORTANT

If you have upgraded from a previous OpenShift Data Foundation version to 4.21, in order to enable consistency groups you must disable DR after the upgrade, and then re-enable DR.

4.10.1. ApplicationSet-based applications

4.10.1.1. Creating ApplicationSet-based applications

Prerequisite

- Ensure that the Red Hat OpenShift GitOps operator is installed on all three clusters: **Hub cluster**, **Primary managed cluster** and **Secondary managed cluster**. For instructions, see [Installing Red Hat OpenShift GitOps Operator in web console](#).
- On the Hub cluster, ensure that both Primary and Secondary managed clusters are registered to GitOps. For registration instructions, see [Registering managed clusters to GitOps](#). Then check if the Placement used by **GitOpsCluster** resource to register both managed clusters, has the tolerations to deal with cluster unavailability. You can verify if the following tolerations are added to the Placement using the command **oc get placement <placement-name> -n openshift-gitops -o yaml**.

tolerations:

- key: cluster.open-cluster-management.io/unreachable
operator: Exists
- key: cluster.open-cluster-management.io/unavailable
operator: Exists

In case the tolerations are not added, see [Configuring application placement tolerations for Red Hat Advanced Cluster Management and OpenShift GitOps](#).

- Ensure that you have created the **ClusterRoleBinding** yaml on both the **Primary** and **Secondary managed** clusters. For instruction, see the [Prerequisites chapter in RHACM documentation](#).

Procedure

1. On the Hub cluster, navigate to **All Clusters → Applications** and click **Create application**.
2. Choose the application type as **Argo CD ApplicationSet - Pull model**
3. In the **General** step, enter your **Application set name**.
4. Select **Argo server openshift-gitops** and **Requeue time** as **180** seconds.
5. Click **Next**.
6. In the Repository location for resources section, select **Repository type Git**.
7. Enter the Git repository URL for the sample application, the github Branch and Path where the resources busybox Pod and PVC will be created.
 - a. Use the sample application repository as <https://github.com/red-hat-storage/ocm-ramen-samples>
 - b. Select **Revision** as **main**
 - c. Choose one of the following Path:
 - **workloads/deployment/odr-regional-rbd** to use RBD Regional-DR.
 - **workloads/deployment/odr-regional-cephfs** to use CephFS Regional-DR.
8. Enter **Remote namespace** value. (example, busybox-sample) and click **Next**.
9. Choose the **Sync policy** settings as per your requirement or go with the default selections, and then click **Next**.
You can choose one or more options.
10. In Label expressions, add a label `<name>` with its value set to the **managed cluster** name.
11. Click **Next**.
12. Review the setting details and click **Submit**.

4.10.1.2. Apply Data policy to sample ApplicationSet-based application

Prerequisites

- Ensure that both managed clusters referenced in the Data policy are reachable. If not, the application will not be protected for disaster recovery until both clusters are online.

Procedure

1. On the Hub cluster, navigate to **All Clusters → Applications**.
2. Click the Actions menu at the end of application to view the list of available actions.

3. Click **Manage disaster recovery**.
4. Click **Enroll application**.
5. Select **Policy name** and click **Next**.
6. Select an **Application resource** and then use **PVC label selector** to select **PVC label** for the selected application resource.

**NOTE**

You can select more than one PVC label for the selected application resources.

7. Click **Next**.
8. In the **Enroll managed application** modal, review the policy configuration details and click **Assign**. The newly assigned Data policy details are displayed on the **Manage disaster recovery** modal.
9. Verify that you can view the assigned policy details on the Applications page.
 - a. On the Applications page, navigate to the **DR Status** column and view the status. The status can be any of the following:
 - **Protecting**: Shown during initial DR setup before data synchronization begins
 - **Healthy**: All volumes and Kubernetes resources are syncing normally
 - **Warning**: Volumes or Kubernetes resources are syncing slower than usual
 - **Critical**: Volumes or Kubernetes resources are not syncing
 - **Protection Error**: An error occurred during application protection setup
You can click the status to view the last synced time for application volumes and the DR policy assigned.

**NOTE**

It may take a few minutes for the **DR Status** to move from **Protecting** to **Healthy** after disaster recovery is enabled. The **DR status** will show **Protecting** during initial setup, then automatically transition to health-based statuses (**Healthy/Warning/Critical**) once data synchronization begins.

- b. Failover and relocate status can be viewed in the **DR Status** column when initiated. These statuses can be clicked to view more details about the action.
DR Status shows either **Protecting**, **Healthy**, **Warning**, or **Critical** when failover or relocate is not taking place. During failover or relocate operations, the status will show the operation-specific status (**FailingOver**, **Relocating**, etc.).
10. In the user interface, verify the volume consistency groups:
 - a. In the **Overview** tab, at the end of the protected application row from the action menu (⋮), select **Manage disaster recovery**.
 - b. In the modal, verify **Volume group replication** is *Enabled*.

- c. In the same modal, select **View volume groups**.
 - d. Verify the volume consistency group(s) has the correct volumes and a *Synced* status.
11. Optional: For Rados block device (RBD), verify **volumereplication**, **volumereplicationgroup**, and **volume group replication** on the primary cluster.

For **volumereplication**:

```
$ oc get volumereplication -A
```

Example output:

NAMESPACE	NAME	AGE	VOLUMEREPLICATIONCLASS	DESIREDSTATE	CURRENTSTATE
SOURCEKIND	SOURCENAME				
busybox-sample	vr-bc4a166e-13bc-45cd-99f0-8bd516f7fcc	112s	rd-		
volumereplicationclass-1625360775	VolumeGroupReplication	vgr-			
07a0050bee1f25720f5cae8d5daa6230-busybox-placement-drpc	primary			Primary	

For **volumereplicationgroup**:

```
$ oc get volumereplicationgroup -A
```

Example output:

NAMESPACE	NAME	DESIREDSTATE	CURRENTSTATE
busybox-sample	busybox-placement-drpc	primary	Primary

For **volume group replication**:

```
$ oc get volume group replication -A
```

Example output:

NAMESPACE	NAME	VOLUME GROUP REPLICATION CLASS	VOLUME GROUP REPLICATION CONTENT	DESIREDSTATE	CURRENTSTATE	AGE
busybox-sample	vgr-07a0050bee1f25720f5cae8d5daa6230-busybox-placement-drpc	rd-				
volume group replication class-1625360775-1602718344	vgrcontent-bc4a166e-13bc-45cd-99f0-8bd516f7fcc	primary	Primary	114s		

12. Optional: For CephFS, verify **ReplicationGroupSource** and **ReplicationSource** on the Primary cluster.

For **ReplicationGroupSource**:

```
$ oc get ReplicationGroupSource -A
```

Example output:

NAMESPACE	NAME	LAST SYNC	DURATION
NEXT SYNC	SOURCE		

LAST SYNC START

```
busybox-cephfs busybox-cephfs-bbd241b65b6589782fddaa82636e358d 2026-03-18T11:20:14Z 1m16.462467415s 2026-03-18T11:25:00Z {"matchLabels":{"ramendr.openshift.io/consistency-group":"busybox-cephfs-bbd241b65b6589782fddaa82636e358d"}}
```

For **ReplicationSource**:

```
$ oc get ReplicationSource -A
```

Example output:

NAMESPACE	NAME	SOURCE	LAST SYNC	DURATION	NEXT SYNC
busybox-cephfs	busybox-pvc	vs-busybox-pvc	2026-03-18T11:20:14Z	16.450725976s	

- Optional: For CephFS, verify **ReplicationGroupDestination** and **ReplicationDestination** on the Secondary cluster.

For **ReplicationGroupDestination**:

```
$ oc get ReplicationGroupDestination -A
```

Example output:

NAMESPACE	NAME	LAST SYNC	DURATION
busybox-cephfs	busybox-cephfs-b1f2b5bac27b3d8805437424895d212c	2026-03-18T11:20:15Z	47.700328282s

For **ReplicationDestination**:

```
$ oc get ReplicationDestination -A
```

Example output:

NAMESPACE	NAME	LAST SYNC	DURATION	NEXT SYNC
busybox-cephfs	busybox-pvc	2026-03-18T11:20:15Z	47.401909636s	

4.10.2. Deleting sample application

This section provides instructions for deleting the sample application **busybox** using the RHACM console.



IMPORTANT

When deleting a DR protected application, access to both clusters that belong to the DRPolicy is required. This is to ensure that all protected API resources and resources in the respective S3 stores are cleaned up as part of removing the DR protection. If access to one of the clusters is not healthy, deleting the **DRPlacementControl** resource for the application, on the hub, would remain in the Deleting state.

Prerequisites

- These instructions to delete the sample application should not be executed until the failover and relocate testing is completed and the application is ready to be removed from RHACM and the managed clusters.

Procedure

1. On the RHACM console, navigate to **Applications**.
2. Search for the sample application to be deleted (for example, **busybox**).
3. Click the Action Menu () next to the application you want to delete.
4. Click **Delete application**.
When the **Delete application** is selected a new screen will appear asking if the application related resources should also be deleted.
5. Select **Remove application related resources** checkbox to delete the Subscription and PlacementRule.
6. Click **Delete**. This will delete the busybox application on the Primary managed cluster (or whatever cluster the application was running on).
7. In addition to the resources deleted using the RHACM console, delete the **DRPlacementControl** if it is not auto-deleted after deleting the **busybox** application.
 - a. Log in to the OpenShift Web console for the Hub cluster and navigate to Installed Operators for the project **busybox-sample**.
For ApplicationSet applications, select the project as **openshift-gitops**.
 - b. Click **OpenShift DR Hub Operator** and then click the **DRPlacementControl** tab.
 - c. Click the Action Menu () next to the **busybox** application DRPlacementControl that you want to delete.
 - d. Click **Delete DRPlacementControl**.
 - e. Click **Delete**.



NOTE

This process can be used to delete any application with a **DRPlacementControl** resource.

4.11. VIRTUAL MACHINE GARBAGE COLLECTION AUTOMATION

Virtual machine (VM) garbage collection automation provides an automated mechanism to clean up VM resources on the source cluster during Disaster Recovery (DR) failover or relocation workflows. This reduces manual effort and helps users perform DR operations more quickly, supporting lower Recovery Time Objectives (RTO).

Prerequisites

- Virtual machine protection must be enabled through the OpenShift Data Foundation Virtualization console.

How it works

During DR failover or relocation, VM resources that are protected and orchestrated by the OpenShift Data Foundation Virtualization console are automatically deleted from the source cluster.

This deletion uses Kubernetes' owner-based garbage collection principles, ensuring that associated child resources, such as disks and pods, are also removed cleanly as part of normal Kubernetes garbage collection.

Known risks and limitations

- Automation applies only to VM resources whose protection is managed by the OpenShift Data Foundation Virtualization console.
- Other discovered application workload types, including GitOps-managed and other non-VM resources, remain out of scope and require manual cleanup.
- Automation is limited to resources directly owned by the VM lifecycle.

4.12. SUBSCRIPTION-BASED APPLICATION FAILOVER BETWEEN MANAGED CLUSTERS

Failover is a process that transitions an application from a primary cluster to a secondary cluster in the event of a primary cluster failure. While failover provides the ability for the application to run on the secondary cluster with minimal interruption, making an uninformed failover decision can have adverse consequences, such as complete data loss in the event of unnoticed replication failure from primary to secondary cluster. If a significant amount of time has gone by since the last successful replication, it's best to wait until the failed primary is recovered.

`LastGroupSyncTime` is a [critical metric](#) that reflects the time since the last successful replication occurred for all PVCs associated with an application. In essence, it measures the synchronization health between the primary and secondary clusters. So, prior to initiating a failover from one cluster to another, check for this metric and only initiate the failover if the `LastGroupSyncTime` is within a reasonable time in the past.



NOTE

During the course of failover the Ceph-RBD mirror deployment on the failover cluster is scaled down to ensure a clean failover for volumes that are backed by Ceph-RBD as the storage provisioner.

Prerequisites

- If your setup has active and passive RHACM hub clusters, see [Hub recovery using Red Hat Advanced Cluster Management](#).
- If you are protecting virtual machines (VMs) with DR, you must ensure that the target VMs have successfully transitioned to the **Running** state before triggering a failover. Do not enroll or initiate a failover for a VM that is in any of the following intermediate or pending phases:
 - **Provisioning**
 - **Scheduling**
 - **CreatingDataVolume**

- **Starting**

- When the primary cluster is in a state other than **Ready**, check the actual status of the cluster as it might take some time to update.
 1. Navigate to the **RHACM console → Infrastructure → Clusters → Cluster list** tab.
 2. Check the status of both the managed clusters individually before performing failover operation.
However, failover operation can still be performed when the cluster you are failing over to is in a *Ready* state.
- Run the following command on the Hub Cluster to check if **lastGroupSyncTime** is within an acceptable data loss window, when compared to current time.

```
$ oc get drpc -o yaml -A | grep lastGroupSyncTime
```

Example output:

```
[...]
lastGroupSyncTime: "2023-07-10T12:40:10Z"
```

Procedure

1. Failover the protected application using either of the following paths:

Option 1: From the **Applications** page

 - On the Hub cluster, navigate to **Applications**. Click **Actions (⋮)** and select **Failover application**.

Option 2: From the **Protected Applications** page

 - On the Hub cluster, navigate to **Data Services → Disaster recovery → Protected Applications → Actions (⋮) → Failover**.
2. After the **Failover application** modal is shown, select **policy** and **target cluster** to which the associated application will failover in case of a disaster.
3. Click the **Select subscription group** dropdown to verify the default selection or modify this setting.
By default, the subscription group that replicates for the application resources is selected.
4. Check the status of the **Failover readiness**.
 - If the status is **Ready** with a green tick, it indicates that the target cluster is ready for failover to start. Proceed to step 7.
 - If the status is **Unknown** or **Not ready**, then wait until the status changes to **Ready**.



IMPORTANT

If there are data inconsistencies caused by synchronization delays, a warning message appears stating **Inconsistent data on target cluster**. This alerts to the possibility of data loss if the failover is initiated. The message is no longer displayed when data synchronization is complete.

5. Click **Initiate**. The busybox application is now failing over to the **Secondary-managed cluster**.
6. In the **DR status** column, you can view the progress of the failover. The modal shows the progress of the failover as **Preparing → Failover → Restoring → Clean up**.
7. Verify that the activity status shows as **FailedOver** for the application.
 - a. Navigate to the **Applications → Overview** tab.
 - b. In the **Data policy** column, click the **policy** link for the application you applied the policy to.
 - c. On the **Data policy** popover, click the **View more details** link.
 - d. Verify that you can see one or more policy names and the ongoing activities (Last sync time and Activity status) associated with the policy in use with the application.

4.13. APPLICATIONSET-BASED APPLICATION FAILOVER BETWEEN MANAGED CLUSTERS

Failover is a process that transitions an application from a primary cluster to a secondary cluster in the event of a primary cluster failure. While failover provides the ability for the application to run on the secondary cluster with minimal interruption, making an uninformed failover decision can have adverse consequences, such as complete data loss in the event of unnoticed replication failure from primary to secondary cluster. If a significant amount of time has gone by since the last successful replication, it's best to wait until the failed primary is recovered.

[LastGroupSyncTime](#) is a critical metric that reflects the time since the last successful replication occurred for all PVCs associated with an application. In essence, it measures the synchronization health between the primary and secondary clusters. So, prior to initiating a failover from one cluster to another, check for this metric and only initiate the failover if the LastGroupSyncTime is within a reasonable time in the past.



NOTE

During the course of failover the Ceph-RBD mirror deployment on the failover cluster is scaled down to ensure a clean failover for volumes that are backed by Ceph-RBD as the storage provisioner.

Prerequisites

- If your setup has active and passive RHACM hub clusters, see [Hub recovery using Red Hat Advanced Cluster Management](#).
- If you are protecting virtual machines (VMs) with DR, you must ensure that the target VMs have successfully transitioned to the **Running** state before triggering a failover.
Do not enroll or initiate a failover for a VM that is in any of the following intermediate or pending phases:
 - **Provisioning**
 - **Scheduling**
 - **CreatingDataVolume**
 - **Starting**

- When the primary cluster is in a state other than **Ready**, check the actual status of the cluster as it might take some time to update.
 1. Navigate to the **RHACM console → Infrastructure → Clusters → Cluster list** tab.
 2. Check the status of both the managed clusters individually before performing failover operation.
However, failover operation can still be performed when the cluster you are failing over to is in a *Ready* state.
- Run the following command on the Hub Cluster to check if **lastGroupSyncTime** is within an acceptable data loss window, when compared to current time.

```
$ oc get drpc -o yaml -A | grep lastGroupSyncTime
```

Example output:

```
[...]
lastGroupSyncTime: "2023-07-10T12:40:10Z"
```

Procedure


1. Failover the protected application using either of the following paths:

Option 1: From the **Applications** page

 - On the Hub cluster, navigate to **Applications**. Click **Actions (⋮)** and select **Failover application**.

Option 2: From the **Protected Applications** page

 - On the Hub cluster, navigate to **Data Services → Disaster recovery → Protected Applications → Actions (⋮) → Failover**.
2. When the **Failover application** modal is shown, verify the details presented are correct and check the status of the Failover **readiness**. If the status is **Ready** with a green tick, it indicates that the target cluster is ready for failover to start.



IMPORTANT

If there are data inconsistencies caused by synchronization delays, a warning message appears stating **Inconsistent data on target cluster**. This alerts to the possibility of data loss if the failover is initiated. The message is no longer displayed when data synchronization is complete.

3. Click **Initiate**. The busybox resources are now created on the target cluster.
4. In the **DR status** column, you can view the progress of the failover. The modal shows the progress of the failover as **Preparing → Failover → Restoring → Clean up**.
5. Verify that the activity status shows as **FailedOver** for the application.
 - a. Navigate to the **Applications → Overview** tab.
 - b. In the **Data policy** column, click the **policy** link for the application you applied the policy to.

- c. On the **Data policy** popover, verify that you can see one or more policy names and the ongoing activities associated with the policy in use with the application.

4.14. RELOCATING SUBSCRIPTION-BASED APPLICATION BETWEEN MANAGED CLUSTERS

Relocate an application to its preferred location when all managed clusters are available.

Prerequisite

- If your setup has active and passive RHACM hub clusters, see [Hub recovery using Red Hat Advanced Cluster Management](#).
- If you are protecting virtual machines (VMs) with DR, you must ensure that the target VMs have successfully transitioned to the **Running** state before triggering a relocate.
Do not enroll or initiate a relocate for a VM that is in any of the following intermediate or pending phases:
 - **Provisioning**
 - **Scheduling**
 - **CreatingDataVolume**
 - **Starting**
- When the primary cluster is in a state other than **Ready**, check the actual status of the cluster as it might take some time to update. Relocate can only be performed when both primary and preferred clusters are up and running.
 1. Navigate to **RHACM console** → **Infrastructure** → **Clusters** → **Cluster list** tab.
 2. Check the status of both the managed clusters individually before performing relocate operation.
- Perform relocate when **lastGroupSyncTime** is within the replication interval (for example, 5 minutes) when compared to current time. This is recommended to minimize the Recovery Time Objective (RTO) for any single application.
Run this command on the Hub Cluster:

```
$ oc get drpc -o yaml -A | grep lastGroupSyncTime
```

Example output:

```
[...]
lastGroupSyncTime: "2023-07-10T12:40:10Z"
```

Compare the output time (UTC) to current time to validate that all **lastGroupSyncTime** values are within their application replication interval. If not, wait to Relocate until this is true for all **lastGroupSyncTime** values.

Procedure

1. Relocate the protected application using either of the following paths:

Option 1: From the **Applications** page

- On the Hub cluster, navigate to **All Clusters → Applications**.
- In the **Overview** tab, at the end of the application row from the action menu (⋮), select **Manage disaster recovery**.
- Click **Relocate**.

Option 2: From the **Protected Applications** page

On the Hub cluster, navigate to **Data Services → Disaster recovery → Protected Applications → Actions (⋮) → Relocate**.

- When the **Relocate application** modal is shown, select **policy** and **target cluster** to which the associated application will relocate to in case of a disaster.
- By default, the subscription group that will deploy the application resources is selected. Click the **Select subscription group** dropdown to verify the default selection or modify this setting.
- Check the status of the **Relocation readiness**.
 - If the status is **Ready** with a green tick, it indicates that the target cluster is ready for relocation to start. Proceed to step 7.
 - If the status is **Unknown** or **Not ready**, then wait until the status changes to **Ready**.

**IMPORTANT**

If there are data inconsistencies caused by synchronization delays, a warning message appears stating **Inconsistent data on target cluster**. This alerts to the possibility of data loss if the relocate is initiated. The message is no longer displayed when data synchronization is complete.

- Click **Initiate**. The busybox resources are now created on the target cluster.
- In the **DR status** column, you can view the progress of the failover. The modal shows the progress of the failover as **Preparing → Syncing → Restoring → Clean up**.
- Verify that the activity status shows as **Relocated** for the application.
 - Navigate to the **Applications → Overview** tab.
 - In the **Data policy** column, click the **policy** link for the application you applied the policy to.
 - On the **Data policy** popover, click the **View more details** link.
 - Verify that you can see one or more policy names and the ongoing activities (Last sync time and Activity status) associated with the policy in use with the application.

4.15. RELOCATING AN APPLICATIONSET-BASED APPLICATION BETWEEN MANAGED CLUSTERS

Relocate an application to its preferred location when all managed clusters are available.

Prerequisite

- If your setup has active and passive RHACM hub clusters, see [Hub recovery using Red Hat Advanced Cluster Management](#).
- If you are protecting virtual machines (VMs) with DR, you must ensure that the target VMs have successfully transitioned to the **Running** state before triggering a relocate.
Do not enroll or initiate a relocate for a VM that is in any of the following intermediate or pending phases:
 - **Provisioning**
 - **Scheduling**
 - **CreatingDataVolume**
 - **Starting**
- When the primary cluster is in a state other than **Ready**, check the actual status of the cluster as it might take some time to update. Relocate can only be performed when both primary and preferred clusters are up and running.
 1. Navigate to **RHACM console** → **Infrastructure** → **Clusters** → **Cluster list** tab.
 2. Check the status of both the managed clusters individually before performing relocate operation.
- Perform relocate when **lastGroupSyncTime** is within the replication interval (for example, 5 minutes) when compared to current time. This is recommended to minimize the Recovery Time Objective (RTO) for any single application.
Run this command on the Hub Cluster:

```
$ oc get drpc -o yaml -A | grep lastGroupSyncTime
```

Example output:

```
[...]
lastGroupSyncTime: "2023-07-10T12:40:10Z"
```

Compare the output time (UTC) to current time to validate that all **lastGroupSyncTime** values are within their application replication interval. If not, wait to Relocate until this is true for all **lastGroupSyncTime** values.

Procedure

1. Relocate the protected application using either of the following paths:

Option 1: From the **Applications** page

 - On the Hub cluster, navigate to **All Clusters** → **Applications**.
 - In the **Overview** tab, at the end of the application row from the action menu (⋮), select **Manage disaster recovery**.
 - Click **Relocate**.

Option 2: From the **Protected Applications** page

On the Hub cluster, navigate to **Data Services → Disaster recovery → Protected Applications → Actions (⚙) → Relocate** .

2. When the **Relocate application** modal is shown, select **policy** and **target cluster** to which the associated application will relocate to in case of a disaster.



IMPORTANT

If there are data inconsistencies caused by synchronization delays, a warning message appears stating **Inconsistent data on target cluster**. This alerts to the possibility of data loss if the relocate is initiated. The message is no longer displayed when data synchronization is complete.

3. Click **Initiate**. The busybox resources are now created on the target cluster.
4. Close the modal window and track the status using the **DR Status** column on the Applications page. The status moves from **Preparing → Syncing → Restoring → Clean up**.
5. Verify that the activity status shows as **Relocated** for the application.
 - a. Navigate to the **Applications → Overview** tab.
 - b. In the **DR Status** column, click the **policy** link for the application you applied the policy to.
 - c. On the **Data policy** popover, verify that you can see one or more policy names and the relocation status associated with the policy in use with the application.

4.16. DISASTER RECOVERY PROTECTION FOR DISCOVERED APPLICATIONS

Red Hat OpenShift Data Foundation now provides disaster recovery (DR) protection and support for workloads that are deployed in one of the managed clusters directly without using Red Hat Advanced Cluster Management (RHACM). These workloads are called discovered applications.

Workloads deployed using RHACM are termed managed applications, while those deployed directly on one of the managed clusters without using RHACM are called discovered applications. Although RHACM displays the details of both types of workloads, it does not manage the lifecycle (create, delete, edit) of discovered applications.

4.16.1. Prerequisites for disaster recovery protection of discovered applications

This section provides instructions to guide you through the prerequisites for protecting discovered applications. This includes tasks such as assigning a data policy and initiating DR actions such as failover and relocate.

1. Ensure that all the DR configurations have been installed on the Primary managed cluster and the Secondary managed cluster.
2. Install the OADP 1.4 operator (or greater).



NOTE

Any version before OADP 1.4 will not work for protecting discovered applications.

- a. On the **Primary** and **Secondary managed cluster**, navigate to **Ecosystem → Software Catalog** and use the keyword filter to search for **OADP**.
- b. Click the **OADP** tile.
- c. Keep all default settings and click **Install**. Ensure that the operator resources are installed in the **openshift-adp** project.



NOTE

If OADP 1.4 is installed after DR configuration has been completed then the **ramen-dr-cluster-operator** pods on the **Primary managed cluster** and the **Secondary managed cluster** in namespace **openshift-dr-system** must be restarted (deleted and recreated).

3. [Optional] Add caCertificates to **ramen-hub-operator-config ConfigMap**.
Configure network (SSL) access between the primary and secondary clusters so that metadata can be stored on the alternate cluster in a Multicloud Gateway (MCG) object bucket using a secure transport protocol and in the Hub cluster for verifying access to the object buckets.



NOTE

If all of your OpenShift clusters are deployed using a signed and valid set of certificates for your environment then this section can be skipped.

If you are using self-signed certificates, then you have already created a **ConfigMap** named **user-ca-bundle** in the **openshift-config** namespace and added this ConfigMap to the default Proxy cluster resource. **This means you need to add the `caCertificates` parameter to the configmap `ramen-hub-operator-config` with the encoded value.**

- a. Find the encoded value for the caCertificates.

```
$ oc get configmap user-ca-bundle -n openshift-config -o jsonpath="{['data']['ca-bundle.crt']}" |base64 -w 0
```

- b. Add this base64 encoded value to the configmap **ramen-hub-operator-config** on the Hub cluster. Example below shows where to add caCertificates.

```
$ oc edit configmap ramen-hub-operator-config -n openshift-operators
```

```
[...]
ramenOpsNamespace: openshift-dr-ops
s3StoreProfiles:
- s3Bucket: odrbucket-36bceb61c09c
  s3CompatibleEndpoint: https://s3-openshift-storage.apps.hyper3.vmw.ibmfusion.eu
  s3ProfileName: s3profile-hyper3-ocs-storagecluster
  s3Region: noobaa
  s3SecretRef:
    name: 60f2ea6069e168346d5ad0e0b5faa59bb74946f
  caCertificates: {input base64 encoded value here}
- s3Bucket: odrbucket-36bceb61c09c
  s3CompatibleEndpoint: https://s3-openshift-storage.apps.hyper4.vmw.ibmfusion.eu
  s3ProfileName: s3profile-hyper4-ocs-storagecluster
```

```
s3Region: noobaa
s3SecretRef:
  name: cc237eba032ad5c422fb939684eb633822d7900
caCertificates: {input base64 encoded value here}
```

4. Verify that there are DR secrets created in the OADP operator default namespace **openshift-adp** on the **Primary managed cluster** and the **Secondary managed cluster**. The DR secrets that were created when the first DRPolicy was created, will be similar to the secrets below. The DR secret name is preceded with the letter **v**.

```
$ oc get secrets -n openshift-adp
NAME                                     TYPE    DATA  AGE
v60f2ea6069e168346d5ad0e0b5faa59bb74946f Opaque  1      3d20h
vcc237eba032ad5c422fb939684eb633822d7900 Opaque  1      3d20h
[...]
```



NOTE

There will be one DR created secret for each managed cluster in the **openshift-adp** namespace.

5. Verify if the Data Protection Application (DPA) is already installed on each managed cluster in the OADP namespace **openshift-adp**. If not already created then follow the next step to create this resource.

- a. Create the DPA by copying the following YAML definition content to **dpa.yaml**.

```
apiVersion: oadp.openshift.io/v1alpha1
kind: DataProtectionApplication
metadata:
  labels:
    app.kubernetes.io/component: velero
  name: velero
  namespace: openshift-adp
spec:
  backupImages: false
  configuration:
    nodeAgent:
      enable: false
      uploaderType: restic
    velero:
      defaultPlugins:
        - openshift
        - aws
        - kubvirt
      noDefaultBackupLocation: true
```

- b. Create the DPA resource.

```
$ oc create -f dpa.yaml -n openshift-adp
```

```
dataprotectionapplication.oadp.openshift.io/velero created
```

- c. Verify that the OADP resources are created and are in **Running** state.

```
$ oc get pods,dpa -n openshift-adp
```

NAME	READY	STATUS	RESTARTS	AGE
pod/openshift-adp-controller-manager-7b64b74fcd-msjbs	1/1	Running	0	5m30s
pod/velero-694b5b8f5c-b4kww	1/1	Running	0	3m31s


```
$ oc get dataprotectionapplication.oadp.openshift.io/velero
```

NAME	AGE
dataprotectionapplication.oadp.openshift.io/velero	3m31s

4.16.2. Creating a sample discovered application

In order to test **failover** from the **Primary managed cluster** to the **Secondary managed cluster** and **relocate** for discovered applications, you need a sample application that is installed without using the RHACM create application capability.



NOTE

Regional-DR now supports leveraging CephRBD volumes using non-default replica-2 storage classes that are managed by OpenShift Data Foundation. Support has also been extended to compression-enabled storage classes.

Procedure

1. Log in to the **Primary managed cluster** and clone the sample application repository.

```
$ git clone https://github.com/red-hat-storage/ocm-ramen-samples.git
```

2. Verify that you are on the **main** branch.

```
$ cd ~/ocm-ramen-samples
$ git branch
* main
```

The correct directory should be used when creating the sample application based on your scenario, metro or regional. To find your directory:

```
$ ls workloads/deployment | egrep -v 'k8s|base'
odr-metro-rbd
odr-regional-rbd
odr-regional-cephfs
```

3. Create a project named **busybox-discovered** on both the **Primary and Secondary managed clusters**.

```
$ oc new-project busybox-discovered
```

4. Create the **busybox** application on the **Primary managed cluster**.

This sample application example is for Regional-DR using a Ceph RBD volume. CephFS volumes can be used as well.

—


```
$ oc apply -k workloads/deployment/odr-regional-rbd -n busybox-discovered
persistentvolumeclaim/busybox-pvc created
deployment.apps/busybox created
```

**NOTE**

OpenShift Data Foundation Disaster Recovery solution now extends protection to discovered applications that span across multiple namespaces.

5. Verify that busybox is running in the correct project on the **Primary managed cluster**.

```
$ oc get pods,pvc,deployment -n busybox-discovered
```

```
NAME                                READY STATUS  RESTARTS  AGE
pod/busybox-796fccbb95-qmxjf  1/1   Running    0         18s
```

```
NAME                                STATUS  VOLUME                                     CAPACITY  ACCESS
MODES STORAGECLASS                  VOLUMEATTRIBUTESCLASS AGE
persistentvolumeclaim/busybox-pvc  Bound  pvc-b20e4129-902d-47c7-b962-
040ad64130c4  1Gi    RWO    ocs-storagecluster-ceph-rbd  <unset>  18s
```

```
NAME                                READY UP-TO-DATE  AVAILABLE  AGE
deployment.apps/busybox  1/1    1             1          18
```

4.16.3. Enrolling a sample discovered application for disaster recovery protection

This section guides you on how to apply an existing DR Policy to a discovered application from the **Protected applications** tab.

Prerequisites

- Ensure that Disaster Recovery has been configured and that at least one DR Policy has been created.

Procedure

1. On RHACM console, navigate to **Disaster recovery** → **Protected applications** tab.
2. Click **Enroll application** to start configuring existing applications for DR protection.
3. Select **ACM discovered applications**.
4. In the Namespace page, choose the **DR cluster** which is the name of the **Primary managed cluster** where **busybox** is installed.
5. Select **namespace** where the application is installed. For example, **busybox-discovered**.

**NOTE**

If you have workload spread across multiple namespaces then you can select all of those namespaces to DR protect.

6. Choose a unique **Name**, for example **busybox-rbd**, for the discovered application and click **Next**.
7. In the Configuration page, select either **Resource label** or **Recipe**.
8. **Resource label** is used to protect your resources where you can set which resources will be included in the **kubernetes-object** backup and what volume's persistent data will be replicated. If you selected **Resource label**, provide label expressions and PVC label selector. Choose the label **appname=busybox** for both the **kubernetes-objects** and for the **PVC(s)**.
9. If you selected **Recipe**, then from the **Recipe list** select the name of the recipe.



IMPORTANT

The recipe resource must be created in the application namespace on both managed clusters before enrolling an application for disaster recovery.

10. [Optional]: If you are using a recipe and want to add the recipe's parameters, under **Recipe parameters**, add the **Key** and **Value (optional)**, then click **Add parameter**.
11. Click **Next**.
12. In the Replication page, select an existing **DR Policy** and the **kubernetes-objects backup interval**.



NOTE

It is recommended to choose the same duration for the PVC data replication and kubernetes-object backup interval (i.e., 5 minutes).

13. Click **Next**.
14. Review the configuration and click **Save**.
Use the **Back** button to go back to the screen to correct any issues.
15. Verify that you can view the assigned policy details on the Applications page.
 - a. On the Applications page, navigate to the **DR Status** column and view the status. The status can be any of the following:
 - **Protecting**: Shown during initial DR setup before data synchronization begins
 - **Healthy**: All volumes and Kubernetes resources are syncing normally
 - **Warning**: Volumes or Kubernetes resources are syncing slower than usual
 - **Critical**: Volumes or Kubernetes resources are not syncing
 - **Protection Error**: An error occurred during application protection setup
You can click the status to view the last synced time for application volumes and the DR policy assigned.

**NOTE**

It may take a few minutes for the **DR Status** to move from **Protecting** to **Healthy** after disaster recovery is enabled. The **DR status** will show **Protecting** during initial setup, then automatically transition to health-based statuses (**Healthy/Warning/Critical**) once data synchronization begins.

16. Verify that the volume consistency groups have been enabled from the **Protected applications** tab:
 - a. Click the arrow next to the DRPC to expand the details.
 - b. Click the number under the **Volume Consistency groups** column to open the modal.
 - c. In the **Volume consistency groups** modal, verify the volume consistency group(s) has the correct volumes and a *Synced* status.
17. View the status of DRPC and VolumeReplicationGroup (VRG).
 - a. To see the status of the DRPC, run the following command on the Hub cluster:

```
$ oc get drpc {drpc_name} -o wide -n openshift-dr-ops
```

The discovered applications store resources such as DRPlacementControl (DRPC) and Placement on the Hub cluster in a new namespace called **openshift-dr-ops**. The DRPC name can be identified by the unique Name configured in prior steps (i.e., **busybox-rbd**).

- b. To see the status of the VRG for discovered applications, run the following command on the managed cluster where the busybox application was manually installed.

```
$ oc get vrg {vrg_name} -n openshift-dr-ops
```

The VRG resource is stored in the namespace **openshift-dr-ops** after a DR Policy is assigned to the discovered application. The VRG name can be identified by the unique Name configured in prior steps (i.e., **busybox-rbd**).

18. Optional: For Rados block device (RBD), verify **volumereplication** and **volumegroupreplication** on the primary cluster.
For **volumereplication**:

```
$ oc get volumereplication -A
```

Example output:

NAMESPACE	NAME	AGE	VOLUMEREPLICATIONCLASS
SOURCEKIND	SOURCENAME		DESIREDSTATE
CURRENTSTATE			
busybox-discovered	vr-f9fc057c-5e95-47cb-b515-b4ee19d4ca8e	109s	rbd-
volumereplicationclass-1625360775	VolumeGroupReplication	vgr-	
e8b7cca8211a41455790dad2707a08c-busybox-rbd	primary	Primary	

For **volumegroupreplication**:

```
$ oc get volumegroupreplication -A
```

Example output:

```
NAMESPACE      NAME
VOLUMEGROUPREPLICATIONCLASS
VOLUMEGROUPREPLICATIONCONTENT      DESIREDSTATE
CURRENTSTATE  AGE
busybox-discovered  vgr-e8b7cca8211a41455790dadc2707a08c-busybox-rbd  rbd-
volumegroupreplicationclass-1625360775-1602718344  vgrcontent-f9fc057c-5e95-47cb-
b515-b4ee19d4ca8e  primary      Primary      113s
```

19. Optional: For CephFS, verify **ReplicationGroupSource** and **ReplicationSource** on the Primary cluster.

For **ReplicationGroupSource**:

```
$ oc get ReplicationGroupSource -A
```

Example output:

```
NAMESPACE      NAME                                LAST SYNC      DURATION
NEXT SYNC      SOURCE
LAST SYNC START
busybox-cephfs  busybox-cephfs-bbd241b65b6589782fddaa82636e358d  2026-03-
18T11:20:14Z  1m16.462467415s  2026-03-18T11:25:00Z  {"matchLabels":
{"ramendr.openshift.io/consistency-group":"busybox-cephfs-
bbd241b65b6589782fddaa82636e358d"}}
```

For **ReplicationSource**:

```
$ oc get ReplicationSource -A
```

Example output:

```
NAMESPACE      NAME      SOURCE      LAST SYNC      DURATION      NEXT
SYNC
busybox-cephfs  busybox-pvc  vs-busybox-pvc  2026-03-18T11:20:14Z  16.450725976s
```

20. Optional: For CephFS, verify **ReplicationGroupDestination** and **ReplicationDestination** on the Secondary cluster.

For **ReplicationGroupDestination**:

```
$ oc get ReplicationGroupDestination -A
```

Example output:

```
NAMESPACE      NAME                                LAST SYNC      DURATION
LAST SYNC START
busybox-cephfs  busybox-cephfs-b1f2b5bac27b3d8805437424895d212c  2026-03-
18T11:20:15Z  47.700328282s  2026-03-18T11:20:15Z
```

For **ReplicationDestination**:

```
$ oc get ReplicationDestination -A
```

Example output:

NAMESPACE	NAME	LAST SYNC	DURATION	NEXT SYNC
busybox-cephfs	busybox-pvc	2026-03-18T11:20:15Z	47.401909636s	

4.16.4. Discovered application failover and relocate

A protected Discovered application can *Failover* or *Relocate* to its peer cluster similar to **managed applications**. However, there are some additional steps for discovered applications since RHACM does not manage the lifecycle of the application as it does for Managed applications.

This section guides you through the Failover and Relocate process for a protected discovered application.



IMPORTANT

Never initiate a Failover or Relocate of an application when one or both resource types are in a Warning or Critical status.

4.16.4.1. Failover disaster recovery protected discovered application

This section guides you on how to failover a discovered application which is disaster recovery protected.

Prerequisites

- Ensure that the application namespace is created in both managed clusters (for example, **busybox-discovered**).

Procedure

1. In the RHACM console, navigate to **Disaster Recovery → Protected applications** tab.
2. At the end of the application row, click on the **Actions** menu and choose to initiate *Failover*.
3. In the Failover application modal window, review the status of the application and the target cluster.



IMPORTANT

If there are data inconsistencies caused by synchronization delays, a warning message appears stating **Inconsistent data on target cluster**. This alerts to the possibility of data loss if the failover is initiated. The message is no longer displayed when data synchronization is complete.

4. Click **Initiate**.
5. In the **DR status** column, you can view the progress of the failover. The modal shows the progress of the failover as **Preparing → Failover → Restoring → Clean up**.
6. Verify that the busybox application is running on the **Secondary managed cluster**.

```
$ oc get pods,pvc,volumereplication -n busybox-discovered
NAME                                READY STATUS RESTARTS AGE
pod/busybox-796fccbb95-qmxjf 1/1   Running 0      2m46s
```

```
NAME                                STATUS VOLUME CAPACITY ACCESS
MODES STORAGECLASS VOLUMEATTRIBUTESCLASS AGE
persistentvolumeclaim/busybox-pvc Bound pvc-b20e4129-902d-47c7-b962-040ad64130c4 1Gi RWO ocs-storagecluster-ceph-rbd <unset>
2m57s
```

```
NAMESPACE NAME AGE VOLUMEREPLICATIONCLASS
SOURCEKIND SOURCENAME DESIREDSTATE
CURRENTSTATE
busybox-sample vr-5a034832-23d0-4376-a7c1-b996979494e7 2m27s rbd-
volumereplicationclass-1625360775 VolumeGroupReplication
c560635344cdf0a880feaabf38c3bc89busybox-placement-drpc primary Primary
```

7. Check the progression status of Failover until the result is **Clean up**. The DRPC name can be identified by the unique Name configured in prior steps (for example, **busybox-rbd**).

```
$ oc get drpc {drpc_name} -n openshift-dr-ops -o jsonpath='{.status.progression}'{"\n"}'
WaitOnUserToCleanUp
```

8. Remove the busybox application from the **Primary managed cluster** to complete the **Failover** process.
 - a. Navigate to the **Protected applications** tab. You will see a message to remove the application.
 - b. Navigate to the cloned repository for **busybox** and run the following commands on the **Primary managed cluster** where you **failed** over from. Use the same directory that was used to create the application (for example, **odr-regional-rbd**).

```
$ cd ~/ocm-ramen-samples/
$ git branch
* main
$ oc delete -k workloads/deployment/odr-regional-rbd -n busybox-discovered
persistentvolumeclaim "busybox-pvc" deleted
deployment.apps "busybox" deleted
```

9. After deleting the application, navigate to the **Protected applications** tab and verify that the busybox resources are both in *Healthy* status.


4.16.4.2. Relocate disaster recovery protected discovered application

This section guides you on how to relocate a discovered application which is disaster recovery protected.

Procedure

1. In the RHACM console, navigate to **Disaster Recovery → Protected applications** tab.
2. At the end of the application row, click on the **Actions** menu and choose to initiate *Relocate*.

3. In the Relocate application modal window, review the status of the application and the target cluster.



IMPORTANT

If there are data inconsistencies caused by synchronization delays, a warning message appears stating **Inconsistent data on target cluster**. This alerts to the possibility of data loss if the relocate is initiated. The message is no longer displayed when data synchronization is complete.

4. Click **Initiate**.
5. In the **DR status** column, you can view the progress of the relocation. The progress shows **Preparing → Clean up**, at which time you need to perform the following clean up activities.
6. The DRPC name can be identified by the unique Name configured in prior steps (for example, **busybox-rbd**).

```
$ oc get drpc {drpc_name} -n openshift-dr-ops -o jsonpath='{.status.progression}{"\n"}'
WaitOnUserToCleanUp
```

7. Remove the busybox application from the **Secondary managed cluster** before *Relocate* to the **Primary managed cluster** is completed.
Navigate to the cloned repository for **busybox** and run the following commands on the **Secondary managed cluster** where you **relocated** from. Use the same directory that was used to create the application (for example, **odr-regional-rbd**).

```
$ cd ~/ocm-ramen-samples/
$ git branch
* main
$ oc delete -k workloads/deployment/odr-regional-rbd -n busybox-discovered
persistentvolumeclaim "busybox-pvc" deleted
deployment.apps "busybox" deleted
```

8. After deleting the application, navigate to the **Protected applications** tab and verify that the busybox resources are both in *Healthy* status.
9. Verify that the **busybox** application is running on the **Primary managed cluster**.

```
$ oc get pods,pvc,volumereplication -n busybox-discovered
NAME                                READY STATUS RESTARTS AGE
pod/busybox-796fccbb95-qmxjf        1/1   Running 0       2m46s
```

```
NAME                                STATUS VOLUME                                CAPACITY ACCESS
MODES STORAGECLASS                  VOLUMEATTRIBUTESCLASS AGE
persistentvolumeclaim/busybox-pvc   Bound  pvc-b20e4129-902d-47c7-b962-
040ad64130c4 1Gi    RWO    ocs-storagecluster-ceph-rbd <unset>
2m57s
```

```
NAMESPACE NAME                                AGE VOLUMEREPLICATIONCLASS
SOURCEKIND SOURCENAME                                DESIREDSTATE
CURRENTSTATE
```

```

busybox-sample  vr-5a034832-23d0-4376-a7c1-b996979494e7  2m27s  rbd-
volumereplicationclass-1625360775  VolumeGroupReplication
c560635344cdf0a880feaabf38c3bc89busybox-placement-drpc  primary      Primary

```

- Once the clean up activities are performed, the **DR Status** column will move from **Syncing** → **Restoring**.

4.16.5. Disable disaster recovery for protected applications

This section guides you to disable disaster recovery resources when you want to delete the protected applications or when the application no longer needs to be protected.

Procedure

- Login to the Hub cluster.
- List the **DRPlacementControl** (DRPC) resources. Each DRPC resource was created when the application was assigned a DR policy.

```
$ oc get drpc -n openshift-dr-ops
```

- Find the **DRPC** that has a name that includes the unique identifier that you chose when assigning a DR policy (for example, **busybox-rbd**) and delete the DRPC.

```
$ oc delete {drpc_name} -n openshift-dr-ops
```

- List the **Placement** resources. Each Placement resource was created when the application was assigned a DR policy.

```
$ oc get placements -n openshift-dr-ops
```

- Find the **Placement** that has a name that includes the unique identifier that you chose when assigning a DR policy (for example, **busybox-rbd-placement-1**) and delete the **Placement**.

```
$ oc delete placements {placement_name} -n openshift-dr-ops
```

4.17. RECOVERING TO A REPLACEMENT CLUSTER WITH REGIONAL-DR

When there is a failure with the primary cluster, you get the options to either repair, wait for the recovery of the existing cluster, or replace the cluster entirely if the cluster is irredeemable. This solution guides you when replacing a failed primary cluster with a new cluster and enables failback (relocate) to this new cluster.

In these instructions, we are assuming that a RHACM managed cluster must be replaced after the applications have been installed and protected. For purposes of this section, the RHACM managed cluster is the **replacement cluster**, while the cluster that is not replaced is the **surviving cluster** and the new cluster is the **recovery cluster**.

Prerequisite

- Ensure that the Regional-DR environment has been configured with applications installed using Red Hat Advance Cluster Management (RHACM).
- Ensure that the applications are assigned a Data policy which protects them against cluster failure.

Procedure

1. On the surviving cluster, backup the **clientProfileMapping**:

```
$ oc get clientprofilemapping ocs-storagecluster -n openshift-storage -o yaml >
clientProfileMappingBeforeRecovery.yaml
```

2. Failover all protected applications on the failed replacement cluster to the surviving cluster on the RHACM console.

- **For managed applications:**

You can failover protected applications using either of the following paths:

Option 1: From the **Applications** page

- On the Hub cluster, navigate to **Applications**. Click **Actions (⋮)** and select **Failover application**.

Option 2: From the **Protected Applications** page

- On the Hub cluster, navigate to **Data Services → Disaster recovery → Protected Applications → Actions (⋮) → Failover**.

- **For discovered applications:**

Navigate to **All Clusters → Data Services → Disaster recovery → Protected applications tab**. Click the **Actions (⋮)** and select **Failover**.

3. Validate that all protected applications are running on the surviving cluster before moving to the next step.



NOTE

The PROGRESSION state for each managed application **DRPlacementControl** shows as **Cleaning Up**. For discovered applications the PROGRESSION will be **WaitOnUserToCleanUp**. This is to be expected if the **replacement cluster** is offline or down.

4. From the Hub cluster, delete the DRCluster for the **replacement cluster**.

```
$ oc delete drcluster <drcluster_name> --wait=false
```



NOTE

Use **--wait=false** since the DRCluster will not be deleted until a later step.

5. Remove disaster recovery for each protected application on the surviving cluster. Perform all the sub-steps on the hub cluster.

- a. For each managed application, edit the Placement and ensure that the surviving cluster is selected:

```
$ oc edit placement <placement_name> -n <namespace>
```

```
apiVersion: cluster.open-cluster-management.io/v1beta1
kind: Placement
metadata:
  annotations:
    cluster.open-cluster-management.io/experimental-scheduling-disable: "true"
[...]
predicates:
- requiredClusterSelector:
    claimSelector: {}
    labelSelector:
      matchExpressions:
      - key: name
        operator: In
        values:
        - cluster1 <-- Modify to be surviving cluster name
[...]

```



NOTE

For Subscription-based applications the associated Placement can be found in the same namespace on the hub cluster similar to the managed clusters. For ApplicationSets-based applications the associated Placement can be found in the **openshift-gitops** namespace on the hub cluster.

- b. For each managed application, edit the PlacementDecision and ensure that the **replacement cluster** is deleted.

```
$ oc edit placementdecision <placementdecision_name> -n <namespace> --
subresource=status
```

```
apiVersion: cluster.open-cluster-management.io/v1beta1
kind: PlacementDecision
metadata:
  [...]
status:
  decisions:
    - clusterName: cluster2 <-- Delete this line
      reason: RetainedForFailover <-- Delete this line
    - clusterName: cluster1
      reason: cluster1
```



NOTE

For Subscription-based applications the associated PlacementDecision can be found in the same namespace on the hub cluster similar to the managed clusters. For ApplicationSets-based applications the associated PlacementDecision can be found in the **openshift-gitops** namespace on the hub cluster

- c. Verify that the **s3Profile** is removed for the replacement cluster by running the following command on the surviving cluster for each protected application's VolumeReplicationGroup (VRG).



NOTE

For managed applications, the associated VRG is found in the application namespace. For Discovered applications the associated VRG is in the **openshift-dr-ops** namespace.

```
$ oc get vrg -n <vrg_namespace> -o jsonpath='{.items[0].spec.s3Profiles}' | jq
```

- d. Remove disaster recovery from the applications on the RHACM console:

- **For managed applications:**

You can remove disaster recovery using either of the following paths:

Option 1: From the **Applications** page

- On the Hub cluster, navigate to **All Clusters → Applications**.
- In the **Overview** tab, at the end of the protected application row from the action menu (⋮), select **Manage disaster recovery**.
- Click **Remove disaster recovery**.
- Click **Confirm remove**.

Option 2: From the **Protected Applications** page

- On the Hub cluster, navigate to **All Clusters → Data Services → Disaster recovery → Protected Applications**.
- At the end of the protected application row from the action menu (⋮), select **Manage disaster recovery**.
- Click **Remove disaster recovery**.
- Click **Confirm remove**.

- **For discovered applications:**

- Navigate to **All Clusters → Data Services → Disaster recovery → Protected applications tab**.
- At the end of the application row, click on the action menu (⋮) and choose **Remove disaster recovery**.
- Click **Remove** in the next prompt.

6. Repeat the process detailed in the last step and the sub-steps for every protected application on the surviving cluster. Removing disaster recovery for protected applications is now completed.
7. Remove all disaster recovery configurations for the **surviving cluster** and the **replacement cluster**:

- a. From the hub cluster, delete DRPolicies associated with the **surviving cluster** and **replacement cluster** slated for DR removal:

```
$ oc delete drppolicy {drppolicy_name1 drppolicy_name2 ...}
```

- b. From the hub cluster, delete the mirrorpeer associated with the **surviving cluster** and **replacement cluster** slated for DR removal:

```
$ oc delete mirrorpeer {mirrorpeer_name}
```

- c. From the hub cluster, patch the mirrorpeer to remove the finalizer:

```
$ oc patch mirrorpeer {mirrorpeer_name} -p '{"metadata":{"finalizers":null}}' --type=merge
```

- d. From the **surviving cluster**, disable mirroring on the **cephblockpoolradosnamespace**:

```
$ oc patch cephblockpoolradosnamespace <cephblockpoolradosnamespace_name> --type merge -p '{"spec":{"mirroring": null}}' -n openshift-storage
```

- e. From the **surviving cluster**, disable mirroring on the **cephblockpool**:

```
$ oc patch cephblockpool <cephblockpool_name> --type json -p '[{"op": "replace", "path": "/spec/mirroring", "value": {}}]' -n openshift-storage
```

8. Uninstall **Submariner** for only the **replacement cluster** (failed cluster) using the RHACM console.


- a. Navigate to **Infrastructure → Clusters → Clustersets → Submariner add-ons** view and uninstall Submariner for only the replacement cluster.



NOTE

The uninstall process of Submariner for the replacement cluster (failed cluster) will stay GREEN and not complete until the replacement cluster has been detached from the RHACM console.

9. Navigate back to **Clusters view** and detach replacement cluster.
10. Create new OpenShift cluster (recovery cluster) and import into **Infrastructure → Clusters view**.
11. Add the new recovery cluster to the **Clusterset** used by Submariner.
12. Install **Submariner add-ons** only for the new **recovery cluster**.



NOTE

If GlobalNet is used for the surviving cluster make sure to enable GlobalNet for the recovery cluster as well.

13. Install **OpenShift Data Foundation** on the **recovery cluster**. The OpenShift Data Foundation version should be OpenShift Data Foundation 4.17 (or greater) and the same version of ODF as on the surviving cluster.

For managed applications using ApplicationSets, the OpenShift GitOps operator must be installed. For instructions, see [Installing Red Hat OpenShift GitOps Operator in web console](#).

For discovered applications, the OADP operator and the Data Protection Application (DPA) must be installed. For instructions, see [Prerequisites for disaster recovery protection of discovered applications](#).



NOTE

Make sure to follow the optional instructions in the documentation to modify the OpenShift Data Foundation storage cluster on the recovery cluster if GlobalNet has been enabled when installing Submariner.

14. Using the RHACM console, navigate to **Data Services → Disaster recovery → Policies** tab.
 - a. Select **Create DRPolicy** and name your policy.
 - b. Select the **recovery cluster** and the **surviving cluster**.
 - c. Create the policy. For instructions see chapter on [Creating Disaster Recovery Policy on Hub cluster](#).

Proceed to the next step only after the status of DRPolicy changes to **Validated**.

15. Apply the DRPolicy to the applications on the surviving cluster that were originally protected before the **replacement cluster** failed.



NOTE

For discovered applications, make sure to create application namespace on the **recovery cluster** before assigning the policy.

16. Ensure that there is a **poolID** mapping between the blockpools of the **replacement cluster** and **surviving cluster**, as well as between the **failed cluster** and **recovery cluster**.
 - a. Open the **clientProfileMappingBeforeRecovery.yaml** created in step 1 for editing:

```
$ vi clientProfileMappingBeforeRecovery.yaml
```

- b. In the **clientProfileMappingBeforeRecovery.yaml** delete the lines and replace the values notes as follows:

```
apiVersion: csi.ceph.io/v1
kind: ClientProfileMapping
metadata:
  creationTimestamp: "2025-11-13T07:43:45Z" <-- Remove this line
  generation: 7 <-- Remove this line
  name: ocs-storagecluster-surviving <-- name must be unique
  namespace: openshift-storage
  ownerReferences: <-- Remove this line
- apiVersion: ocs.openshift.io/v1alpha1 <-- Remove this line
  blockOwnerDeletion: true <-- Remove this line
  controller: true <-- Remove this line
  kind: StorageClient <-- Remove this line
  name: ocs-storagecluster <-- Remove this line
```

```

uid: c9b227a4-88d7-4667-82c8-f54286f498a9 <-- Remove this line
resourceVersion: "1392181" <-- Remove this line
uid: ba0676ab-d1f0-4b13-9193-db4b732280d3 <-- Remove this line
spec:
  mappings:
  - blockPoolIdMapping:
    - "3" <-- Sample blockPoolId, your values may be different
    - "3" <-- Sample blockPoolId, your values may be different
    - "16" <-- Sample blockPoolId, your values may be different
    - "14" <-- Sample blockPoolId, your values may be different
  localClientProfile: openshift-storage
  remoteClientProfile: openshift-storage

```

- c. Save the **clientProfileMappingBeforeRecovery.yaml** file to **clientProfileMappingBeforeRecovery-surviving.yaml**.
- d. Create the **clientProfileMappingBeforeRecovery-surviving.yaml** resource on the **surviving cluster**:

```
$ oc create -f clientProfileMappingBeforeRecovery-surviving.yaml -n openshift-storage
```

- e. On the **recovery cluster**, get the **blockpoolIdMapping** between the **recovery cluster** and the **surviving cluster**:

```
$ oc get clientprofilemapping ocs-storagecluster -n openshift-storage -o yaml | yq
'.spec.mappings[].blockPoolIdMapping'
```

Example output for two **cephblockpools**. The default is 1 **cephblockpool**. Your **blockPoolIds** may be different:

```

- "3" <-- Note this value for a future step
- "3"
- "13" <-- Note this value for a future step
- "14"

```

- f. Copy the contents of the **clientProfileMappingBeforeRecovery-surviving.yaml** into a new yaml:

```
$ cp clientProfileMappingBeforeRecovery-surviving.yaml
clientProfileMappingAfterRecovery.yaml
```

- g. Edit the **clientProfileMappingAfterRecovery.yaml** with the correct **blockPoolIds**:

```
$ vi clientProfileMappingAfterRecovery.yaml
```

The following as an example yaml with the required changes. Replace the **blockPoolIds** with the values you noted earlier:

```

apiVersion: csi.ceph.io/v1
kind: ClientProfileMapping
metadata:
  name: ocs-storagecluster-recovery <-- name must be unique
  namespace: openshift-storage

```

```
spec:
  mappings:
    - blockPoolIdMapping:
      - - "3" <-- First index from the ClientProfileMapping on recovery cluster
        - "3"
      - - "13" <-- Second index from the ClientProfileMapping on recovery cluster
        - "14"
    localClientProfile: openshift-storage
    remoteClientProfile: openshift-storage
```

- h. Create the **clientProfileMappingAfterRecovery.yaml** resource on the **recovery cluster**:

```
$ oc create -f clientProfileMappingAfterRecovery.yaml -n openshift-storage
```

17. Relocate the newly protected applications on the **surviving cluster** back to the new **recovery cluster** on the RHCAM console.

- **For managed applications:**

You can relocate protected applications using either of the following paths:

Option 1: From the **Applications** page

- On the Hub cluster, navigate to **All Clusters** → **Applications**.
- In the **Overview** tab, at the end of the application row from the action menu (⋮), select **Manage disaster recovery**.
- Click **Relocate**.

Option 2: From the **Protected Applications** page

- On the Hub cluster, navigate to **Data Services** → **Disaster recovery** → **Protected Applications** → **Actions** (⋮) → **Relocate**.

- **For discovered applications:**

- Navigate to **All Clusters** → **Data Services** → **Disaster recovery** → **Protected applications tab**.
- At the end of the application row, click on the action menu (⋮) and choose **Relocate**.

4.18. VIEWING RECOVERY POINT OBJECTIVE VALUES FOR DISASTER RECOVERY ENABLED APPLICATIONS

Recovery Point Objective (RPO) value is the most recent sync time of persistent data from the cluster where the application is currently active to its peer. This sync time helps determine duration of data lost during failover.



NOTE

This RPO value is applicable only for Regional-DR during failover. Relocation ensures there is no data loss during the operation, as all peer clusters are available.

You can view the Recovery Point Objective (RPO) value of all the protected volumes for their workload on the Hub cluster.

Procedure

1. On the Hub cluster, navigate to **Applications → Overview** tab.
2. In the **Data policy** column, click the **policy** link for the application you applied the policy to. A **Data Policies** modal page appears with the number of disaster recovery policies applied to each application along with failover and relocation status.
3. On the **Data Policies** modal page, click the **View more details** link.

A detailed **Data Policies** modal page is displayed that shows the policy names and the ongoing activities (Last sync, Activity status) associated with the policy that is applied to the application.

The **Last sync time** reported in the modal page, represents the most recent sync time of all volumes that are DR protected for the application.

4.19. HUB RECOVERY USING RED HAT ADVANCED CLUSTER MANAGEMENT

When your setup has active and passive Red Hat Advanced Cluster Management for Kubernetes (RHACM) hub clusters, and in case where the active hub is down, you can use the passive hub to failover or relocate the disaster recovery protected workloads.

4.19.1. Configuring passive hub cluster

To perform hub recovery in case the active hub is down or unreachable, follow the procedure in this section to configure the passive hub cluster and then failover or relocate the disaster recovery protected workloads.

Procedure

1. Ensure that RHACM operator and **MultiClusterHub** is installed on the passive hub cluster. See [RHACM installation guide](#) for instructions.
After the operator is successfully installed, the web console automatically reloads to apply the changes. During this process, a temporary error message might appear on the page and this is expected and disappears after the refresh completes.
2. Before hub recovery, configure backup and restore. See [Backup and restore](#) topics of *RHACM Business continuity* guide.
3. Install the multicluster orchestrator (MCO) operator along with Red Hat OpenShift GitOps operator on the passive RHACM hub prior to the restore. For instructions to restore your RHACM hub, see [Installing OpenShift Data Foundation Multicluster Orchestrator operator](#).
4. Ensure that **.spec.cleanupBeforeRestore** is set to **None** for the **Restore.cluster.open-cluster-management.io** resource. For details, see [Restoring passive resources while checking for backups](#) chapter of RHACM documentation.
5. If SSL access across clusters was configured manually during setup, then re-configure SSL access across clusters. For instructions, see [Configuring SSL access across clusters](#) chapter.
6. On the passive hub, add a label to **openshift-operators** namespace to enable basic monitoring of **VolumeSynchronizationDelay** alert using this command. For alert details, see [Disaster recovery alerts](#) chapter.

```
$ oc label namespace openshift-operators openshift.io/cluster-monitoring='true'
```


4.19.2. Switching to passive hub cluster

Use this procedure when the active hub is down or unreachable.


Procedure

1. During the restore procedure, to avoid eviction of resources when **ManifestWorks** are not regenerated correctly, you can enlarge the **AppliedManifestWork eviction** grace period. On the passive hub cluster, check for existing global **KlusterletConfig**.
 - If global **KlusterletConfig** exists then edit and set the value for **appliedManifestWorkEvictionGracePeriod** parameter to a larger value. For example, 24 hours or more.
 - If global **KlusterletConfig** does not exist, then create the **Klusterletconfig** using the following yaml:

```
apiVersion: config.open-cluster-management.io/v1alpha1
kind: KlusterletConfig
metadata:
  name: global
spec:
  appliedManifestWorkEvictionGracePeriod: "24h"
```

The configuration will be propagated to all the managed clusters automatically.

2. Restore the backups on the passive hub cluster. For information, see [Restoring a hub cluster from backup](#).



IMPORTANT

Recovering a failed hub to its passive instance will only restore applications and their DR protected state to its last scheduled backup. Any application that was DR protected after the last scheduled backup would need to be protected again on the new hub.

3. Verify that the restore is complete.

```
$ oc -n <restore-namespace> wait restore <restore-name> --
for=jsonpath='{.status.phase}'=Finished --timeout=120s
```

4. Verify that the Primary and Secondary managed clusters are successfully imported into the RHACM console and they are accessible. If any of the managed clusters are down or unreachable then they will not be successfully imported.
Wait until DRPolicy validation succeeds before performing any DR operation.

If the previously down managed cluster comes back online in the future after the hub recovery process, ensure that it is successfully imported and tied to the new hub. If the cluster doesn't import and tie to the new hub, delete and recreate the **restore-acm** resource:

```
$ oc get restore restore-acm -o yaml -n open-cluster-management-backup
```

Then verify that the cluster is properly connected to the new hub.

**NOTE**

Submariner is automatically installed once the managed clusters are imported on the passive hub.

5. Verify that the **DRPolicy** is created successfully. Run this command on the **Hub cluster** for each of the DRPolicy resources created, where *<drpolicy_name>* is replaced with a unique name.

```
$ oc get drpolicy <drpolicy_name> -o jsonpath='{.status.conditions[].reason}'
```

Example output:

```
Succeeded
```

6. Refresh the RHACM console to make the DR monitoring dashboard tab accessible if it was enabled on the Active hub cluster.
7. Verify the DRPC output using the following command on the new hub cluster:

```
$ oc get drpc -A -o wide
```

If **PROGRESSION** shows a status of **PAUSED**, administrative intervention is required to unpause it. **PROGRESSION** enters **PAUSED** state under the following conditions:

- **Cluster Query Failure:** None of the clusters were successfully queried during the DRPC reconciliation. This situation can occur during hub recovery.
- **Action Mismatch:** The DRPC action differs from the queried VRG action.
- **Cluster Mismatch:** The DRPC action and the VRG action are the same, but the Primary VRG is found in a different cluster than the one expected by the DRPC.

**IMPORTANT**

If you cannot diagnose and resolve the cause of the pause, contact [Red Hat Customer Support](#).

If **PROGRESSION** is in either **Completed** or **Cleaning up**, it is safe to proceed.

8. Edit the global **KlusterletConfig** on the new hub and remove the parameter **appliedManifestWorkEvictionGracePeriod** and its value.
9. Depending on whether the active hub cluster, or both the active hub cluster along with the primary managed cluster had been down, follow the next steps based on your scenario:
 - a. If only the active hub cluster had been down, and if the managed clusters are still accessible, no further action is required.
 - b. If the primary managed cluster had been down, along with the active hub cluster, you need to fail over the workloads from the primary managed cluster to the secondary managed cluster.
For failover instructions, based on your workload type, see [Subscription-based applications](#) or [ApplicationSet-based applications](#).

10. Verify that the failover is successful. If the Primary managed cluster is also down, then the PROGRESSION status for the workload would be in the **Cleaning Up** phase until the down Primary managed cluster is back online and successfully imported into the RHACM console. On the passive hub cluster, run the following command to check the PROGRESSION status.

```
$ oc get drpc -o wide -A
```

4.20. ENABLING GRANULAR DISASTER RECOVERY FOR INDIVIDUAL OR GROUPS OF VIRTUAL MACHINES IN A NAMESPACE

You can enable granular disaster recovery (DR) for individual virtual machines (VMs) or VM groups within the same namespace, allowing independent failover and relocation actions.

Key highlights

- **Granular VM Control** Each discovered and ACM-managed VM can have its own DR policy, enabling independent DR operations without impacting other VMs in the namespace.
- **Namespace Partitioning** Supports multiple DRPCs within a namespace, avoiding the limitations of namespace-level DR protection.
- **Scenarios Supported:**
 - Independent DR for discovered and ACM-managed VMs.
 - Namespace-level DR remains for non-VM applications.
- **Upgrade Considerations:**
 - Existing DRPCs are treated as namespace-level.
 - Users must disable DR before switching to VM-specific protection.
- **Manage disaster recovery of virtual machines from** [Virtual Machines page](#)



NOTE

Existing UI for namespace-level DR remains unchanged.

IMPORTANT

During failover of a CNV discovered application, if the primary managed cluster is down and other CNV applications are running in different namespaces on either of the managed clusters, there have been cases where the VM fails to start on the failover cluster. This issue is due to MAC address conflicts.

To prevent this issue, it is recommended to configure disjoint MAC address ranges for each cluster from the Day 1. For example:

```
$ oc patch configmap kubemacpool-mac-range-config -n openshift-cnv --type merge -
p '{"data":
{"RANGE_START":"02:00:00:00:00:00","RANGE_END":"02:00:FF:ff:FF:FF"}}'
```

```
$ oc patch configmap kubemacpool-mac-range-config -n openshift-cnv --type merge -
p '{"data":
{"RANGE_START":"02:01:00:00:00:00","RANGE_END":"02:01:FF:FF:FF:FF"}}'
```

Then restart the **kubemacpool-mac-controller-manager**:

```
$ oc get pods -n openshift-cnv|grep kubemacpool-mac-controller-manager
kubemacpool-mac-controller-manager-544ff96b4f-dgb7x 2/2 Running 3
30h
```

4.21. ENABLING GRANULAR DISASTER RECOVERY FOR VIRTUAL MACHINES USING THE CLI

This section describes how to configure granular, CLI driven Regional Disaster Recovery (Regional-DR) for individual or groups of KubeVirt virtual machines (VMs) using the built-in Ramen **vm-recipe**. This approach allows cluster administrators to protect specific VM workloads within a shared namespace without using the OpenShift Web Console or defining custom Recipe CRs.

Prerequisites

- OpenShift Data Foundation DR is installed and configured on the hub cluster.
- A **DRPolicy** defining the peer managed clusters is created.
- The **kubectrl** and **oc** CLI tools are installed and configured with the hub cluster context.

IMPORTANT

If you are protecting virtual machines (VMs) with Disaster Recovery, you must ensure that the target VMs have successfully transitioned to the **Running** state before triggering any DR protection workflows, failover, or relocate operations.

Do not enroll or initiate DR operations for a VM that is in any of the following intermediate or pending phases:

- **Provisioning**
- **Scheduling**
- **CreatingDataVolume**
- **Starting**

IMPORTANT

If the target VMs were migrated using the Migration Toolkit for Virtualization (MTV), you must archive the associated migration plans before enrolling the VMs in a DR policy. Active or unarchived migration plans leave residual validation and import pods that retain PVC mounts. This will cause DR operations to stall in a **WaitingForCleanup** or **Terminating** state.

For archiving instructions, see the [Migration Toolkit for Virtualization Documentation](#).

Procedure

1. Determine the Ramen operations namespace configured in your environment by inspecting the operator configuration map on the hub cluster:

```
$ oc get cm -n openshift-operators ramen-hub-operator-config -o yaml | grep
"ramenOpsNamespace"
```

Example output:

```
ramenOpsNamespace: openshift-dr-ops
```

Use this namespace value (e.g., **openshift-dr-ops**) in all subsequent steps where it is referenced.

2. Label Your VM resources

The built-in **vm-recipe** utilizes label selectors to identify resources. You must apply an identical, unique tracking label to both the **VirtualMachine** resource and its associated **PersistentVolumeClaims** (PVCs).

.Label the VirtualMachine resource:

+

```
$ oc label vm -n <vm-namespace> \
  ramendr.openshift.io/k8s-resource-selector= <unique-selector-value>
```

.Label all PVCs attached to the VM:

+

```
$ oc label pvc -n <vm-namespace> \
  ramendr.openshift.io/k8s-resource-selector=<unique-selector-value>
```

**NOTE**

Use the same selector value for all resources belonging to the same VM application. For protecting multiple VMs as a group, use the same label value on all VMs and their PVCs.

1. Create the Placement Resource

A RHACM **Placement** resource must exist on the hub cluster **before** configuring the DR scheduling controls.

a. Create a file named **vm-placement.yaml**:

```
apiVersion: cluster.open-cluster-management.io/v1beta1
kind: Placement
metadata:
  name: vm-placement
  namespace: <ramenOpsNamespace> # Use the value identified in Step 1
  annotations:
    cluster.open-cluster-management.io/experimental-scheduling-disable: "true"
spec:
  numberOfClusters: 1
  clusterSets:
    - <your-clusterset-name>
```

b. Apply the configuration to the hub cluster:

```
$ oc apply -f vm-placement.yaml
```

2. Create the DRPlacementControl with vm-recipe reference

a. Create a file named **vm-drpc.yaml** with the embedded vm-recipe reference:

```
apiVersion: ramendr.openshift.io/v1alpha1
kind: DRPlacementControl
metadata:
  name: vm-drpc
  namespace: <ramenOpsNamespace> # Use the value identified in Step 1
spec:
  preferredCluster: <primary-cluster-name>
  drPolicyRef:
    name: <your-dr-policy-name>
  placementRef:
    kind: Placement
    name: vm-placement
  protectedNamespaces:
    - <vm-namespace>
  pvcSelector:
    matchLabels:
```

```

    ramendr.openshift.io/k8s-resource-selector: <unique-selector-value>
  kubeObjectProtection:
    recipeRef:
      namespace: <ramenOpsNamespace> # Use the value identified in Step 1
      name: vm-recipe # Explicitly references the built-in Ramen VM recipe
    recipeParameters:
      K8S_RESOURCE_SELECTOR:
        - <unique-selector-value>
      PVC_RESOURCE_SELECTOR:
        - <unique-selector-value>

```

- b. Apply the DRPC configuration:

```
$ oc apply -f vm-drpc.yaml
```

3. Verify protection status.

- a. Check the **DRPlacementControl** status on the hub cluster:

```
$ oc get drpc vm-drpc -n <ramen-ops-namespace> -o yaml
```

Ensure that the output reflects the following parameters: * **status.phase** reports as **Deployed**. * **status.conditions** contains a type of **Protected** with a status of **True**. * **status.lastGroupSyncTime** displays a valid, recent timestamp indicating continuous delta replication.

- b. Verify that the VR) is successfully initialized on the primary managed cluster:

```
$ oc get vrg -n <vm-namespace> --context <managed-cluster-context>
```

4. Test DR operations

Once protected, you can execute disaster recovery orchestration directly via patch commands from your terminal.

- a. Execute a relocate (planned migration).

To migrate your virtual machine to a targeted peer cluster during planned maintenance windows, update the action parameter:

```
$ oc patch drpc vm-drpc -n $RAMEN_OPS_NS --type merge \
-p '{"spec":{"action":"Relocate","failoverCluster":"<target-cluster-name>"}}'
```

- b. Execute a failover (unplanned outage)

In the event of a primary cluster failure, force immediate disaster recovery procedures on the surviving destination cluster:

```
$ oc patch drpc vm-drpc -n $RAMEN_OPS_NS --type merge \
-p '{"spec":{"action":"Failover","failoverCluster":"<target-cluster-name>"}}'
```

- c. Monitor operation progress

Track the synchronization, teardown, and spin-up transitions of your VM workloads in real-time:

```
$ oc get drpc vm-drpc -n <ramen-ops-namespace> -w
```

4.21.1. Key points about granular virtual machine protection

Review the following key structural behaviors to plan and scale your command-line virtual machine configurations effectively:

Individual virtual machine protection

You can label and protect a single virtual machine explicitly by isolating it with its own unique selector value.

Group virtual machine protection

You can protect multiple independent virtual machines collectively by applying the exact same label value across all target VMs and their respective PVCs within that group.

Namespace isolation

Multiple DRPCs can protect different VM groups in the same namespace using different labels.

No recipe CR creation needed

The **vm-recipe** is embedded directly in the Ramen operator code, eliminating the need to create or manage external custom Recipe CRs.

Mandatory operations namespace alignment

Every DRPC and recipe reference evaluation must use the namespace from Step 1.

This approach enables fine-grained control over which VMs are protected, allowing you to create separate DR policies for different VM groups within the same namespace.

CHAPTER 5. DISASTER RECOVERY WITH STRETCH CLUSTER FOR OPENSHIFT DATA FOUNDATION

Red Hat OpenShift Data Foundation deployment can be stretched between two different geographical locations to provide the storage infrastructure with disaster recovery capabilities. When faced with a disaster, such as one of the two locations is partially or totally not available, OpenShift Data Foundation deployed on the OpenShift Container Platform deployment must be able to survive. This solution is available only for metropolitan spanned data centers with specific latency requirements between the servers of the infrastructure.




NOTE

The stretch cluster solution is designed for deployments where latencies do not exceed 10 ms maximum round-trip time (RTT) between the zones containing data volumes. For Arbiter nodes follow the latency requirements specified for etcd, see [Guidance for Red Hat OpenShift Container Platform Clusters - Deployments Spanning Multiple Sites\(Data Centers/Regions\)](#). Contact [Red Hat Customer Support](#) if you are planning to deploy with higher latencies.

The following diagram shows the simplest deployment for a stretched cluster:

OpenShift nodes and OpenShift Data Foundation daemons



In the diagram the OpenShift Data Foundation monitor pod deployed in the Arbiter zone has a built-in tolerance for the master nodes. The diagram shows the master nodes in each Data Zone which are required for a highly available OpenShift Container Platform control plane. Also, it is important that the OpenShift Container Platform nodes in one of the zones have network connectivity with the OpenShift Container Platform nodes in the other two zones.



IMPORTANT

You can now easily set up disaster recovery with stretch cluster for workloads based on OpenShift virtualization technology using OpenShift Data Foundation. For more information, see [OpenShift Virtualization in OpenShift Container Platform](#) guide.

5.1. REQUIREMENTS FOR ENABLING STRETCH CLUSTER

- Ensure you have addressed OpenShift Container Platform requirements for deployments spanning multiple sites. For more information, see [knowledgebase article on cluster deployments spanning multiple sites](#).
- Ensure that you have at least three OpenShift Container Platform master nodes in three different zones. One master node in each of the three zones.
- Ensure that you have at least four OpenShift Container Platform worker nodes evenly distributed across the two Data Zones.
- For stretch clusters on bare metall, use the SSD drive as the root drive for OpenShift Container Platform master nodes.
- Ensure that each node is pre-labeled with its zone label. For more information, see the [Applying topology zone labels to OpenShift Container Platform node](#) section.
- The stretch cluster solution is designed for deployments where latencies do not exceed 10 ms between zones. Contact [Red Hat Customer Support](#) if you are planning to deploy with higher latencies.



NOTE

Flexible scaling and Arbiter both cannot be enabled at the same time as they have conflicting scaling logic. With Flexible scaling, you can add one node at a time to your OpenShift Data Foundation cluster. Whereas in an Arbiter cluster, you need to add at least one node in each of the two data zones.

5.2. APPLYING TOPOLOGY ZONE LABELS TO OPENSIFT CONTAINER PLATFORM NODES

During a site outage, the zone that has the arbiter function makes use of the arbiter label. These labels are arbitrary and must be unique for the three locations.

For example, you can label the nodes as follows:

```
topology.kubernetes.io/zone=arbiter for Master0
```

```
topology.kubernetes.io/zone=datacenter1 for Master1, Worker1, Worker2
```

```
topology.kubernetes.io/zone=datacenter2 for Master2, Worker3, Worker4
```

- To apply the labels to the node:

```
$ oc label node <NODENAME> topology.kubernetes.io/zone=<LABEL>
```

<NODENAME>

Is the name of the node

<LABEL>

Is the topology zone label

- To validate the labels using the example labels for the three zones:

```
$ oc get nodes -l topology.kubernetes.io/zone=<LABEL> -o name
```

<LABEL>

Is the topology zone label

Alternatively, you can run a single command to see all the nodes with its zone.

```
$ oc get nodes -L topology.kubernetes.io/zone
```

The stretch cluster topology zone labels are now applied to the appropriate OpenShift Container Platform nodes to define the three locations.

Next step

- [Install the local storage operator from the OpenShift Container Platform web console](#) .

5.3. INSTALLING LOCAL STORAGE OPERATOR

Install the Local Storage Operator from the Software Catalog before creating Red Hat OpenShift Data Foundation clusters on local storage devices.

Procedure

1. Log in to the OpenShift Web Console.
2. Click **Ecosystem → Software Catalog**
3. Type **local storage** in the **Filter by keyword** box to find the **Local Storage Operator** from the list of operators, and click on it.
4. Set the following options on the **Install Operator** page:
 - a. Update channel as **stable**.
 - b. Installation mode as **A specific namespace on the cluster**
 - c. Installed Namespace as **Operator recommended namespace openshift-local-storage**.
 - d. Update approval as **Automatic**.
5. Click **Install**.

Verification steps

- Verify that the Local Storage Operator shows a green tick indicating successful installation.

5.4. INSTALLING RED HAT OPENSIFT DATA FOUNDATION OPERATOR

You can install Red Hat OpenShift Data Foundation Operator using the Red Hat OpenShift Container Platform Software Catalog.

Prerequisites

- Access to an OpenShift Container Platform cluster using an account with cluster-admin and Operator installation permissions.
- You must have at least four worker nodes evenly distributed across two data centers in the Red Hat OpenShift Container Platform cluster.
- For additional resource requirements, see [Planning your deployment](#).



IMPORTANT

- When you need to override the cluster-wide default node selector for OpenShift Data Foundation, you can use the following command in command-line interface to specify a blank node selector for the **openshift-storage** namespace (create openshift-storage namespace in this case):


```
$ oc annotate namespace openshift-storage openshift.io/node-selector=
```
- Taint a node as **infra** to ensure only Red Hat OpenShift Data Foundation resources are scheduled on that node. This helps you save on subscription costs. For more information, see [How to use dedicated worker nodes for Red Hat OpenShift Data Foundation](#) chapter in the *Managing and Allocating Storage Resources* guide.

Procedure

1. Log in to the OpenShift Web Console.
2. Click **Ecosystem → Software Catalog**
3. Scroll or type **OpenShift Data Foundation** into the **Filter by keyword** box to search for the **OpenShift Data Foundation** Operator.
4. Click **Install**.
5. Set the following options on the **Install Operator** page:
 - a. Update Channel as **stable-4.21**.
 - b. Installation Mode as **A specific namespace on the cluster**
 - c. Installed Namespace as **Operator recommended namespace openshift-storage**. If Namespace **openshift-storage** does not exist, it is created during the operator installation.
 - d. Select **Approval Strategy** as **Automatic** or **Manual**.
 If you select **Automatic** updates, then the Operator Lifecycle Manager (OLM) automatically upgrades the running instance of your Operator without any intervention.

 If you selected **Manual** updates, then the OLM creates an update request. As a cluster administrator, you must then manually approve that update request to update the Operator to a newer version.
6. Ensure that the **Enable** option is selected for the **Console plugin**.
7. Click **Install**.

Verification steps

- After the operator is successfully installed, the web console automatically reloads to apply the changes. During this process, a temporary error message might appear on the page and this is expected and disappears after the refresh completes.
- In the Web Console:
 - Navigate to Installed Operators and verify that the **OpenShift Data Foundation** Operator shows a green tick indicating successful installation.
 - Navigate to **Storage** and verify if the **Data Foundation** dashboard is available.

Next steps

- [Create an OpenShift Data Foundation cluster](#).


5.5. CREATING OPENSIFT DATA FOUNDATION CLUSTER

Prerequisites

- Ensure that you have met all the requirements in [Requirements for enabling stretch cluster](#) section.

Procedure

1. In the OpenShift Web Console, click **Storage** → **Data Foundation** → **Storage Systems** → **Create StorageSystem**.
2. In the Backing storage page, select the **Create a new StorageClass using the local storage devices** option.
3. Click **Next**.



IMPORTANT

You are prompted to install the Local Storage Operator if it is not already installed. Click **Install**, and follow the procedure as described in [Installing Local Storage Operator](#).

4. In the **Create local volume set** page, provide the following information:
 - a. Enter a name for the **LocalVolumeSet** and the **StorageClass**.
By default, the local volume set name appears for the storage class name. You can change the name.
 - b. Choose one of the following:
 - **Disks on all nodes**
Uses the available disks that match the selected filters on all the nodes.
 - **Disks on selected nodes**
Uses the available disks that match the selected filters only on selected nodes.



IMPORTANT

If the nodes selected do not match the OpenShift Data Foundation cluster requirement of an aggregated 30 CPUs and 72 GiB of RAM, a minimal cluster is deployed.

For minimum starting node requirements, see the [Resource requirements](#) section in the *Planning* guide.

- c. Select **SSD** or **NVMe** to build a supported configuration. You can select **HDDs** for unsupported test installations.

- d. Expand the **Advanced** section and set the following options:

Volume Mode	Block is selected by default.
Device Type	Select one or more device types from the dropdown list.
Disk Size	Set a minimum size of 100GB for the device and maximum available size of the device that needs to be included.
Maximum Disks Limit	This indicates the maximum number of PVs that can be created on a node. If this field is left empty, then PVs are created for all the available disks on the matching nodes.

- e. Click **Next**.

A pop-up to confirm the creation of LocalVolumeSet is displayed.

- f. Click **Yes** to continue.

5. In the **Capacity and nodes** page, configure the following:

- a. **Available raw capacity** is populated with the capacity value based on all the attached disks associated with the storage class. This takes some time to show up.
The **Selected nodes** list shows the nodes based on the storage class.

- b. Select **Enable arbiter** checkbox if you want to use the stretch clusters. This option is available only when all the prerequisites for arbiter are fulfilled and the selected nodes are populated. For more information, see Arbiter stretch cluster requirements in [Requirements for enabling stretch cluster](#).
Select the **arbiter zone** from the dropdown list.

6. Choose a performance profile for **Configure performance**.

You can also configure the performance profile after the deployment using the **Configure performance** option from the options menu of the **StorageSystems** tab.



IMPORTANT

Before selecting a resource profile, make sure to check the current availability of resources within the cluster. Opting for a higher resource profile in a cluster with insufficient resources might lead to installation failures. For more information about resource requirements, see [Resource requirement for performance profiles](#).

- a. Click **Next**.
7. Optional: In the Security and network page, configure the following based on your requirement:
 - a. To enable encryption, select **Enable data encryption for block and file storage**
 - b. Select one of the following **Encryption level**:
 - **Cluster-wide encryption** to encrypt the entire cluster (block and file).
 - **StorageClass encryption** to create encrypted persistent volume (block only) using encryption enabled storage class.
 - c. Optional: Select the **Connect to an external key management service** checkbox. This is optional for cluster-wide encryption.
 - i. From the **Key Management Service Provider** drop-down list, either select **Vault** or **Thales CipherTrust Manager (using KMIP)**. If you selected **Vault**, go to the next step. If you selected **Thales CipherTrust Manager (using KMIP)**, go to step iii.
 - ii. Select an **Authentication Method**.

Using Token authentication method

- Enter a unique **Connection Name**, host **Address** of the Vault server ('https://<hostname or ip>'), **Port** number and **Token**.
- Expand **Advanced Settings** to enter additional settings and certificate details based on your **Vault** configuration:
 - Enter the Key Value secret path in the **Backend Path** that is dedicated and unique to OpenShift Data Foundation.
 - Optional: Enter **TLS Server Name** and **Vault Enterprise Namespace**
 - Upload the respective PEM encoded certificate file to provide the **CA Certificate**, **Client Certificate** and **Client Private Key**.
 - Click **Save** and skip to step iv.

Using Kubernetes authentication method

- Enter a unique Vault **Connection Name**, host **Address** of the Vault server ('https://<hostname or ip>'), **Port** number and **Role** name.
- Expand **Advanced Settings** to enter additional settings and certificate details based on your **Vault** configuration:
 - Enter the Key Value secret path in the **Backend Path** that is dedicated and unique to OpenShift Data Foundation.

- Optional: Enter **TLS Server Name** and **Authentication Path** if applicable.
 - Upload the respective PEM encoded certificate file to provide the **CA Certificate**, **Client Certificate** and **Client Private Key**.
 - Click **Save** and skip to step iv.
- iii. To use **Thales CipherTrust Manager (using KMIP)** as the KMS provider, follow the steps below:
- A. Enter a unique **Connection Name** for the Key Management service within the project.
 - B. In the **Address** and **Port** sections, enter the IP of Thales CipherTrust Manager and the port where the KMIP interface is enabled. For example:
 - **Address:** 123.34.3.2
 - **Port:** 5696
 - C. Upload the **Client Certificate**, **CA certificate**, and **Client Private Key**.
 - D. If StorageClass encryption is enabled, enter the Unique Identifier to be used for encryption and decryption generated above.
 - E. The **TLS Server** field is optional and used when there is no DNS entry for the KMIP endpoint. For example, **kmip_all_<port>.ciphertrustmanager.local**.
- d. **Network** is set to Default (OVN) if you are using a single network. You can switch to Custom (Multus) if you are using multiple network interfaces and then choose any one of the following:
- i. Select a Public Network Interface from the dropdown.
 - ii. Select a Cluster Network Interface from the dropdown.



NOTE

If you are using only one additional network interface, select the single **NetworkAttachmentDefinition**, that is, **ocs-public-cluster** for the Public Network Interface, and leave the Cluster Network Interface blank.

- a. Click **Next**.
8. In the **Review and create** page, review the configuration details.
To modify any configuration settings, click **Back** to go back to the previous configuration page.
9. Click **Create StorageSystem**.

Verification steps

- To verify the final Status of the installed storage cluster:
 - a. On the OpenShift Web Console, navigate to **Storage → Data Foundation → Storage System**.

- b. In the **Status** card of the **Overview** tab, click **Storage System** and then click the storage system link from the pop up that appears.
 - c. In the **Status** card of the **Block and File** tab, verify that the *Storage Cluster* has a green tick.
- For arbiter mode of deployment:
 - a. In the OpenShift Web Console, navigate to **Storage → Data Foundation → View Storage** in the **Storage cluster card**
 - b. From the actions menu, select **Edit Storage Cluster** to view the YAML.
 - c. Search for the **arbiter** key in the **spec** section and ensure **enable** is set to **true**.

```
spec:
  arbiter:
    enable: true
  [..]
  nodeTopologies:
    arbiterLocation: arbiter #arbiter zone
  storageDeviceSets:
  - config: {}
    count: 1
    [..]
    replica: 4
  status:
    conditions:
    [..]
    failureDomain: zone
```

- To verify that all the components for OpenShift Data Foundation are successfully installed, see [Verifying your OpenShift Data Foundation installation](#).

5.5.1. Creating NetworkFenceClass

NetworkFenceClass is required to prevent corruption of volume contents due to non-graceful node shutdowns.

Procedure

1. Create the **NetworkFenceClass** from the following YAML and verify that the **csiaddonsnode** objects are populated to the IP address and Ceph cluster ID.



NOTE

You need to create one or more **NetworkFenceClass** depending on the amount of storage clusters connected to the OpenShift Container Platform cluster.

```
apiVersion: csiaddons.openshift.io/v1alpha1
kind: NetworkFenceClass
metadata:
  name: odf-networkfenceclass
spec:
  provisioner: openshift-storage.rbd.csi.ceph.com
```

```
parameters:
  clusterID: openshift-storage
  csiaddons.openshift.io/networkfence-secret-name: rook-csi-rbd-node
  csiaddons.openshift.io/networkfence-secret-namespace: openshift-storage
```

- **provisioner**: specifies the name of storage provisioner.
 - **parameters**: specifies storage provider specific parameters.
 - **csiaddons.openshift.io/networkfence-secret-name**: specifies the name of the secret required for network fencing operation. This can be fetched from **parameters.csi.storage.k8s.io/node-stage-secret-name** in the **ocs-storagecluster-ceph-rbd storageclass**.
 - **csiaddons.openshift.io/networkfence-secret-namespace**: specifies the namespace in which the secret is located. This can be fetched from **parameters.csi.storage.k8s.io/node-stage-secret-namespace** in **ocs-storagecluster-ceph-rbd storageclass**.
2. Run the following command to return the **csiaddonsnode** objects. The node objects belongs to the daemonset pod (RBD) that will have the IP address that needs to be fenced.

```
$ oc get csiaddonsnode -n openshift-storage
```

5.6. VERIFYING OPENSIFT DATA FOUNDATION DEPLOYMENT

To verify that OpenShift Data Foundation is deployed correctly:

- [Verify the state of the pods.](#)
- [Verify that the OpenShift Data Foundation cluster is healthy.](#)
- [Verify that the Multicloud Object Gateway is healthy.](#)
- [Verify that the OpenShift Data Foundation specific storage classes exist.](#)

5.6.1. Verifying the state of the pods

Procedure

1. Click **Workloads** → **Pods** from the OpenShift Web Console.
2. Select **openshift-storage** from the **Project** drop-down list.



NOTE

If the **Show default projects** option is disabled, use the toggle button to list all the default projects.

For more information about the expected number of pods for each component and how it varies depending on the number of nodes, see [Table 5.1, “Pods corresponding to OpenShift Data Foundation cluster”](#).

- Click the **Running** and **Completed** tabs to verify that the following pods are in **Running** and **Completed** state:



NOTE

The available pods depend on the cluster configuration. When the cluster is deployed as a standalone Multicloud Object Gateway, the **rook-ceph-operator-*** pods are not available. Similarly, when the cluster is deployed without the Multicloud Object Gateway, **noobaa-*** pods are not available.

Table 5.1. Pods corresponding to OpenShift Data Foundation cluster

Component	Corresponding pods
OpenShift Data Foundation Operator	● ocs-operator-* (1 pod on any worker node) ● ocs-metrics-exporter-* (1 pod on any worker node) ● odf-operator-controller-manager-* (1 pod on any worker node) ● odf-console-* (1 pod on any worker node) ● csi-addons-controller-manager-* (1 pod on any worker node)
Rook-ceph Operator	rook-ceph-operator-* (1 pod on any worker node)
Multicloud Object Gateway	● noobaa-operator-* (1 pod on any worker node) ● noobaa-core-* (1 pod on any storage node) ● noobaa-db-pg-cluster-1 and noobaa-db-pg-cluster-2 (2 instances of MCG DB pod on any storage node) ● noobaa-endpoint-* (1 pod on any storage node) ● cnpg-controller-manager-* (1 pod on any storage node)
MON	rook-ceph-mon-* (5 pods are distributed across 3 zones, 2 per data-center zones and 1 in arbiter zone)

Component	Corresponding pods
MGR	rook-ceph-mgr-* (2 pods on any storage node)
MDS	rook-ceph-mds-ocs-storagecluster-cephfilesystem-* (2 pods are distributed across 2 data-center zones)
RGW	rook-ceph-rgw-ocs-storagecluster-cephobjectstore-* (2 pods are distributed across 2 data-center zones)
CSI	● cephfs ○ openshift-storage.cephfs.csi.ceph.com-ctrlplugin-* (2 pods distributed across storage nodes) ○ openshift-storage.cephfs.csi.ceph.com-nodeplugin-* (1 pod on each storage node) ● nfs ○ openshift-storage.nfs.csi.ceph.com-ctrlplugin-* (2 pods distributed across storage nodes) ○ openshift-storage.nfs.csi.ceph.com-nodeplugin-* (1 pod on each storage node) ● rbd ○ openshift-storage.rbd.csi.ceph.com-ctrlplugin-* (2 pods distributed across storage nodes) ○ openshift-storage.rbd.csi.ceph.com-nodeplugin-* (1 pod on each storage node)

Component	Corresponding pods
rook-ceph-crashcollector	rook-ceph-crashcollector-* (1 pod on each storage node and 1 pod in arbiter zone)
OSD	● rook-ceph-osd-* (1 pod for each device) ● rook-ceph-osd-prepare-* (1 pod for each device)

5.6.2. Verifying the OpenShift Data Foundation cluster is healthy

Procedure

1. In the OpenShift Web Console, click **Storage → Data Foundation**.
2. In the **Status** card of the **Overview** tab, click **Storage System** and then click the storage system link from the pop up that appears.
3. In the **Status** card of the **Block and File** tab, verify that the *Storage Cluster* has a green tick.
4. In the **Details** card, verify that the cluster information is displayed.


For more information on the health of the OpenShift Data Foundation cluster using the **Block and File** dashboard, see [Monitoring OpenShift Data Foundation](#).

5.6.3. Verifying the Multicloud Object Gateway is healthy

Procedure

1. In the OpenShift Web Console, click **Storage → Data Foundation**.
2. In the **Status** card of the **Overview** tab, click **Storage System** and then click the storage system link from the pop up that appears.
 - a. In the **Status card** of the **Object** tab, verify that both *Object Service* and *Data Resiliency* have a green tick.
 - b. In the **Details** card, verify that the MCG information is displayed.

For more information on the health of the OpenShift Data Foundation cluster using the object service dashboard, see [Monitoring OpenShift Data Foundation](#).



IMPORTANT

To avoid data loss, it is recommended to take a backup of NooBaa DB PVC regularly. If NooBaa DB fails and cannot be recovered, then you can revert to the latest backed-up version. For instructions on backing up your NooBaa DB, follow the steps in the knowledgebase article, [Perform a One-Time Backup of the Database for the Multicloud Object Gateway](#).

5.6.4. Verifying that the specific storage classes exist

Procedure

1. Click **Storage** → **Storage Classes** from the left pane of the OpenShift Web Console.
2. Verify that the following storage classes are created with the OpenShift Data Foundation cluster creation:
 - **ocs-storagecluster-ceph-rbd**
 - **ocs-storagecluster-cephfs**
 - **openshift-storage.noobaa.io**
 - **ocs-storagecluster-ceph-rgw**

5.7. INSTALL ZONE AWARE SAMPLE APPLICATION

Deploy a zone aware sample application to validate whether an OpenShift Data Foundation, stretch cluster setup is configured correctly.



IMPORTANT

With latency between the data zones, you can expect to see performance degradation compared to an OpenShift cluster with low latency between nodes and zones (for example, all nodes in the same location). The rate of or amount of performance degradation depends on the latency between the zones and on the application behavior using the storage (such as heavy write traffic). Ensure that you test the critical applications with stretch cluster configuration to ensure sufficient application performance for the required service levels.

A ReadWriteMany (RWX) Persistent Volume Claim (PVC) is created using the **ocs-storagecluster-cephfs** storage class. Multiple pods use the newly created RWX PVC at the same time. The application used is called File Uploader.

Demonstration on how an application is spread across topology zones so that it is still available in the event of a site outage:



NOTE

This demonstration is possible since this application shares the same RWX volume for storing files. It works for persistent data access as well because Red Hat OpenShift Data Foundation is configured as a stretched cluster with zone awareness and high availability.

1. Create a new project.

```
$ oc new-project my-shared-storage
```

2. Deploy the example PHP application called file-uploader.

```
$ oc new-app openshift/php:latest~https://github.com/mashetty330/openshift-php-upload-demo --name=file-uploader
```

Example Output:

```
Found image 4f2dcc0 (9 days old) in image stream "openshift/php" under tag "7.2-ubi8" for
"openshift/php:7.2-ubi8"
```

```
Apache 2.4 with PHP 7.2
-----
```

```
PHP 7.2 available as container is a base platform for building and running various PHP 7.2
applications and frameworks. PHP is an HTML-embedded scripting language. PHP attempts
to make it easy for developers to write dynamically generated web pages. PHP also offers
built-in database integration for several commercial and non-commercial database
management systems, so writing a database-enabled webpage with PHP is fairly simple.
The most common
use of PHP coding is probably as a replacement for CGI scripts.
```

```
Tags: builder, php, php72, php-72
```

```
* A source build using source code from https://github.com/christianh814/openshift-php-
upload-demo will be cr
eated
* The resulting image will be pushed to image stream tag "file-uploader:latest"
* Use 'oc start-build' to trigger a new build
```

```
--> Creating resources ...
  imagestream.image.openshift.io "file-uploader" created
  buildconfig.build.openshift.io "file-uploader" created
  deployment.apps "file-uploader" created
  service "file-uploader" created
```

```
--> Success
  Build scheduled, use 'oc logs -f buildconfig/file-uploader' to track its progress.
```

```
Application is not exposed. You can expose services to the outside world by executing one
or more of the commands below:
  'oc expose service/file-uploader'
```

```
Run 'oc status' to view your app.
```

3. View the build log and wait until the application is deployed.

```
$ oc logs -f bc/file-uploader -n my-shared-storage
```

Example Output:

```
Cloning "https://github.com/christianh814/openshift-php-upload-demo" ...
```

```
[...]
```

```
Generating dockerfile with builder image image-registry.openshift-image-regis
try.svc:5000/openshift/php@sha256:d97466f33999951739a76bce922ab17088885db610c
0e05b593844b41d5494ea
```

```
STEP 1: FROM image-registry.openshift-image-registry.svc:5000/openshift/php@s
ha256:d97466f33999951739a76bce922ab17088885db610c0e05b593844b41d5494ea
```

```
STEP 2: LABEL "io.openshift.build.commit.author"="Christian Hernandez <christ
ian.hernandez@yahoo.com>" "io.openshift.build.commit.date"="Sun Oct 1 1
7:15:09 2017 -0700" "io.openshift.build.commit.id"="288eda3dff43b02f7f7"
```

```

b6b6b6f93396ffdf34cb2"      "io.openshift.build.commit.ref"="master"      "
io.openshift.build.commit.message"="trying to modularize"      "io.openshift
.build.source-location"="https://github.com/christianh814/openshift-php-uploa
d-demo"      "io.openshift.build.image"="image-registry.openshift-image-regi
stry.svc:5000/openshift/php@sha256:d97466f33999951739a76bce922ab17088885db610
c0e05b593844b41d5494ea"
STEP 3: ENV OPENSIFT_BUILD_NAME="file-uploader-1"
OPENSIFT_BUILD_NAMESP
ACE="my-shared-storage"  OPENSIFT_BUILD_SOURCE="https://github.com/christ
ianh814/openshift-php-upload-demo"  OPENSIFT_BUILD_COMMIT="288eda3dff43b0
2f7f7b6b6b6f93396ffdf34cb2"
STEP 4: USER root
STEP 5: COPY upload/src /tmp/src
STEP 6: RUN chown -R 1001:0 /tmp/src
STEP 7: USER 1001
STEP 8: RUN /usr/libexec/s2i/assemble
--> Installing application source...
=> sourcing 20-copy-config.sh ...
--> 17:24:39   Processing additional arbitrary httpd configuration provide
d by s2i ...
=> sourcing 00-documentroot.conf ...
=> sourcing 50-mpm-tuning.conf ...
=> sourcing 40-ssl-certs.sh ...
STEP 9: CMD /usr/libexec/s2i/run
STEP 10: COMMIT temp.builder.openshift.io/my-shared-storage/file-uploader-1:3
b83e447
Getting image source signatures

[...]
```

The command prompt returns out of the tail mode after you see **Push successful**.



NOTE

The new-app command deploys the application directly from the git repository and does not use the OpenShift template, hence the OpenShift route resource is not created by default. You need to create the route manually.

5.7.1. Scaling the application after installation

Procedure

1. Scale the application to four replicas and expose its services to make the application zone aware and available.

```
$ oc expose svc/file-uploader -n my-shared-storage
```

```
$ oc scale --replicas=4 deploy/file-uploader -n my-shared-storage
```

```
$ oc get pods -o wide -n my-shared-storage
```

You should have four file-uploader pods in a few minutes. Repeat the above command until there are 4 file-uploader pods in the **Running** status.

2. Create a PVC and attach it into an application.

```
$ oc set volume deploy/file-uploader --add --name=my-shared-storage \
-t pvc --claim-mode=ReadWriteMany --claim-size=10Gi \
--claim-name=my-shared-storage --claim-class=ocs-storagecluster-cephfs \
--mount-path=/opt/app-root/src/uploaded \
-n my-shared-storage
```

This command:

- Creates a PVC.
- Updates the application deployment to include a volume definition.
- Updates the application deployment to attach a volume mount into the specified mount-path.
- Creates a new deployment with the four application pods.

3. Check the result of adding the volume.

```
$ oc get pvc -n my-shared-storage
```

Example Output:

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES
STORAGECLASS	AGE			
my-shared-storage	Bound	pvc-5402cc8a-e874-4d7e-af76-1eb05bd2e7c7	10Gi	RWX
ocs-storagecluster-cephfs	52s			

Notice the **ACCESS MODE** is set to RWX.

All the four **file-uploader** pods are using the same RWX volume. Without this access mode, OpenShift does not attempt to attach multiple pods to the same Persistent Volume (PV) reliably. If you attempt to scale up the deployments that are using ReadWriteOnce (RWO) PV, the pods may get colocated on the same node.

5.7.2. Modify Deployment to be Zone Aware

Currently, the **file-uploader** Deployment is not zone aware and can schedule all the pods in the same zone. In this case, if there is a site outage then the application is unavailable. For more information, see [Controlling pod placement by using pod topology spread constraints](#) .

1. Add the pod placement rule in the application deployment configuration to make the application zone aware.
 - a. Run the following command, and review the output:

```
$ oc get deployment file-uploader -o yaml -n my-shared-storage | less
```

Example Output:

```
[...]
spec:
  progressDeadlineSeconds: 600
```

```

replicas: 4
revisionHistoryLimit: 10
selector:
  matchLabels:
    deployment: file-uploader
strategy:
  rollingUpdate:
    maxSurge: 25%
    maxUnavailable: 25%
  type: RollingUpdate
template:
  metadata:
    annotations:
      openshift.io/generated-by: OpenShiftNewApp
    creationTimestamp: null
  labels:
    deployment: file-uploader
  spec: # <-- Start inserted lines after here
    containers: # <-- End inserted lines before here
      - image: image-registry.openshift-image-registry.svc:5000/my-shared-storage/file-
        uploader@sha256:a458ea62f990e431ad7d5f84c89e2fa27bdebdd5e29c5418c70c56eb81f
        0a26b
        imagePullPolicy: IfNotPresent
        name: file-uploader
[...]
```

- b. Edit the deployment to use the topology zone labels.

```
$ oc edit deployment file-uploader -n my-shared-storage
```

Add the following new lines between the **Start** and **End** (shown in the output in the previous step):

```

[...]
```

```

spec:
  topologySpreadConstraints:
    - labelSelector:
        matchLabels:
          deployment: file-uploader
      maxSkew: 1
      topologyKey: topology.kubernetes.io/zone
      whenUnsatisfiable: DoNotSchedule
    - labelSelector:
        matchLabels:
          deployment: file-uploader
      maxSkew: 1
      topologyKey: kubernetes.io/hostname
      whenUnsatisfiable: ScheduleAnyway
  nodeSelector:
    node-role.kubernetes.io/worker: ""
  containers:
[...]
```

Example output:

```
deployment.apps/file-uploader edited
```

2. Scale down the deployment to **zero** pods and then back to **four** pods. This is needed because the deployment changed in terms of pod placement.

Scaling down to zero pods

```
$ oc scale deployment file-uploader --replicas=0 -n my-shared-storage
```

Example output:

```
deployment.apps/file-uploader scaled
```

Scaling up to four pods

```
$ oc scale deployment file-uploader --replicas=4 -n my-shared-storage
```

Example output:

```
deployment.apps/file-uploader scaled
```

3. Verify that the four pods are spread across the four nodes in datacenter1 and datacenter2 zones.

```
$ oc get pods -o wide -n my-shared-storage | egrep '^file-uploader' | grep -v build | awk '{print $7}' | sort | uniq -c
```

Example output:

```
1 perf1-mz8bt-worker-d2hdm
1 perf1-mz8bt-worker-k68rv
1 perf1-mz8bt-worker-ntkp8
1 perf1-mz8bt-worker-qpwsr
```

Search for the zone labels used.

```
$ oc get nodes -L topology.kubernetes.io/zone | grep datacenter | grep -v master
```

Example output:

```
perf1-mz8bt-worker-d2hdm Ready worker 35d v1.20.0+5fbfd19 datacenter1
perf1-mz8bt-worker-k68rv Ready worker 35d v1.20.0+5fbfd19 datacenter1
perf1-mz8bt-worker-ntkp8 Ready worker 35d v1.20.0+5fbfd19 datacenter2
perf1-mz8bt-worker-qpwsr Ready worker 35d v1.20.0+5fbfd19 datacenter2
```

4. Use the file-uploader web application using your browser to upload new files.
 - a. Find the route that is created.

```
$ oc get route file-uploader -n my-shared-storage -o jsonpath --
template="http://{.spec.host}{\n}"
```

Example Output:

```
http://file-uploader-my-shared-storage.apps.cluster-ocs4-abdf.ocs4-abdf.sandbox744.opentlc.com
```

- b. Point your browser to the web application using the route in the previous step.
The web application lists all the uploaded files and offers the ability to upload new ones as well as you download the existing data. Right now, there is nothing.
- c. Select an arbitrary file from your local machine and upload it to the application.
 - i. Click **Choose file** to select an arbitrary file.
 - ii. Click **Upload**.

Figure 5.1. A simple PHP-based file upload tool

- d. Click **List uploaded files** to see the list of all currently uploaded files.



NOTE

The OpenShift Container Platform image registry, ingress routing, and monitoring services are not zone aware.

5.8. RECOVERING OPENSIFT DATA FOUNDATION STRETCH CLUSTER

Given that the stretch cluster disaster recovery solution is to provide resiliency in the face of a complete or partial site outage, it is important to understand the different methods of recovery for applications and their storage.

How the application is architected determines how soon it becomes available again on the active zone.

There are different methods of recovery for applications and their storage depending on the site outage. The recovery time depends on the application architecture. The different methods of recovery are as follows:

- [Recovering zone-aware HA applications with RWX storage](#) .
- [Recovering HA applications with RWX storage](#) .
- [Recovering applications with RWO storage](#) .
- [Recovering StatefulSet pods](#) .

5.8.1. Understanding zone failure

For the purpose of this section, zone failure is considered as a failure where all OpenShift Container Platform, master and worker nodes in a zone are no longer communicating with the resources in the second data zone (for example, powered down nodes). If communication between the data zones is still partially working (intermittently up or down), the cluster, storage, and network admins should disconnect the communication path between the data zones for recovery to succeed.



IMPORTANT

When you install the sample application, power off the OpenShift Container Platform nodes (at least the nodes with OpenShift Data Foundation devices) to test the failure of a data zone in order to validate that your file-uploader application is available, and you can upload new files.




NOTE

To prevent corruption of volume contents during a non-graceful node shutdown, follow the steps in [Fencing and tainting nodes to prevent volume corruption during non-graceful node shutdown](#).

5.8.2. Recovering zone-aware HA applications with RWX storage

Applications that are deployed with **topologyKey: topology.kubernetes.io/zone** have one or more replicas scheduled in each data zone, and are using shared storage, that is, ReadWriteMany (RWX) CephFS volume, terminate themselves in the failed zone after few minutes and new pods are rolled in and stuck in pending state until the zones are recovered.

An example of this type of application is detailed in the [Install Zone Aware Sample Application](#) section.



IMPORTANT

During zone recovery if application pods go into CrashLoopBackOff (CLBO) state with permission denied error while mounting the CephFS volume, then restart the nodes where the pods are scheduled. Wait for some time and then check if the pods are running again.

**NOTE**

To prevent corruption of volume contents during a non-graceful node shutdown, follow the steps in [Fencing and tainting nodes to prevent volume corruption during non-graceful node shutdown](#).

5.8.3. Recovering HA applications with RWX storage

Applications that are using **topologyKey: kubernetes.io/hostname** or no topology configuration have no protection against all of the application replicas being in the same zone.

**NOTE**

This can happen even with *podAntiAffinity* and **topologyKey: kubernetes.io/hostname** in the **Pod** spec because this anti-affinity rule is host-based and not zone-based.

If this happens and all replicas are located in the zone that fails, the application using ReadWriteMany (RWX) storage takes 6-8 minutes to recover on the active zone. This pause is for the OpenShift Container Platform nodes in the failed zone to become **NotReady** (60 seconds) and then for the default pod eviction timeout to expire (300 seconds).

**IMPORTANT**

In order to prevent corruption of volume contents during a non-graceful node shutdown, follow the steps in [Fencing and tainting nodes to prevent volume corruption during non-graceful node shutdown](#).

5.8.4. Recovering applications with RWO storage

Applications that use ReadWriteOnce (RWO) storage have a known behavior described in this [Kubernetes issue](#). Because of this issue, if there is a data zone failure, any application pods in that zone mounting RWO volumes (for example, **cephrbd** based volumes) are stuck with **Terminating** status after 6-8 minutes and are not re-created on the active zone without manual intervention.

To prevent corruption of volume contents during a non-graceful node shutdown, follow the steps in [Fencing and tainting nodes to prevent volume corruption during non-graceful node shutdown](#).

5.8.5. Recovering StatefulSet pods

Pods that are part of a StatefulSet have a similar issue as pods mounting ReadWriteOnce (RWO) volumes. More information is referenced in the Kubernetes resource [StatefulSet considerations](#).

To prevent corruption of volume contents during a non-graceful node shutdown, follow the steps in [Fencing and tainting nodes to prevent volume corruption during non-graceful node shutdown](#).

5.8.6. Fencing and tainting nodes to prevent volume corruption during non-graceful node shutdown

The following procedure needs to be executed during a non-graceful node shutdown to prevent corruption of volume contents.

Prerequisites

- Ensure that the **NetworkFenceClass** is created by following the steps in the *Create `NetworkFenceClass`* section of [Creating OpenShift Data Foundation cluster](#).

Procedure

1. Create **NetworkFence** CR to fence the IP of the storage nodes that are down.

- a. Get the **nodeName** that needs to be fenced:

```
$ oc get node
```

- b. Get the **csiaddonsnode** object corresponding to the node found in the previous step:

```
$ oc get csiaddonsnode -n openshift-storage |grep -i <nodename> |grep -i daemonset
```

- c. Get the IP address from the **csiaddonsnode** object found in the previous step, and create the **NetworkFence** CR and ensure that it is in the **Fenced** state:

```
$ oc get csiaddonsnode <csi-addons-name> -o
jsonpath='{.status.networkFenceClientStatus}'
[{"ClientDetails":{"cidrs":["10.244.0.1/32"],"id":"a815fe8e-eabd-4e87-a6e8-
78cebfb67d08"},"networkFenceClassName":"networkfenceclass-sample"}]
```

Replace **<csi-addons-name>** with the **csiaddonsnode** object found in the earlier step.

- d. Create the **NetworkFence** CR for above CIDR with the OpenShift node name for easy identification:

```
apiVersion: csiaddons.openshift.io/v1alpha1
kind: NetworkFence
metadata:
  name: <openshift-node-name>
spec:
  cidrs:
    - 10.244.0.1/32
  driver: openshift-storage.rbd.csi.ceph.com
  fenceState: Fenced
  parameters:
    clusterID: openshift-storage
  secret:
    name: rook-csi-rbd-provisioner
    namespace: openshift-storage
```

- **parameters.clusterID**: specifies the cluster id and can be fetched from **parameters.clusterID** in **ocs-storagecluster-ceph-rbd** storageclass.
- **secret.name**: specifies the name of the secret required for network fencing operation. This can be fetched from **parameters.csi.storage.k8s.io/provisioner-secret-name** in the **ocs-storagecluster-ceph-rbd** storageclass.
- **secret.namespace**: specifies the namespace in which the secret is located. This can be fetched from **parameters.csi.storage.k8s.io/provisioner-secret-namespace** in **ocs-storagecluster-ceph-rbd** storageclass.

- e. Verify that network fencing has been done successfully:

```
$ oc get networkfence
```

Example output:

```
NAME          DRIVER          CIDRS          FENCESTATE  AGE  RESULT
<openshift-node-name> openshift-storage.rbd.csi.ceph.com ["10.244.0.1/32"] Fenced
42h Succeeded
```

2. Add taint for the nodes that are down:

```
$ oc adm taint nodes <node-name> node.kubernetes.io/out-of-
service=nodeshutdown:NoExecute
```

3. Monitor application relocation to surviving nodes.
Perform the above steps for all the nodes that are down.

Recovery after the node is back

After the nodes are back online, you need to unfence the node and remove the taints on the node to schedule the pods on these nodes.

1. Change the state of **NetworkFence** to **UnFenced** in the CR and wait for it to become unfenced
 - a. Set the **NetworkFence** to **UnFenced**:

```
$ oc patch networkfence <network-fence-object-name> -p '{"spec":
{"fenceState":"Unfenced"}}' --type=merge
```

- b. Verify the Network UnFencing is done successfully:

```
$ oc get networkfence
```

Example output:

```
NAME          DRIVER          CIDRS          FENCESTATE  AGE  RESULT
<openshift-node-name> openshift-storage.rbd.csi.ceph.com ["10.244.0.1/32"]
unfenced 42h Succeeded
```

2. Reboot the node.
3. Untaint the nodes that were previously tainted:

```
$ oc adm taint nodes <node-name> node.kubernetes.io/out-of-
service=nodeshutdown:NoExecute-
```

Repeat these steps for all nodes.

CHAPTER 6. DISABLING DISASTER RECOVERY FOR A DISASTER RECOVERY ENABLED APPLICATION

This section guides you to disable disaster recovery (DR) for an application deployed using Red Hat Advanced Cluster Management (RHACM).

- [Disabling DR managed applications](#).
- [Disabling DR discovered applications](#).


6.1. DISABLING DR MANAGED APPLICATIONS

1. In the RHACM console, navigate to **All Clusters** → **Data Services** → **Protected applications** tab.
2. At the end of the application row, click on the **Actions** menu and choose **Remove disaster recovery**.
3. Click **Remove** in the next prompt.



WARNING

Your application will lose disaster recovery protection, preventing volume synchronization (replication) between clusters.



NOTE

The application is no longer in the **Protected applications** tab once the DR is removed.


6.2. DISABLING DR DISCOVERED APPLICATIONS

1. In the RHACM console, navigate to **All Clusters** → **Data Services** → **Protected applications** tab.
2. At the end of the application row, click on the **Actions** menu and choose **Remove disaster recovery**.
3. Click **Remove** in the next prompt.



WARNING

Your application will lose disaster recovery protection, preventing volume synchronization (replication) between clusters.



IMPORTANT

Due to a [known issue](#), CephFS PVCs can become orphaned after disabling DR protection for a discovered workload on the cluster where the application is not running. To clean up the orphaned CephFS PVCs, manually delete them using the following command:

```
$ oc delete pvc <pvc-name> -n <application-namespace>
```



NOTE

The application is no longer in the **Protected applications** tab once the DR is removed.

CHAPTER 7. UNINSTALLING DISASTER RECOVERY FROM HUB AND MANAGED CLUSTERS

7.1. UNINSTALLING DISASTER RECOVERY FROM HUB AND MANAGED CLUSTERS

There are two levels of uninstallation for both Regional-DR and Metro-DR:


Partial Uninstallation This involves removing all DR resources and configurations from two or more peer managed clusters defined in a DRPolicy.

Complete Uninstallation This is a full uninstall of the DR setup and includes the following steps:

- Remove all DRPolicies and mirrorpeers from all peer clusters.
- Remove the DR operators and associated resources from the hub cluster.

7.1.1. Before uninstalling disaster recovery

Before uninstalling DR, you must first remove disaster recovery protection from all protected applications, both managed and discovered. Follow these steps to remove disaster recovery protection.



IMPORTANT

Both managed clusters must be online and healthy before DR Uninstall is attempted.

1. Find all DRPolicies associated with the peer clusters you want to remove DR from using one of the following options:
 - a. **Using the RHACM console**
 - i. On the hub cluster, navigate to **All Clusters → Data Services → Disaster Recovery → Policies**.
 - ii. Identify any DRPolicies associated with the peer clusters you want to remove DR from.
 - b. **Using the CLI**
 - i. Get the names of the peer clusters from the RHACM console. On the hub cluster, navigate to **All Clusters → Infrastructure → Clusters**.
 - ii. Run the following commands on the hub cluster, replacing the cluster names with the peer clusters found in the previous step:

```
$ CLUSTER_A="{cluster1_name}"
```

```
$ CLUSTER_B="{cluster2_name}"
```

```
$ oc get drpolicies -o json | jq -r --arg c1 "$CLUSTER_A" --arg c2 "$CLUSTER_B"
'.items[] |select((.spec.drClusters | sort) == ([ $c1, $c2 ] | sort)) .metadata.name'
```

- iii. Identify any DRPolicies associated with the peer clusters you want to remove DR from.

2. Find any existing DRPlacementControl (DRPC) resources associated with your peer managed clusters. Run the following on the hub cluster to for each DRPolicy identified in step 1:

```
$ DRPOLICY_NAME="{drpolicy_name}"
```

Replace **drpolicy_name** with the DRPolicy.

```
$ oc get drplacementcontrols -A -o json | jq -r --arg drpc "$DRPOLICY_NAME" \ '.items[] |
select(.spec.drPolicyRef.name == $drpc) | "\(.metadata.namespace)/\(.metadata.name)'"
```

3. If any DRPCs were found, remove DR protection from the associated applications.
 - a. Remove DR protection from managed applications:
 - i. On the Hub cluster, navigate to **All Clusters → Applications**.
 - ii. In the **Overview** tab, at the end of the protected application row from the action menu (**⋮**), select **Manage disaster recovery**.
 - iii. Click **Remove disaster recovery**.
 - iv. Click **Confirm remove**.
 - b. Remove DR protection from discovered applications:
 - i. In the RHACM console, navigate to **All Clusters → Data Services → Protected applications** tab.
 - ii. At the end of the application row, click on the action menu (**⋮**) and choose **Remove disaster recovery**.
 - iii. Click **Remove** in the next prompt.

4. Run the following on the hub cluster to find the associated mirrorpeer:

```
$ export DR_POLICY={drpolicy_name}
```


Replace **drpolicy_name** with one of the DRPolicy names found in step 1.

```
$ DR_CLUSTERS=$(oc get drpolicy $DR_POLICY -o json | jq -r '.spec.drClusters | sort |
join(",")')
```

```
$ oc get mirrorpeers -A -o json | jq -r --arg dr_clusters "$DR_CLUSTERS" '.items[] | {name:
.metadata.name,clusters: (.spec.items | map(clusterName) | sort | join(",")) } |select(.clusters
== $dr_clusters) | .name'
```

7.1.2. Uninstalling disaster recovery for peer managed clusters

Follow these steps to uninstall DR resources for the peer managed clusters.



NOTE

You will need the cluster names, DRPolicies, and mirrorpeer identified in the previous section.

Procedure

1. From the hub cluster, delete DRPolicies associated with the peer managed clusters slated for DR removal.

There will be one or more DRPolicies found in the prior section that need to be deleted.


```
$ oc delete drpolicy {drpolicy_name1 drpolicy_name2 ...}
```

Replace **drpolicy_name1** etc. with the DRPolicy names found in the previous section.

2. From the hub cluster, delete the mirrorpeer associated with the peer managed clusters slated for DR removal.

There will be one mirrorpeer per DRPolicy found in the prior section that needs to be deleted.

```
$ oc delete mirrorpeer {mirrorpeer_name}
```



IMPORTANT

If you are only removing DR for specific peer clusters, do not proceed to uninstall DR operators or other shared resources. This concludes the partial DR uninstall process.

7.1.3. Uninstalling disaster recovery operators

If your goal is to completely remove DR from your OpenShift environment, make sure you have followed all the previous sections to uninstall DR from every peer managed cluster.

Procedure

1. On the hub cluster, delete the **subscriptions.operators** and **clusterserviceversion** (csv) for the OpenShift Data Foundation Multicloud Orchestrator and OpenShift DR Hub Operator:

```
$ oc delete subscriptions.operators.coreos.com -l operators.coreos.com/odr-hub-operator.openshift-operators -n openshift-operators
$ oc delete subscriptions.operators.coreos.com -l operators.coreos.com/odf-multicloud-orchestrator.openshift-operators -n openshift-operators
$ oc delete csv -l operators.coreos.com/odr-hub-operator.openshift-operators -n openshift-operators
$ oc delete csv -l operators.coreos.com/odf-multicloud-orchestrator.openshift-operators -n openshift-operators
```

2. Alternatively, you can delete the ODF Multicloud Orchestrator and DR Hub Operator from the RHACM console:
 - a. Go to **Ecosystem → Installed Operators**.
 - b. Locate the DR Hub Operator.
 - c. Click the action menu (⋮) and select **Uninstall Operator**.
 - d. Repeat the same for the Multicloud Orchestrator.
3. From the hub cluster, delete the **managedclusterview** for each peer cluster:

```
$ oc delete managedclusterview -l multicluster.odf.openshift.io/created-by=odf-multicluster-managedcluster-controller -A
```

4. From the hub cluster, delete the dynamic console plugin for the Multicluster Orchestrator operator:

```
$ oc delete consoleplugins.console.openshift.io odf-multicluster-console
```

5. From the hub cluster, delete the configmap **odf-client-info**:

```
$ oc delete configmap odf-client-info -n openshift-operators
```

6. From the hub cluster, delete the **odf-multicluster-console** service:

```
$ oc delete service odf-multicluster-console -n openshift-operators
```

At this point, all DR-related operator components and configurations should be fully removed from the hub cluster and all managed clusters.

CHAPTER 8. MONITORING DISASTER RECOVERY HEALTH

8.1. ENABLE MONITORING FOR DISASTER RECOVERY

Use this procedure to enable basic monitoring for your disaster recovery setup.

Procedure

1. On the Hub cluster, open a terminal window
2. Add the following label to **openshift-operator** namespace.

```
$ oc label namespace openshift-operators openshift.io/cluster-monitoring='true'
```



NOTE

You must always add this label for Regional-DR solution.

8.2. ENABLING DISASTER RECOVERY DASHBOARD ON HUB CLUSTER

This section guides you to enable the disaster recovery dashboard for advanced monitoring on the Hub cluster.

For Regional-DR, the dashboard shows monitoring status cards for operator health, cluster health, metrics, alerts and application count.

For Metro-DR, you can configure the dashboard to only monitor the ramen setup health and application count.

Prerequisites

- Ensure that you have already installed the following
 - OpenShift Container Platform version 4.17 and have administrator privileges.
 - ODF Multicluster Orchestrator with the console plugin enabled.
 - Red Hat Advanced Cluster Management for Kubernetes 2.11 (RHACM) from Software Catalog. For instructions on how to install, see [Installing RHACM](#).
- Ensure you have enabled observability on RHACM. See [Enabling observability guidelines](#).

Procedure

1. On the Hub cluster, open a terminal window and perform the next steps.
2. Create the configmap file named **observability-metrics-custom-allowlist.yaml**.
You can use the following YAML to list the disaster recovery metrics on Hub cluster. For details, see [Adding custom metrics](#). To know more about ramen metrics, see [Disaster recovery metrics](#).

```
kind: ConfigMap
apiVersion: v1
metadata:
```

```

name: observability-metrics-custom-allowlist
namespace: open-cluster-management-observability
data:
  metrics_list.yaml: |
    names:
      - ceph_rbd_mirror_snapshot_sync_bytes
      - ceph_rbd_mirror_snapshot_snapshots
    matches:
      - __name__="csv_succeeded",exported_namespace="openshift-dr-system",name=~"odr-cluster-operator.*"
      - __name__="kube_running_pod_ready",namespace=~"volsync-system|openshift-operators",pod=~"volsync.*"

```

3. In the **open-cluster-management-observability** namespace, run the following command:

```
$ oc apply -n open-cluster-management-observability -f observability-metrics-custom-allowlist.yaml
```

4. After **observability-metrics-custom-allowlist** yaml is created, RHACM starts collecting the listed OpenShift Data Foundation metrics from all the managed clusters.
To exclude a specific managed cluster from collecting the observability data, add the following cluster label to the **clusters: observability: disabled**.

8.3. VIEWING HEALTH STATUS OF DISASTER RECOVERY REPLICATION RELATIONSHIPS

Prerequisites

Ensure that you have enabled the disaster recovery dashboard for monitoring. For instructions, see chapter [Enabling disaster recovery dashboard on Hub cluster](#).

Procedure

1. On the Hub cluster, ensure **All Clusters** option is selected.
2. Refresh the console to make the DR monitoring dashboard tab accessible.
3. Navigate to **Data Services** and click **Data policies**.
4. On the Overview tab, you can view the health status of the operators, clusters and applications. Green tick indicates that the operators are running and available..
5. Click the **Disaster recovery** tab to view a list of DR policy details and connected applications.

8.4. VIEWING HEALTH STATUS OF APPLICATIONSET-BASED APPLICATIONS

You can view the health status of your ApplicationSet-based applications on the user interface. This health status can alert you to replication delays and make for easier troubleshooting.

Procedure

1. On the Hub cluster, navigate to **All Clusters → Applications**.

2. In the **DR Status** column of the application, view the status:

- **Healthy** The last group sync time is less than 2X that of the sync interval.
- **Warning** The last group sync time is greater than 2X and less than 3X of the sync interval.
- **Critical** The last group sync time is greater than 3X of the sync interval.

You can also view the application status for virtual machine (VM) based workloads:

- On the Hub cluster, navigate to **All Clusters → Infrastructure → Virtual Machines → Status** column.

8.5. DISASTER RECOVERY METRICS

These are the ramen metrics that are scrapped by prometheus.

- `ramen_last_sync_timestamp_seconds`
- `ramen_policy_schedule_interval_seconds`
- `ramen_last_sync_duration_seconds`
- `ramen_last_sync_data_bytes`
- `ramen_workload_protection_status`

Run these metrics from the Hub cluster where Red Hat Advanced Cluster Management for Kubernetes (RHACM operator) is installed.

8.5.1. Last synchronization timestamp in seconds

This is the time in seconds which gives the time of the most recent successful synchronization of all PVCs per application.

Metric name

`ramen_last_sync_timestamp_seconds`

Metrics type

Gauge

Labels

- **ObjType:** Type of the object, here its DRPC
- **ObjName:** Name of the object, here it is DRPC-Name
- **ObjNamespace:** DRPC namespace
- **Policynamespace:** Name of the DRPolicy
- **SchedulingInterval:** Scheduling interval value from DRPolicy

Metric value

Value is set as Unix seconds which is obtained from **lastGroupSyncTime** from DRPC status.

8.5.2. Policy schedule interval in seconds

This gives the scheduling interval in seconds from DRPolicy.

Metric name

ramen_policy_schedule_interval_seconds

Metrics type

Gauge

Labels

- **Policynamespace:** Name of the DRPolicy

Metric value

This is set to a scheduling interval in seconds which is taken from DRPolicy.

8.5.3. Last synchronization duration in seconds

This represents the longest time taken to sync from the most recent successful synchronization of all PVCs per application.

Metric name

ramen_last_sync_duration_seconds

Metrics type

Gauge

Labels

- **obj_type:** Type of the object, here it is DRPC
- **obj_name:** Name of the object, here it is DRPC-Name
- **obj_namespace:** DRPC namespace
- **scheduling_interval:** Scheduling interval value from DRPolicy

Metric value

The value is taken from **lastGroupSyncDuration** from DRPC status.

8.5.4. Total bytes transferred from most recent synchronization

This value represents the total bytes transferred from the most recent successful synchronization of all PVCs per application.

Metric name

ramen_last_sync_data_bytes

Metrics type

Gauge

Labels

- **obj_type:** Type of the object, here it is DRPC

- **obj_name**: Name of the object, here it is DRPC-Name
- **obj_namespace**: DRPC namespace
- **scheduling_interval**: Scheduling interval value from DRPolicy

Metric value

The value is taken from **lastGroupSyncBytes** from DRPC status.

8.5.5. Workload protection status

This value provides the application protection status per application that is DR protected.

Metric name

ramen_workload_protection_status

Metrics type

Gauge

Labels

- **ObjType**: Type of the object, here its DRPC
- **ObjName**: Name of the object, here it is DRPC-Name
- **ObjNamespace**: DRPC namespace

Metric value

The value is either a "1" or a "0", where "1" indicates application DR protection is healthy and a "0" indicates application protection degraded and potentially unprotected.

8.6. DISASTER RECOVERY ALERTS

This section provides a list of all supported alerts associated with Red Hat OpenShift Data Foundation within a disaster recovery environment.

Recording rules

- Record: **ramen_sync_duration_seconds**

Expression

```
sum by (obj_name, obj_namespace, obj_type, job, policyname)(time() -
(ramen_last_sync_timestamp_seconds > 0))
```

Purpose

The time interval between the volume group's last sync time and the time now in seconds.

- Record: **ramen_rpo_difference**

Expression

```
ramen_sync_duration_seconds{job="ramen-hub-operator-metrics-service"} /
on(policyname, job) group_left() (ramen_policy_schedule_interval_seconds{job="ramen-
```

```
hub-operator-metrics-service"))
```

Purpose

The difference between the expected sync delay and the actual sync delay taken by the volume replication group.

- **Record: count_persistentvolumeclaim_total**

Expression

```
count(kube_persistentvolumeclaim_info)
```

Purpose

Sum of all PVC from the managed cluster.

Alerts

- **Alert: VolumeSynchronizationDelay**

Impact

Critical

Purpose

Actual sync delay taken by the volume replication group is thrice the expected sync delay.

YAML

```
alert: VolumeSynchronizationDelay
expr: ramen_rpo_difference >= 3
for: 5s
labels:
  severity: critical
annotations:
  description: "The syncing of volumes is exceeding three times the scheduled snapshot
interval, or the volumes have been recently protected. (DRPC: {{ $labels.obj_name }},
Namespace: {{ $labels.obj_namespace }})"
  alert_type: "DisasterRecovery"
```

- **Alert: VolumeSynchronizationDelay**

Impact

Warning

Purpose

Actual sync delay taken by the volume replication group is twice the expected sync delay.

YAML

```
alert: VolumeSynchronizationDelay
expr: ramen_rpo_difference > 2 and ramen_rpo_difference < 3
for: 5s
labels:
  severity: warning
annotations:
  description: "The syncing of volumes is exceeding two times the scheduled snapshot
```

```
interval, or the volumes have been recently protected. (DRPC: {{ $labels.obj_name }},  
Namespace: {{ $labels.obj_namespace }})"  
alert_type: "DisasterRecovery"
```

- **Alert: WorkloadUnprotected**

Impact

Critical

Purpose

Application protection status is degraded for more than 10 minutes

YAML

```
alert: WorkloadUnprotected  
expr: ramen_workload_protection_status == 0  
for: 10m  
labels:  
  severity: critical  
annotations:  
  description: "Workload is not protected for disaster recovery (DRPC: {{ $labels.obj_name  
}}, Namespace: {{ $labels.obj_namespace }})."  
  alert_type: "DisasterRecovery"
```

CHAPTER 9. TROUBLESHOOTING DISASTER RECOVERY

9.1. TROUBLESHOOTING METRO-DR

9.1.1. A statefulset application stuck after failover

Problem

While relocating to a preferred cluster, DRPlacementControl is stuck reporting PROGRESSION as "MovingToSecondary".

Previously, before Kubernetes v1.23, the Kubernetes control plane never cleaned up the PVCs created for StatefulSets. This activity was left to the cluster administrator or a software operator managing the StatefulSets. Due to this, the PVCs of the StatefulSets were left untouched when their Pods were deleted. This prevents Ramen from relocating an application to its preferred cluster.

Resolution

1. If the workload uses StatefulSets, and relocation is stuck with PROGRESSION as "MovingToSecondary", then run:

```
$ oc get pvc -n <namespace>
```

2. For each bounded PVC for that namespace that belongs to the StatefulSet, run

```
$ oc delete pvc <pvcname> -n namespace
```

Once all PVCs are deleted, Volume Replication Group (VRG) transitions to secondary, and then gets deleted.

3. Run the following command

```
$ oc get drpc -n <namespace> -o wide
```

After a few seconds to a few minutes, the PROGRESSION reports "Completed" and relocation is complete.

Result

The workload is relocated to the preferred cluster

BZ reference: [\[2118270\]](#)

9.1.2. DR policies protect all applications in the same namespace

Problem

While only a single application is selected to be used by a DR policy, all applications in the same namespace will be protected. This results in PVCs, that match the **DRPlacementControl spec.pvcSelector** across multiple workloads or if the selector is missing across all workloads, replication management to potentially manage each PVC multiple times and cause data corruption or invalid operations based on individual **DRPlacementControl** actions.

Resolution

Label PVCs that belong to a workload uniquely, and use the selected label as the DRPlacementControl **spec.pvcSelector** to disambiguate which DRPlacementControl protects and

manages which subset of PVCs within a namespace. It is not possible to specify the **spec.pvcSelector** field for the DRPlacementControl using the user interface, hence the DRPlacementControl for such applications must be deleted and created using the command line.

BZ reference: [\[2128860\]](#)

9.1.3. Relocate or failback might be stuck in Initiating state

Problem

When a primary cluster is down and comes back online while the secondary goes down, **relocate** or **failback** might be stuck in the **Initiating** state.

Resolution

To avoid this situation, cut off all access from the old active hub to the managed clusters. Alternatively, you can scale down the ApplicationSet controller on the old active hub cluster either before moving workloads or when they are in the clean-up phase.

On the old active hub, scale down the two deployments using the following commands:

```
$ oc scale deploy -n openshift-gitops-operator openshift-gitops-operator-controller-manager --replicas=0

$ oc scale statefulset -n openshift-gitops openshift-gitops-application-controller --replicas=0
```

BZ reference: [\[2243804\]](#)

9.2. TROUBLESHOOTING REGIONAL-DR

9.2.1. rbd-mirror daemon health is in warning state

Problem

There appears to be numerous cases where WARNING gets reported if mirror service **::get_mirror_service_status** calls **Ceph** monitor to get service status for **rbd-mirror**. Following a network disconnection, **rbd-mirror** daemon health is in the **warning** state while the connectivity between both the managed clusters is fine.

Resolution

Run the following command in the toolbox and look for **leader:false**

```
rbd mirror pool status --verbose ocs-storagecluster-cephblockpool | grep 'leader:'
```

If you see the following in the output:

**leader:
false**

It indicates that there is a daemon startup issue and the most likely root cause could be due to problems reliably connecting to the secondary cluster.

Workaround: Move the **rbd-mirror** pod to a different node by simply deleting the pod and verify that it has been rescheduled on another node.

leader: true or no output	Contact Red Hat Support
---	---

BZ reference: [\[2118627\]](#)

9.2.2. volsync-rsync-src pod is in error state as it is unable to resolve the destination hostname

Problem

VolSync source pod is unable to resolve the hostname of the VolSync destination pod. The log of the VolSync Pod consistently shows an error message over an extended period of time similar to the following log snippet.

```
$ oc logs -n busybox-workloads-3-2 volsync-rsync-src-dd-io-pvc-1-p25rz
```

Example output

```
VolSync rsync container version: ACM-0.6.0-ce9a280
Syncing data to volsync-rsync-dst-dd-io-pvc-1.busybox-workloads-3-2.svc.clusterset.local:22 ...
ssh: Could not resolve hostname volsync-rsync-dst-dd-io-pvc-1.busybox-workloads-3-2.svc.clusterset.local: Name or service not known
```

Resolution

Restart **submariner-lighthouse-agent** on both nodes.

```
$ oc delete pod -l app=submariner-lighthouse-agent -n submariner-operator
```

9.2.3. Cleanup and data sync for ApplicationSet workloads remain stuck after older primary managed cluster is recovered post failover

Problem

ApplicationSet based workload deployments to managed clusters are not garbage collected in cases when the hub cluster fails. It is recovered to a standby hub cluster, while the workload has been failed over to a surviving managed cluster. The cluster that the workload was failed over from, rejoins the new recovered standby hub.

ApplicationSets that are DR protected, with a regional DRPolicy, hence starts firing the VolumeSynchronizationDelay alert. Further such DR protected workloads cannot be failed over to the peer cluster or relocated to the peer cluster as data is out of sync between the two clusters.

Resolution

The workaround requires that **openshift-gitops** operators can own the workload resources that are orphaned on the managed cluster that rejoined the hub post a failover of the workload was performed from the new recovered hub. To achieve this the following steps can be taken:

1. Determine the Placement that is in use by the ArgoCD ApplicationSet resource on the hub cluster in the **openshift-gitops** namespace.
2. Inspect the placement label value for the ApplicationSet in this field:
spec.generators.clusterDecisionResource.labelSelector.matchLabels
This would be the name of the Placement resource *<placement-name>*
3. Ensure that there exists a **PlacemenDecision** for the ApplicationSet referenced **Placement**.

```
$ oc get placementdecision -n openshift-gitops --selector cluster.open-cluster-management.io/placement=<placement-name>
```

This results in a single **PlacementDecision** that places the workload in the currently desired failover cluster.

4. Create a new **PlacementDecision** for the ApplicationSet pointing to the cluster where it should be cleaned up.
For example:

```
apiVersion: cluster.open-cluster-management.io/v1beta1
kind: PlacementDecision
metadata:
  labels:
    cluster.open-cluster-management.io/decision-group-index: "1" # Typically one higher
    cluster.open-cluster-management.io/decision-group-name: ""
    cluster.open-cluster-management.io/placement: cephfs-appset-busybox10-placement
    name: <placemen-name>-decision-<n> # <n> should be one higher than the existing
    namespace: openshift-gitops
```

5. Update the newly created **PlacementDecision** with a status **subresource**.

```
decision-status.yaml:
status:
  decisions:
    - clusterName: <managedcluster-name-to-clean-up> # This would be the cluster from
      where the workload was failed over, NOT the current workload cluster
      reason: FailoverCleanup
```

```
$ oc patch placementdecision -n openshift-gitops <placemen-name>-decision-<n> --
patch-file=decision-status.yaml --subresource=status --type=merge
```

6. Watch and ensure that the Application resource for the ApplicationSet has been placed on the desired cluster

```
$ oc get application -n openshift-gitops <applicationset-name>-<managedcluster-name-
to-clean-up>
```

In the output, check if the SYNC STATUS shows as **Synced** and the HEALTH STATUS shows as **Healthy**.

7. Delete the PlacementDecision that was created in step (3), such that ArgoCD can garbage collect the workload resources on the *<managedcluster-name-to-clean-up>*

```
$ oc delete placementdecision -n openshift-gitops <placement-name>-decision-<n>
```

ApplicationSets that are DR protected, with a regional DRPolicy, stops firing the **VolumeSynchronizationDelay** alert.

BZ reference: [\[2268594\]](#)

9.2.4. Secondary transition failed as PVC is potentially in use by a pod

Problem

If you see the error **Secondary transition failed as PVC is potentially in use by a pod** after a failover and the failed cluster has recovered, it usually means application cleanup wasn't completed on the primary cluster. This step is critical to release the PVC and allow replication to resume.

Resolution

Complete cleanup on the primary cluster. Replication should resume automatically.

If replication doesn't resume, perform a Relocate action to move the application to the secondary cluster. This ensures recovery using the last synchronized data.

9.2.5. vm-recipe related known issues

9.2.5.1. ramendr.openshift.io/k8s-resource-selector label is not automatically applied when a PVC is added to an already protected workload NS

Problem

When an additional PV or PVC is attached to an already protected VM on Day 2, the newly created PV/PVC is not automatically protected.

Resolution

To ensure the new PV/PVC is protected, manually add the following label on the new PV/PVC:

```
("ramendr.openshift.io/k8s-resource-selector": "<value same as that of base PVC of associated VM>")
```

This label links the new PV/PVC to the existing protected group of resources.

9.2.5.2. DRPolicy without flattening applied to a cloned volume leads to errors

Problem

Regional-DR currently does not support volumes created from CSI snapshots or clones in protected applications. These volumes rely on a parent volume that isn't replicated.

Resolution

To avoid these issues, use a DRPolicy that flattens PVs before replication if cloning is required. For more information, see this [knowledgebase article](#).

9.3. TROUBLESHOOTING 2-SITE STRETCH CLUSTER WITH ARBITER

9.3.1. Recovering workload pods stuck in **ContainerCreating** state post zone recovery

Problem

After performing complete zone failure and recovery, the workload pods are sometimes stuck in **ContainerCreating** state with the any of the below errors:

- MountDevice failed to create newCsiDriverClient: driver name openshift-storage.rbd.csi.ceph.com not found in the list of registered CSI drivers
- MountDevice failed for volume `<volume_name>` : rpc error: code = Aborted desc = an operation with the given Volume ID `<volume_id>` already exists
- MountVolume.SetUp failed for volume `<volume_name>` : rpc error: code = Internal desc = staging path `<path>` for volume `<volume_id>` is not a mountpoint

Resolution

If the workload pods are stuck with any of the above mentioned errors, perform the following workarounds:

- For **ceph-fs** workload stuck in **ContainerCreating**:
 1. Restart the nodes where the stuck pods are scheduled
 2. Delete these stuck pods
 3. Verify that the new pods are running
- For **ceph-rbd** workload stuck in **ContainerCreating** that do not self recover after sometime
 1. Restart csi-rbd plugin pods in the nodes where the stuck pods are scheduled
 2. Verify that the new pods are running